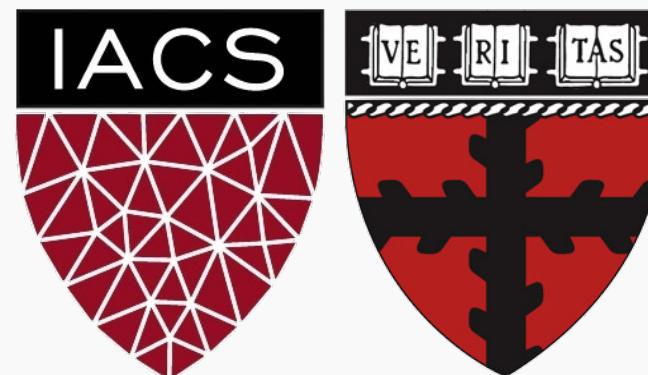


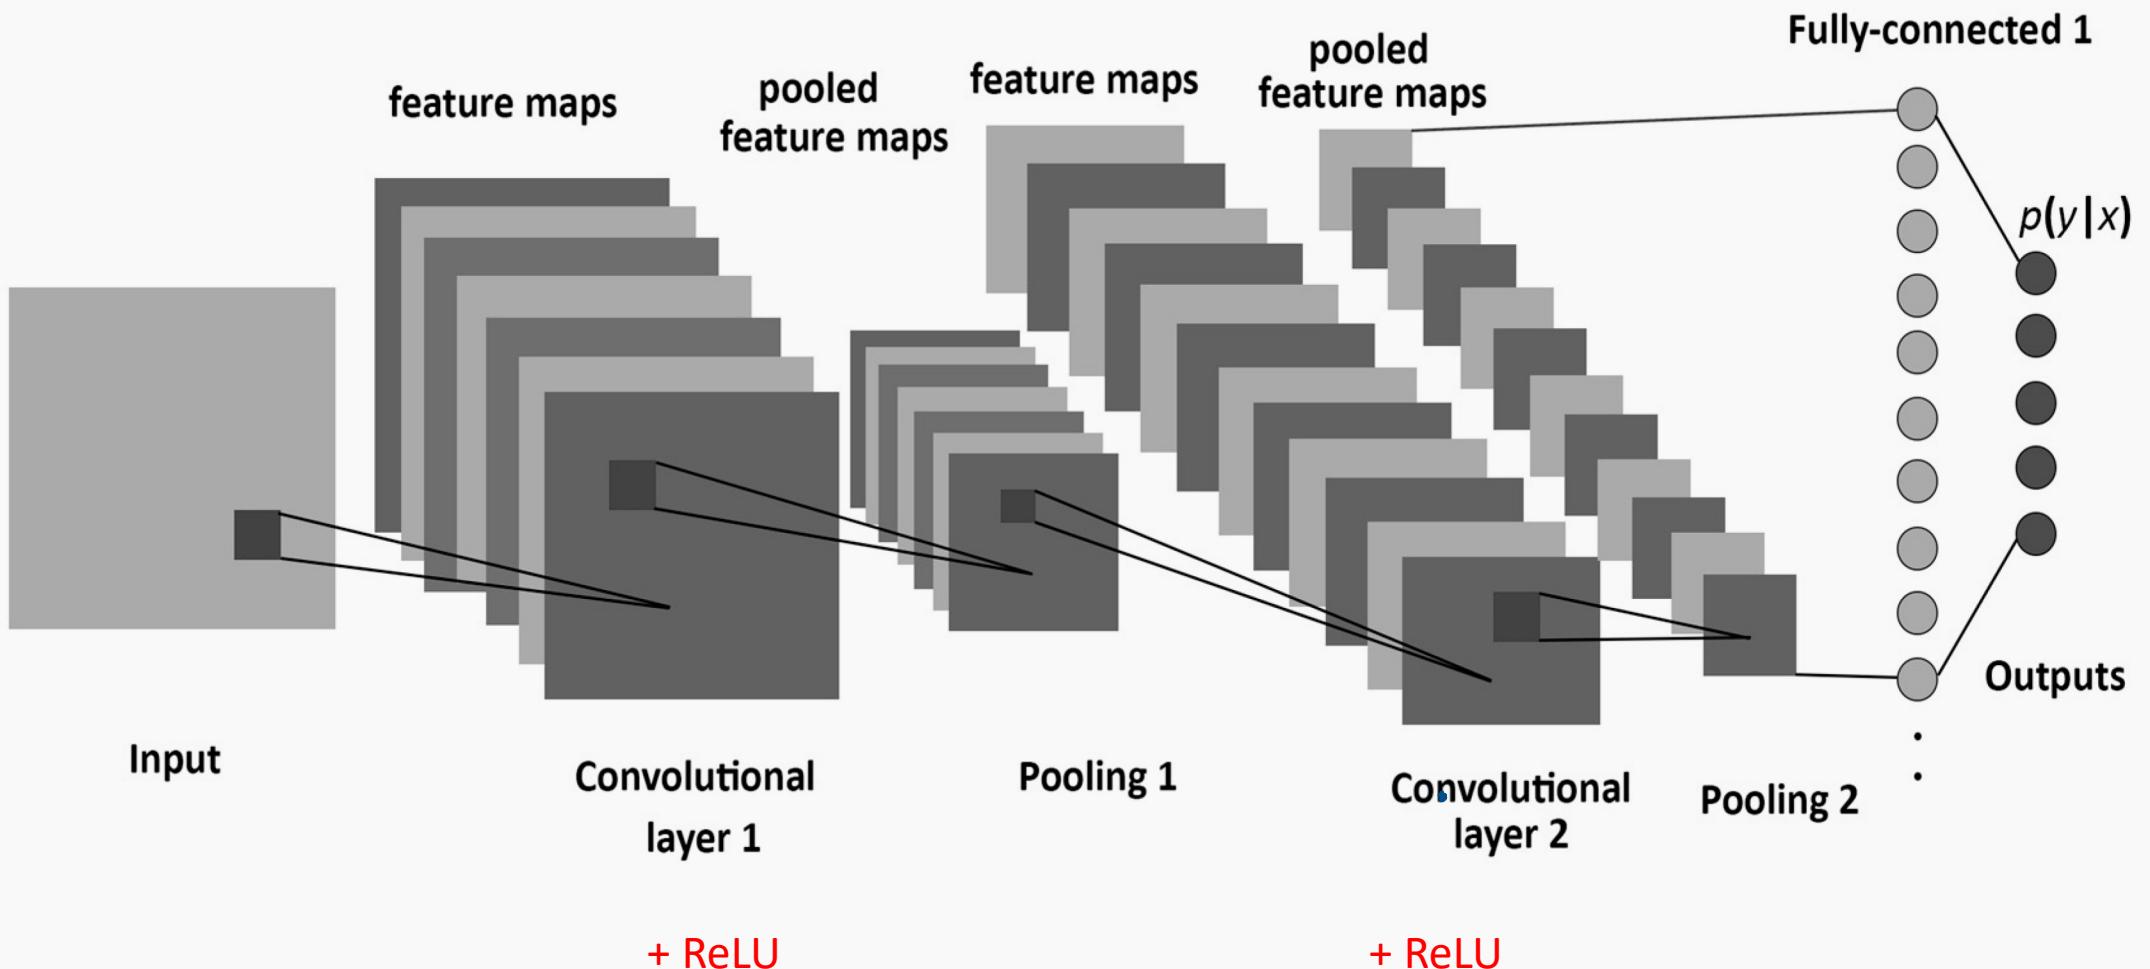
Convolutional Neural Networks 1

CS109B Data Science 2

Pavlos Protopapas, Mark Glickman, and Chris Tanner



A Convolutional Network



The code

In []:

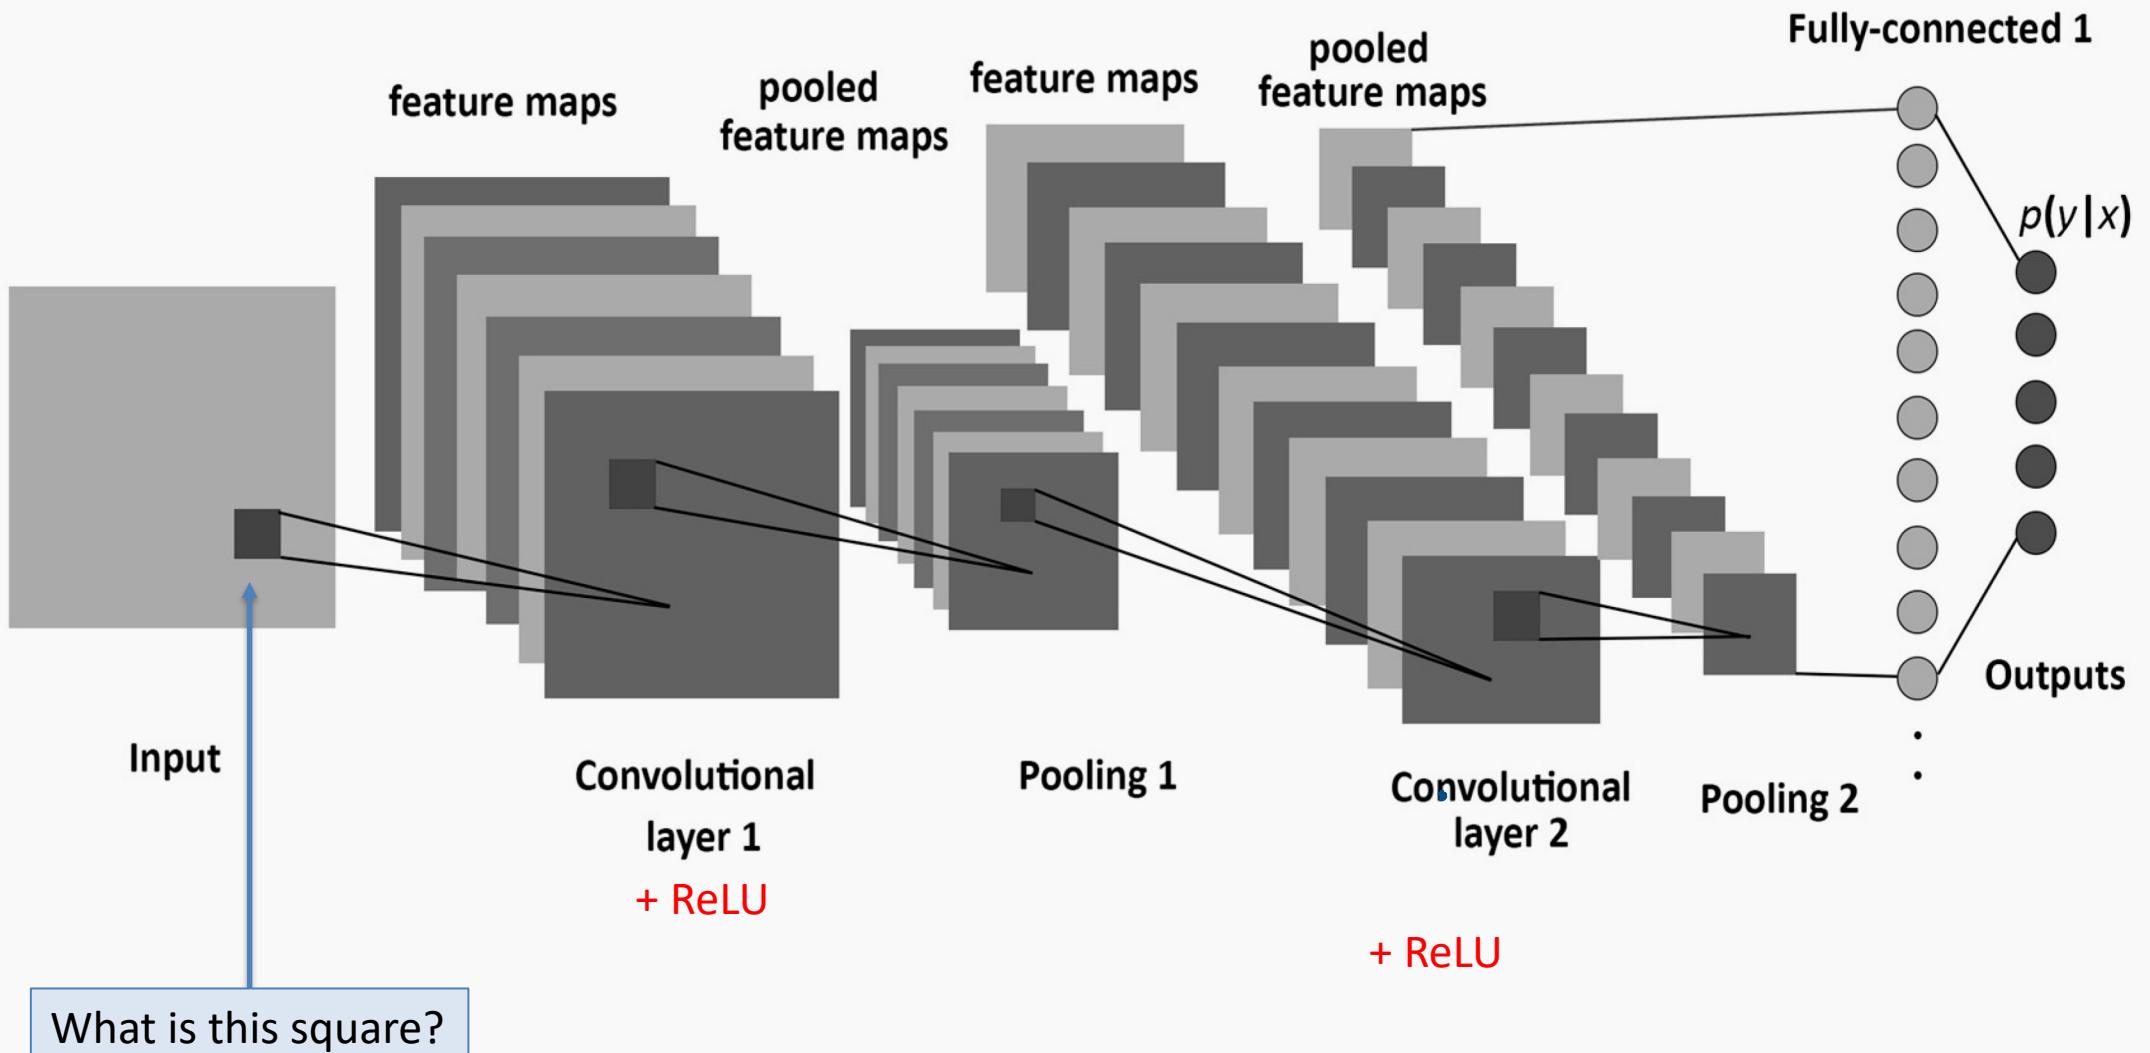
```
1 mnist_cnn_model = Sequential() # Create sequential model
2
3
4 # Add network layers
5 mnist_cnn_model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
6 mnist_cnn_model.add(MaxPooling2D((2, 2)))
7 mnist_cnn_model.add(Conv2D(64, (3, 3), activation='relu'))
8 mnist_cnn_model.add(MaxPooling2D((2, 2)))
9 mnist_cnn_model.add(Conv2D(64, (3, 3), activation='relu'))
10
11 mnist_cnn_model.add(Flatten())
12 mnist_cnn_model.add(Dense(64, activation='relu'))
13
14 mnist_cnn_model.add(Dense(10, activation='softmax'))
15
16 mnist_cnn_model.compile(optimizer=optimizer,
17                         loss=loss,
18                         metrics=metrics)
19
20 history = mnist_cnn_model.fit(train_images, train_labels,
21                               epochs=epochs,
22                               batch_size=batch_size,
23                               verbose=verbose,
24                               validation_split=0.2,
25                               # validation_data=(X_val, y_val) # IF you have val data
26                               shuffle=True)
```



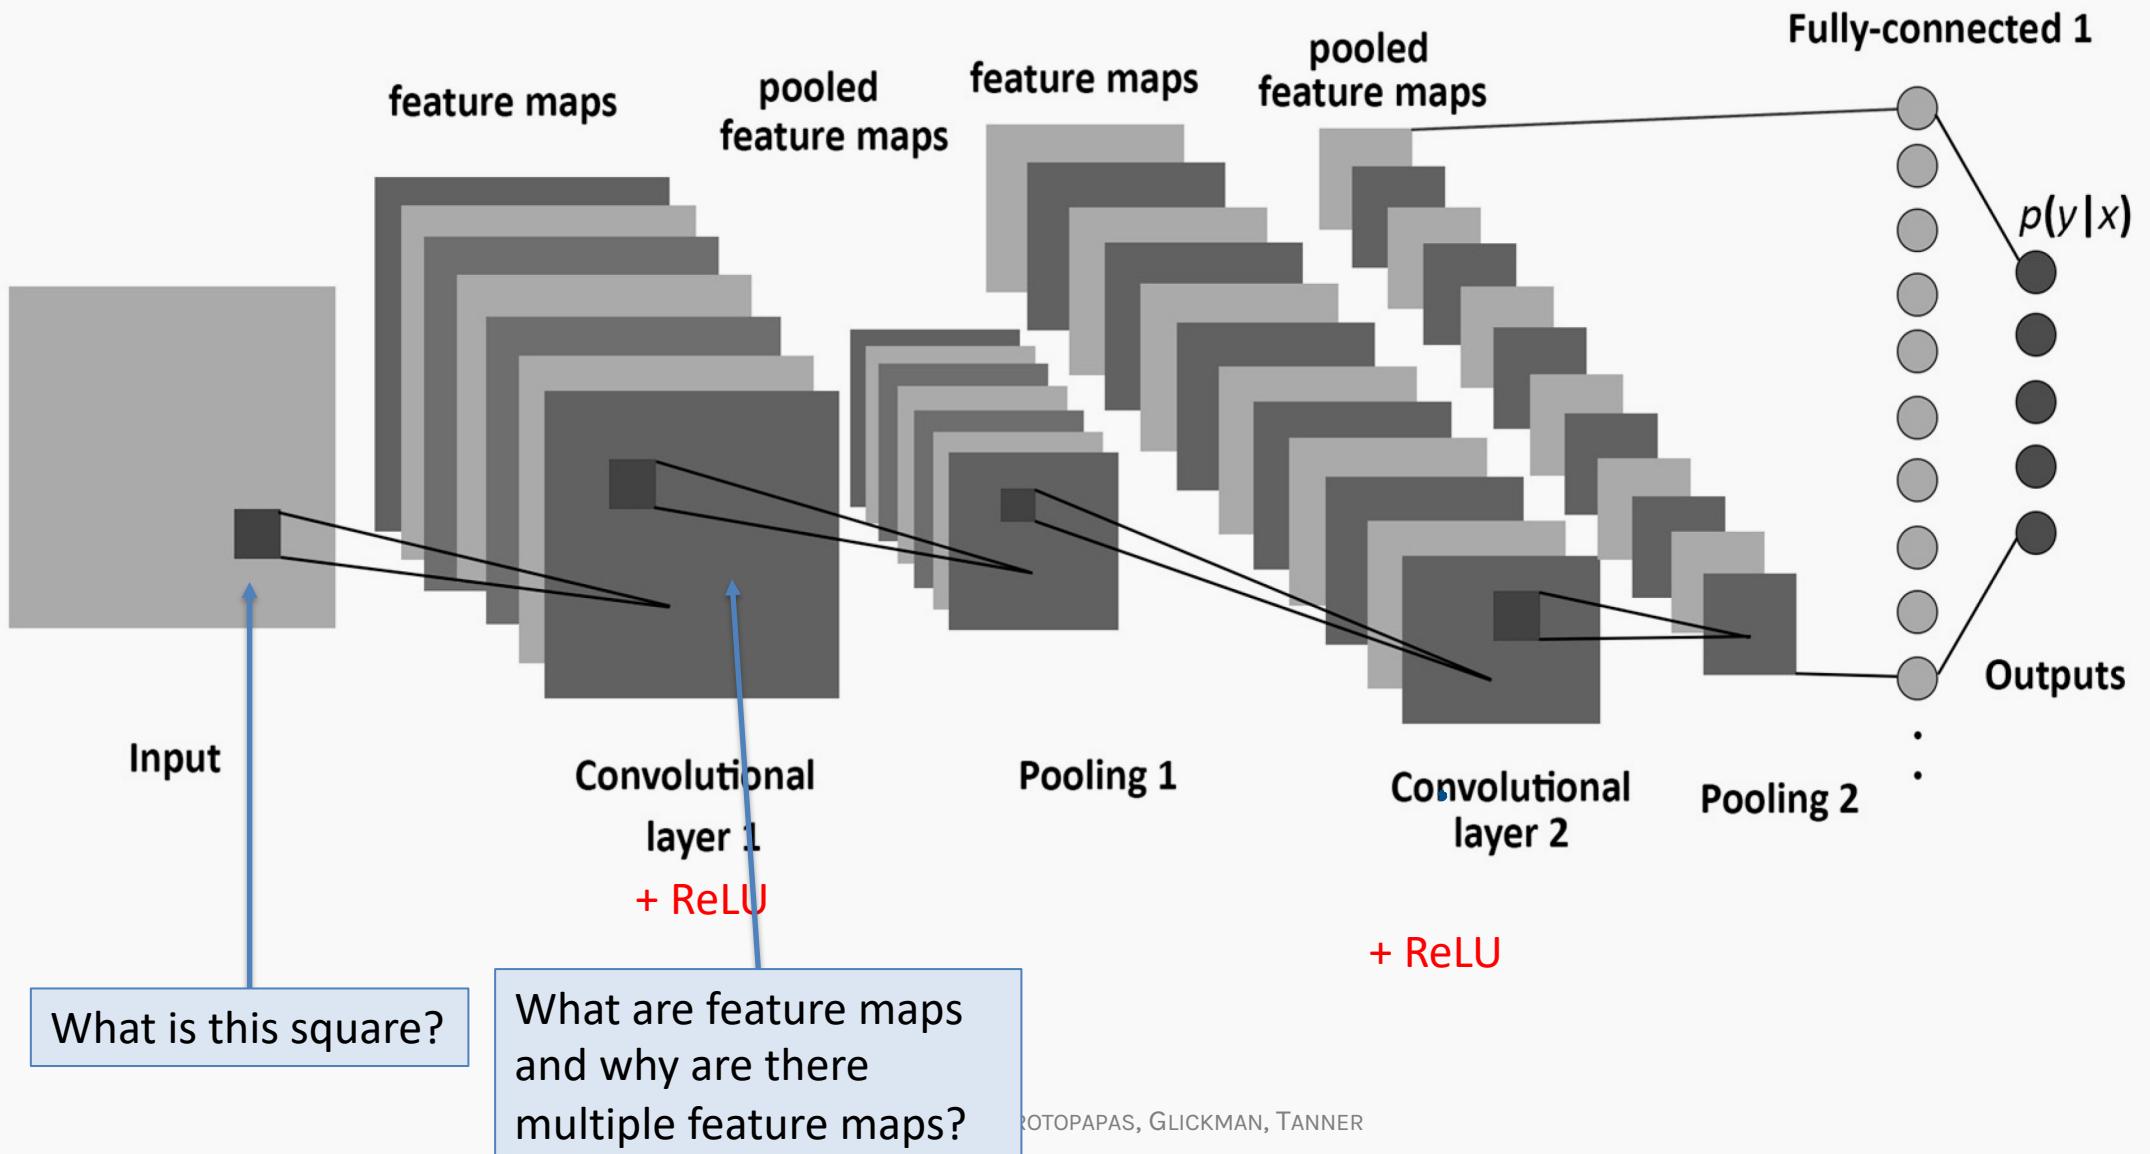
DONE



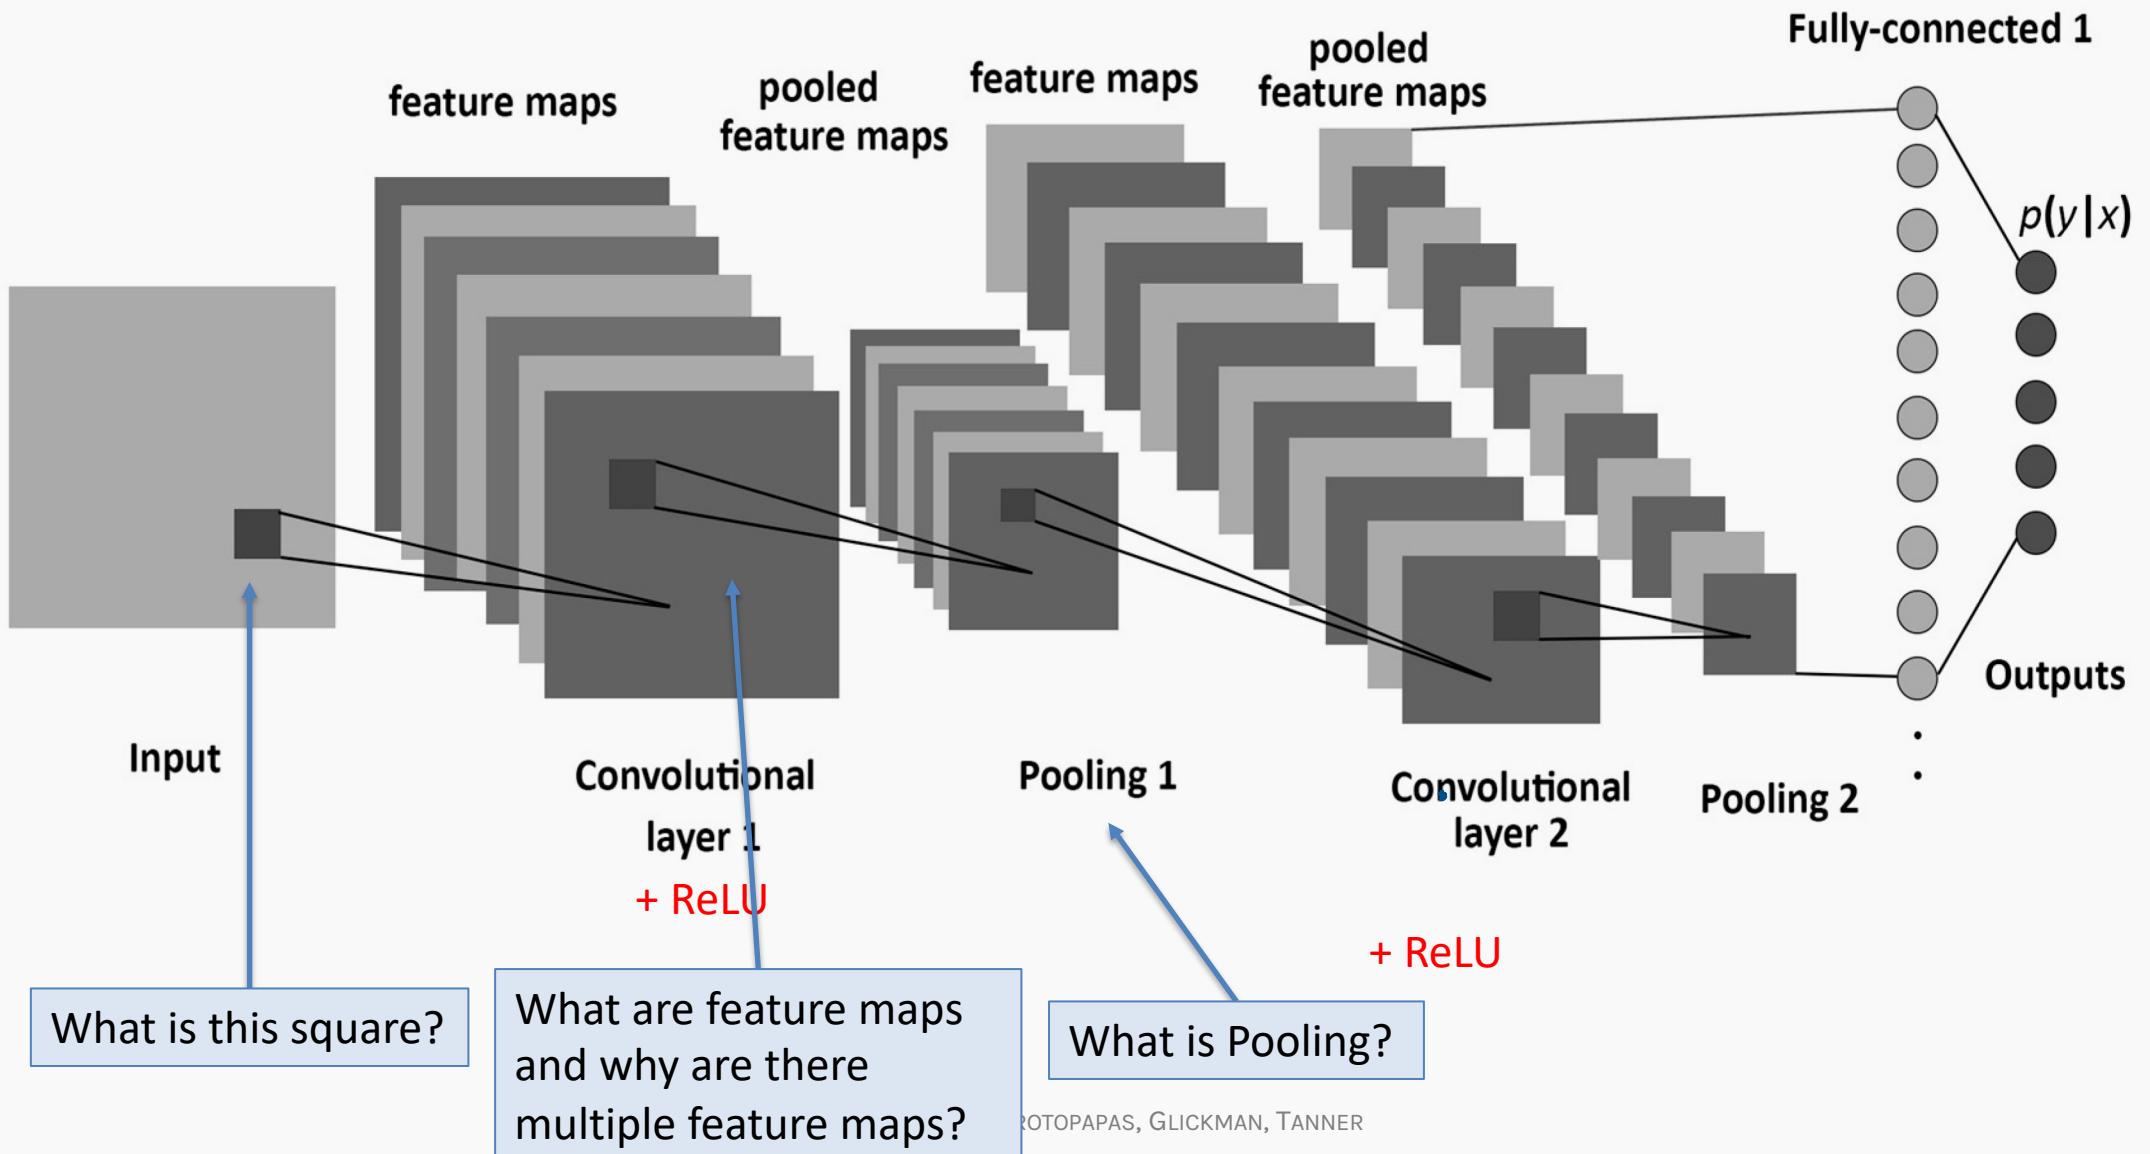
A Convolutional Network



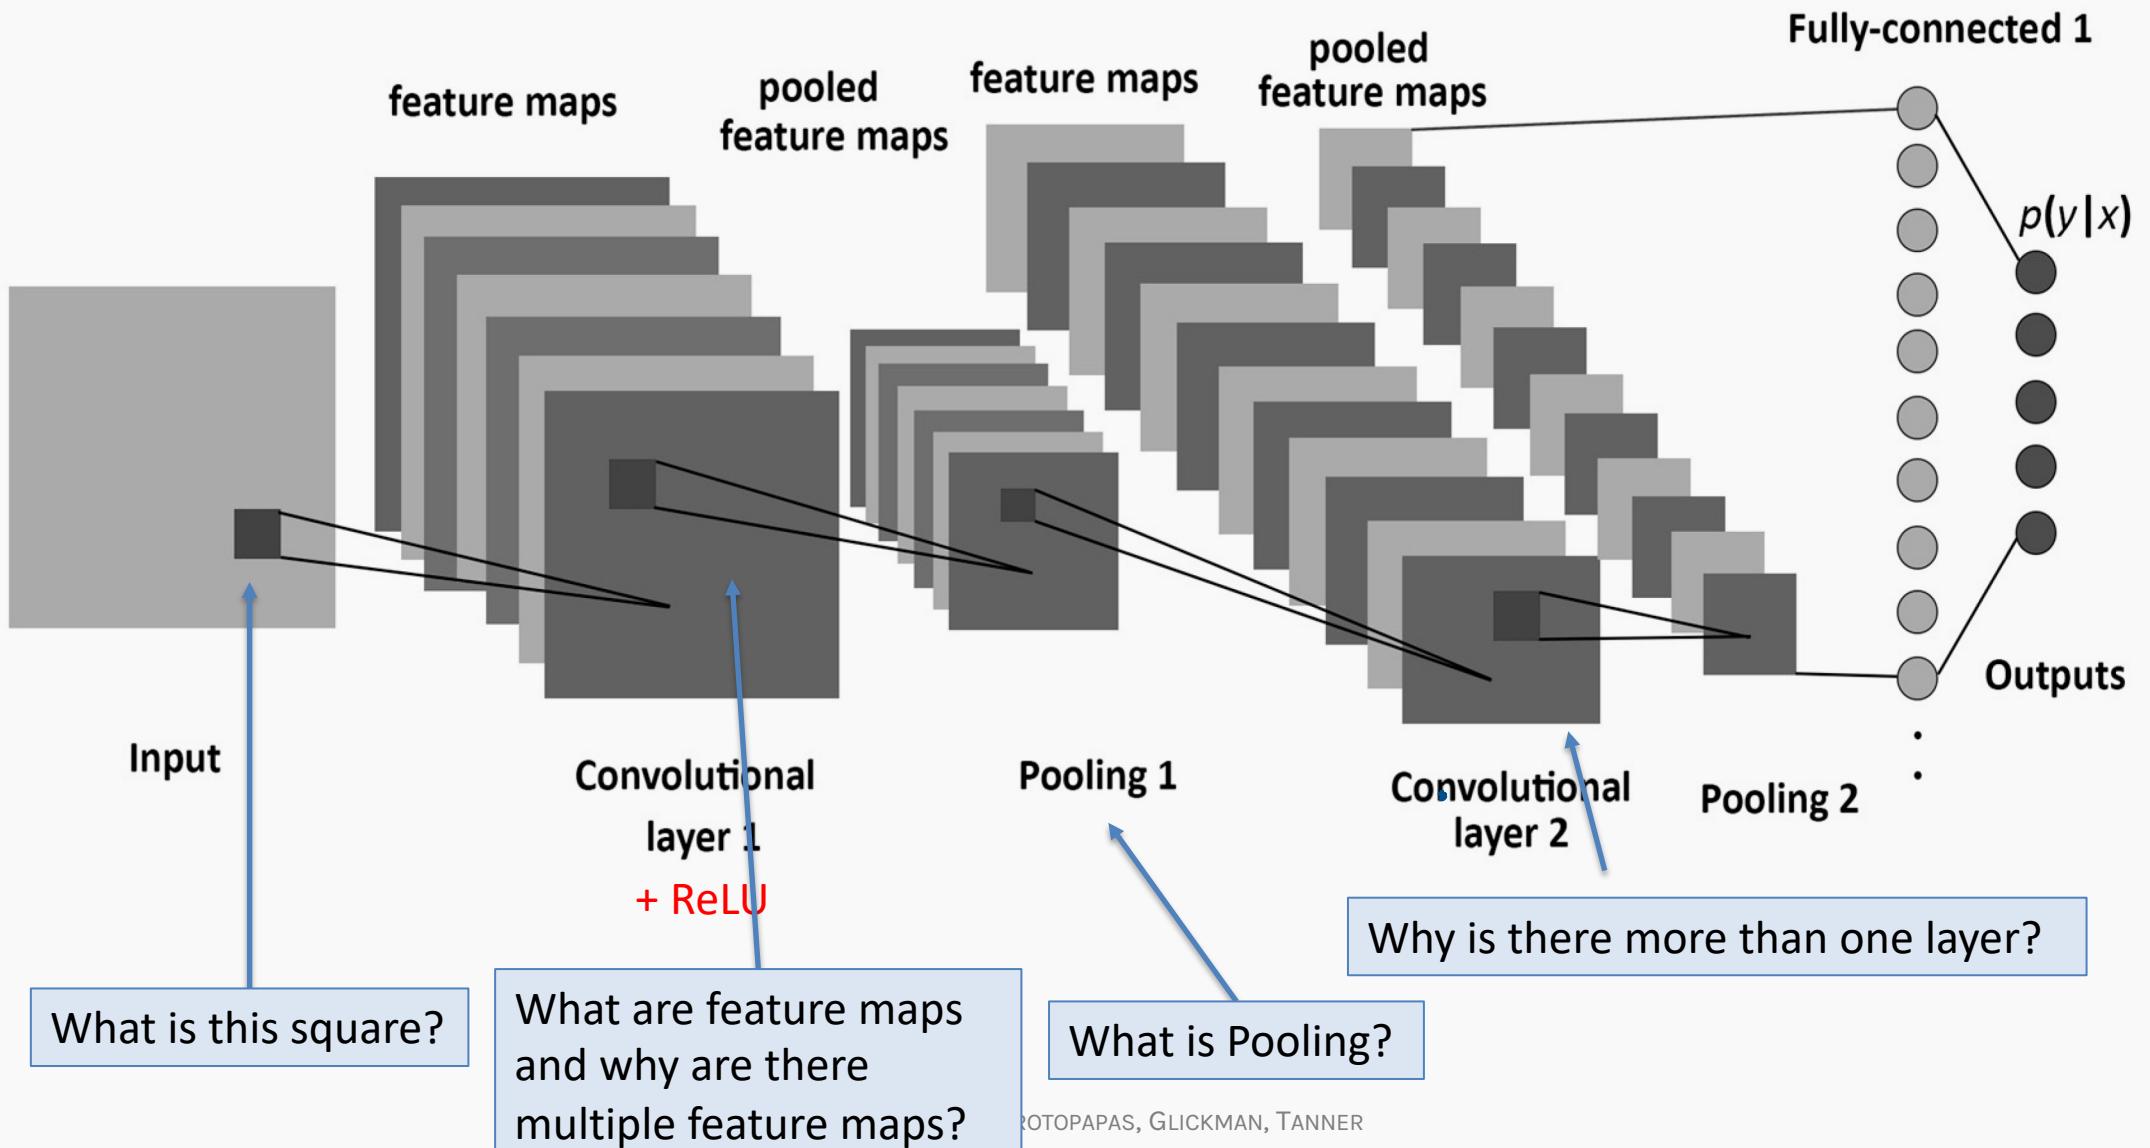
A Convolutional Network



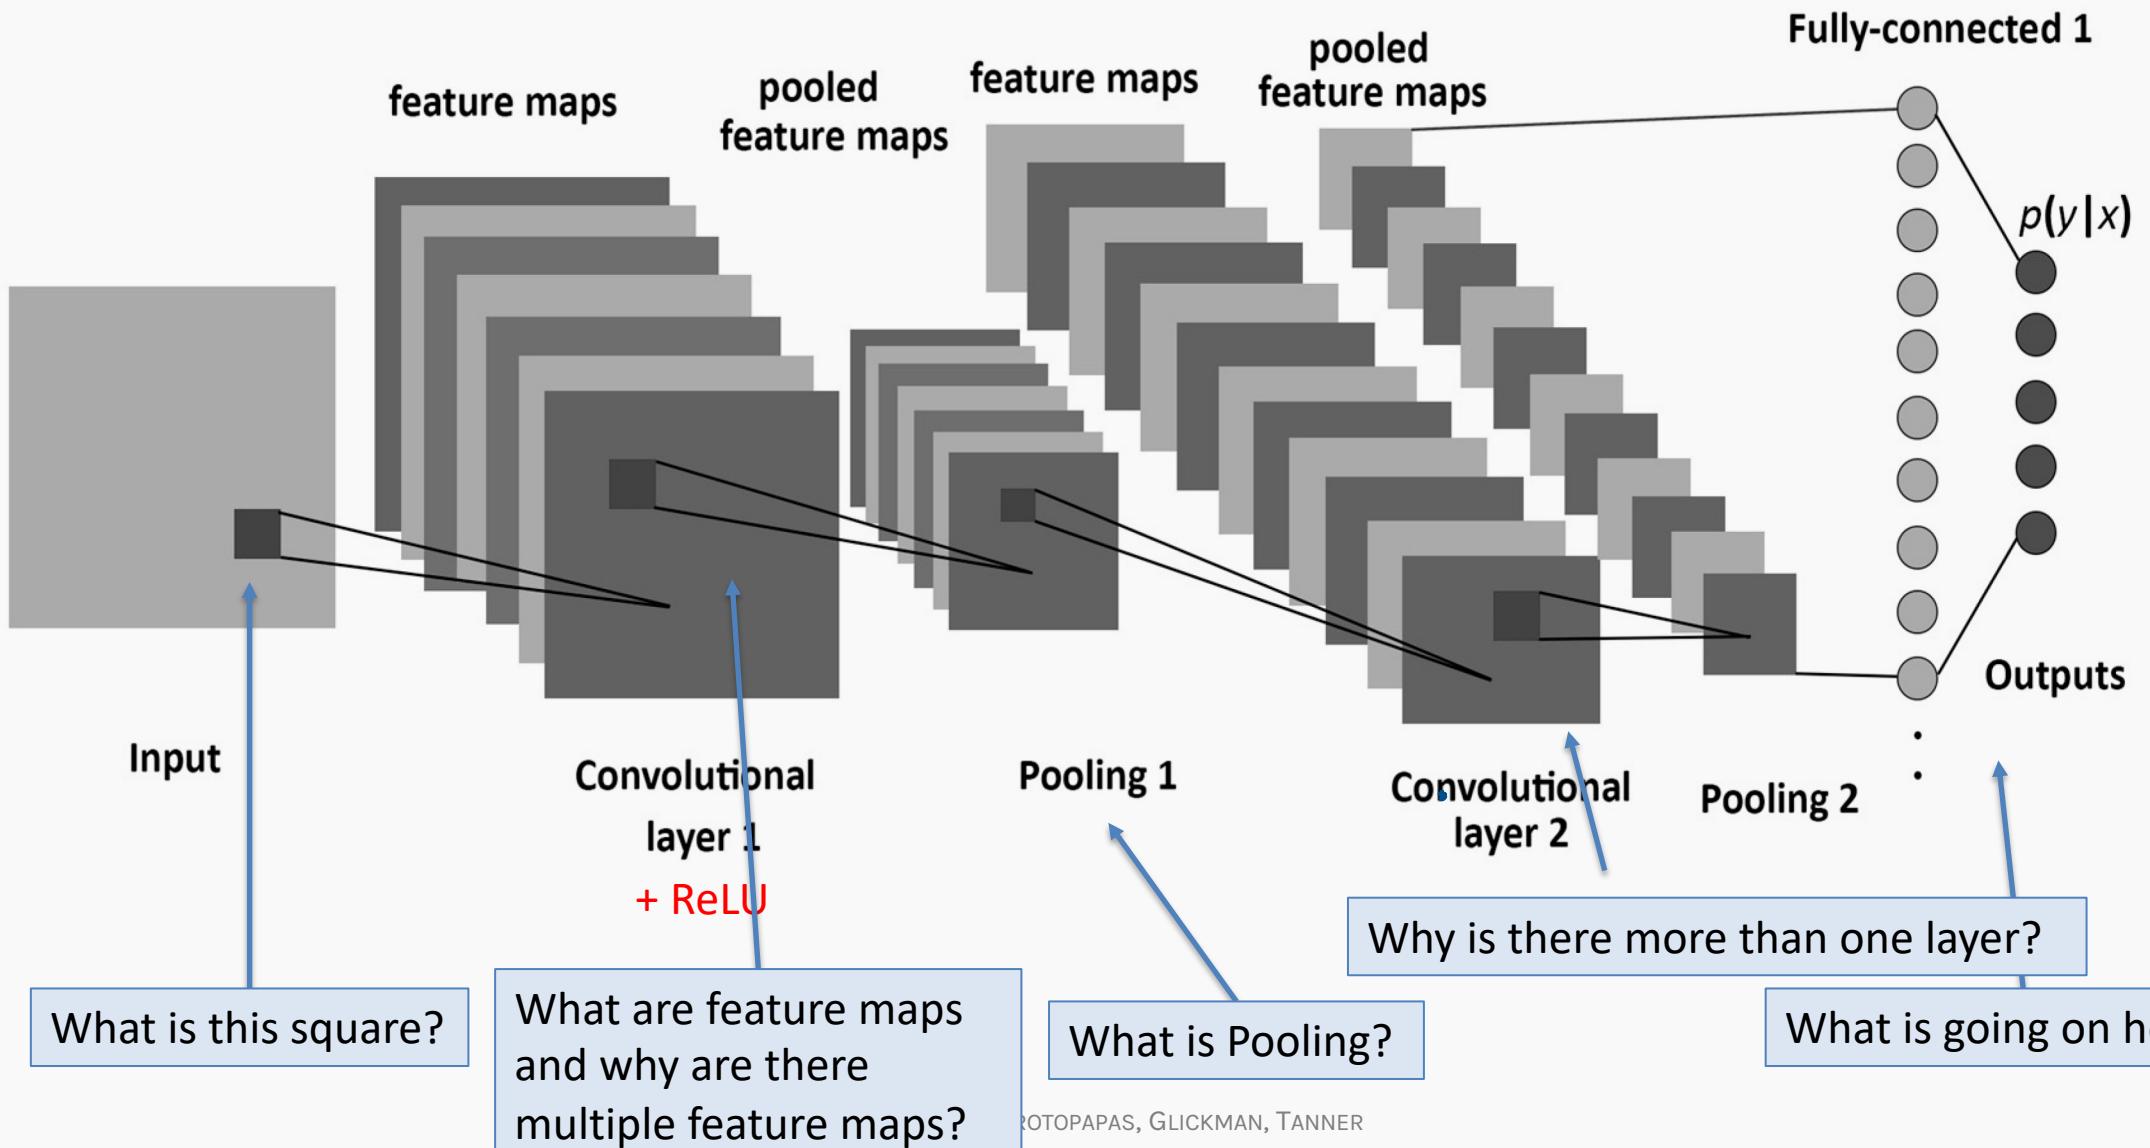
A Convolutional Network



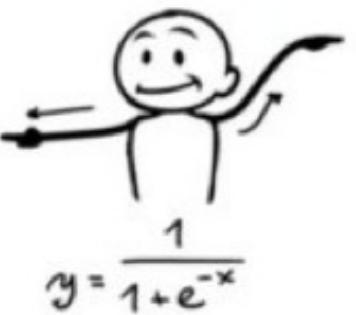
A Convolutional Network



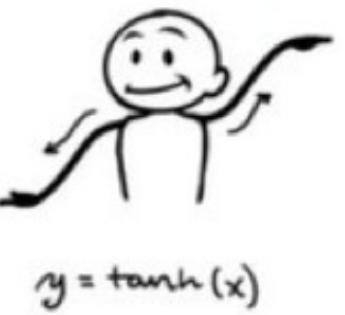
A Convolutional Network



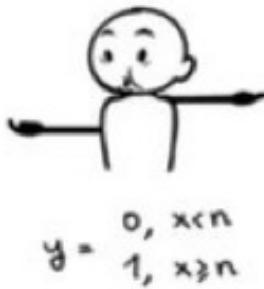
Sigmoid



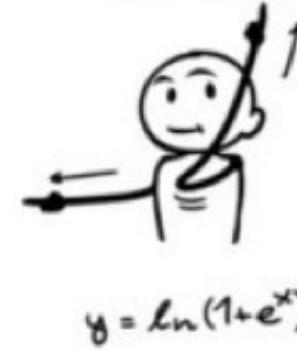
Tanh



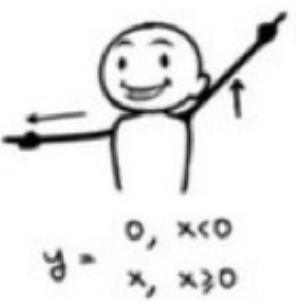
Step Function



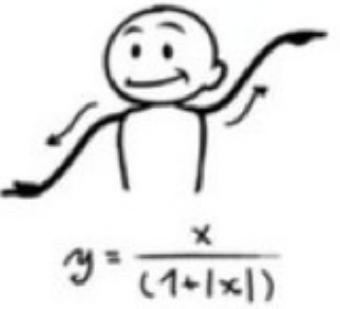
Softplus



ReLU



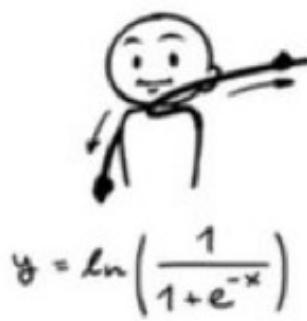
Softsign



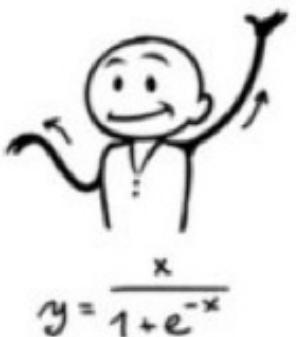
ELU



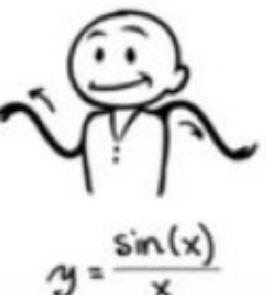
Log of Sigmoid



Swish



Sinc



Leaky ReLU



Mish



Outline

1. Motivation
2. CNN basic ideas
3. Building a CNN



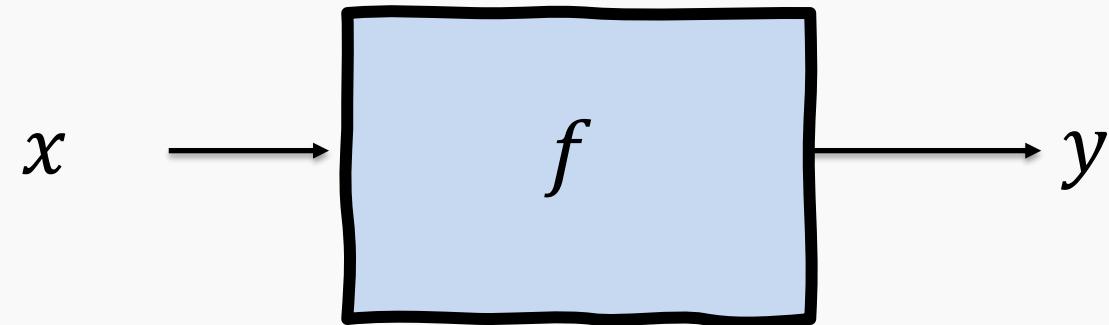
Outline

1. Motivation
2. CNN basic ideas
3. Building a CNN



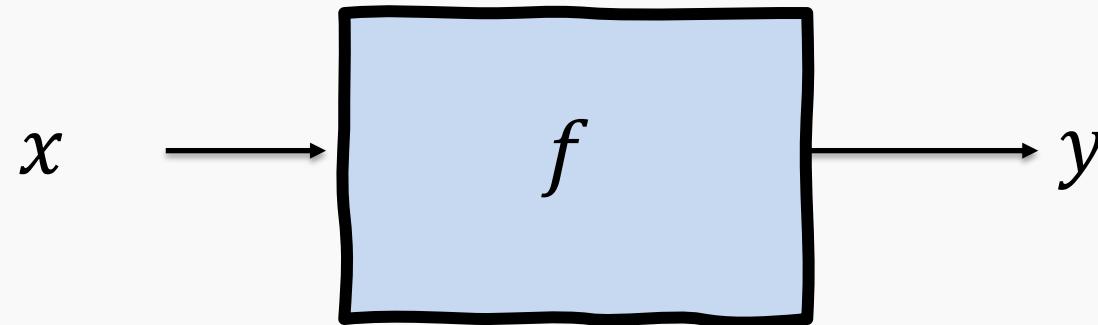
Feed forward Neural Network, Multilayer Perceptron (MLP)

A **function** is a relation that associates each element x of a set X to a single element y of a set Y

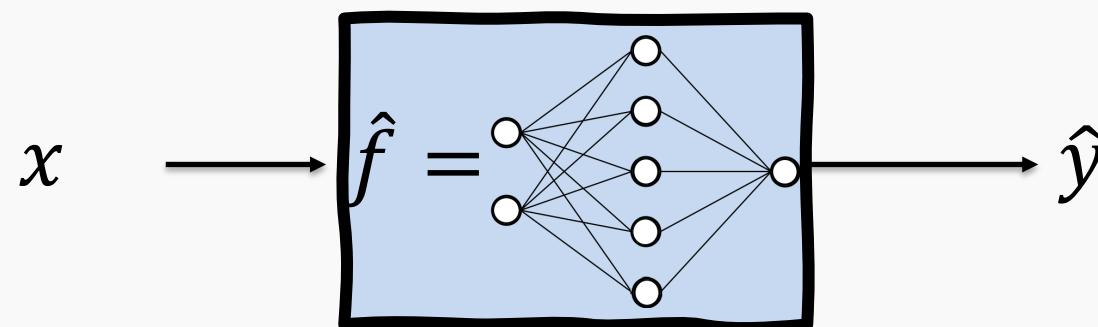


Feed forward Neural Network, Multilayer Perceptron (MLP)

A **function** is a relation that associates each element x of a set X to a single element y of a set Y



Neural networks can approximate a wide variety of functions



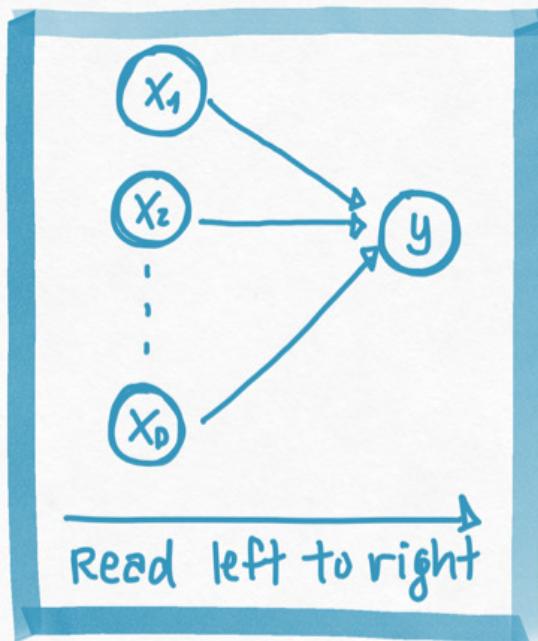
Graphical representation of simple functions

We build these complex functions by composing simple functions of the form:

$$h_w(x) = f(XW + b)$$

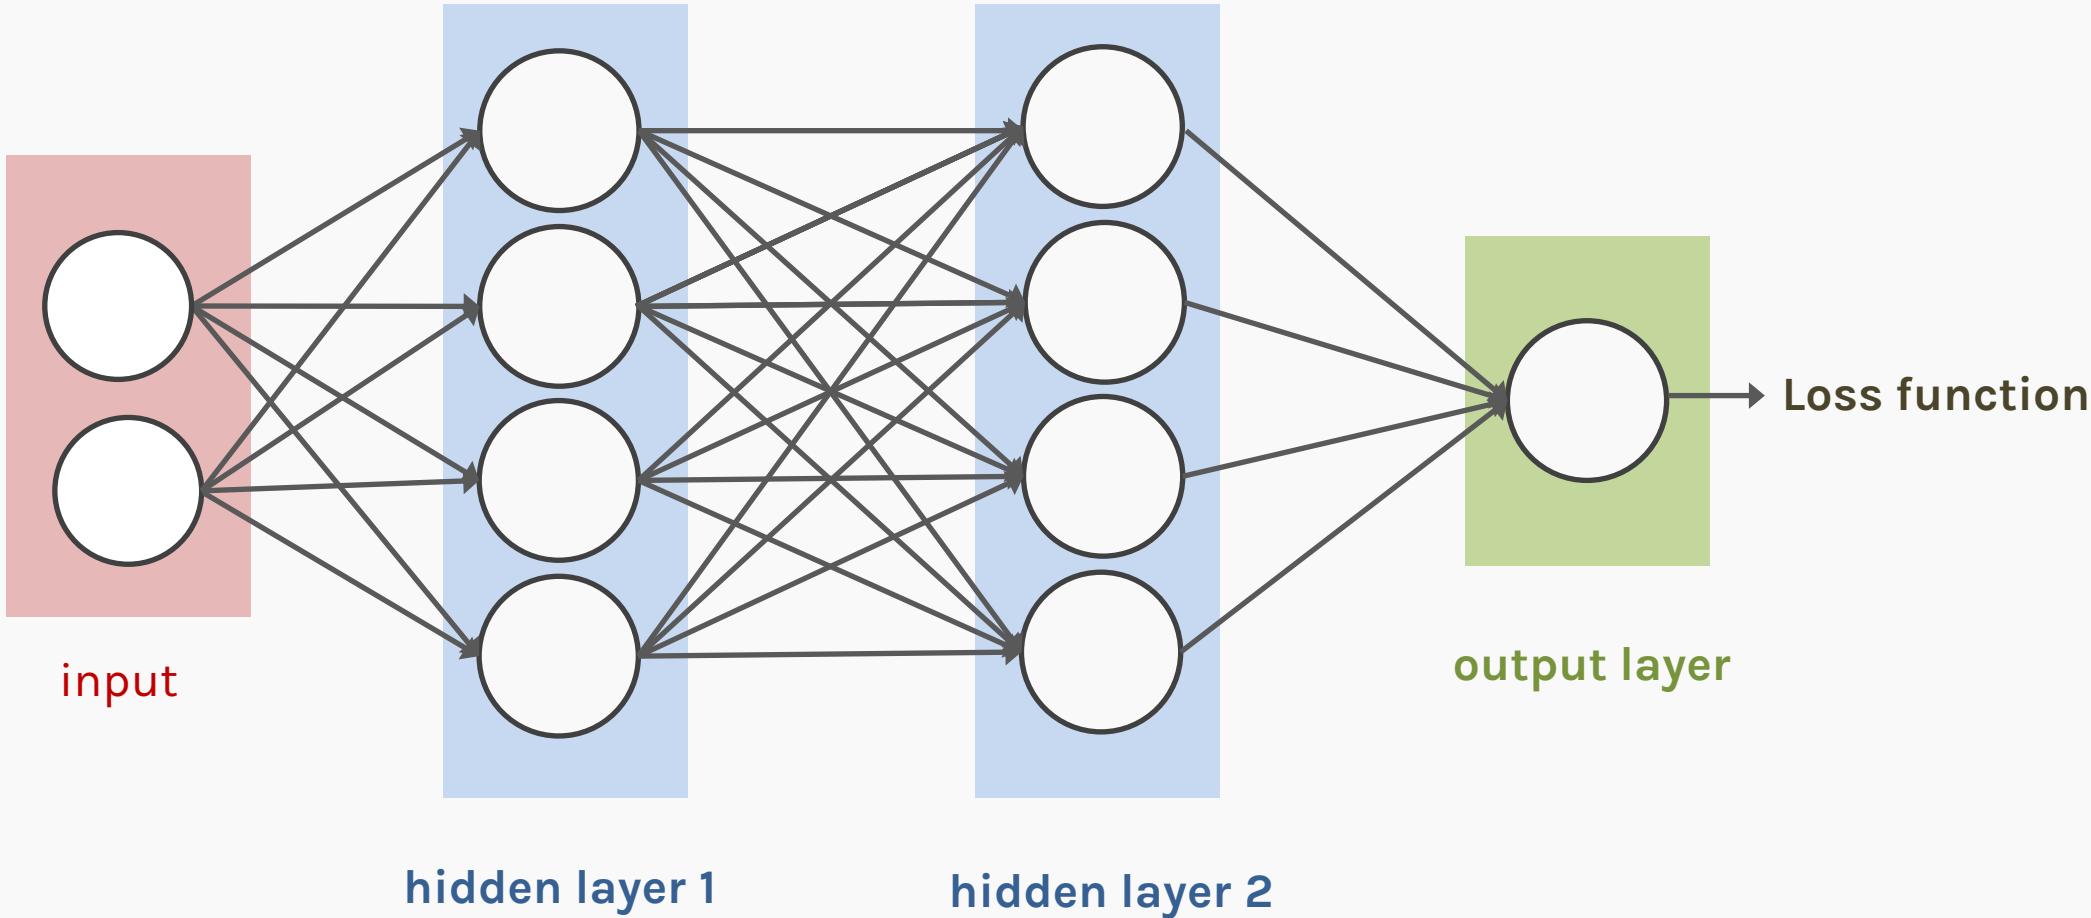
where f is the activation function.

We represent our simple function as a **graph**



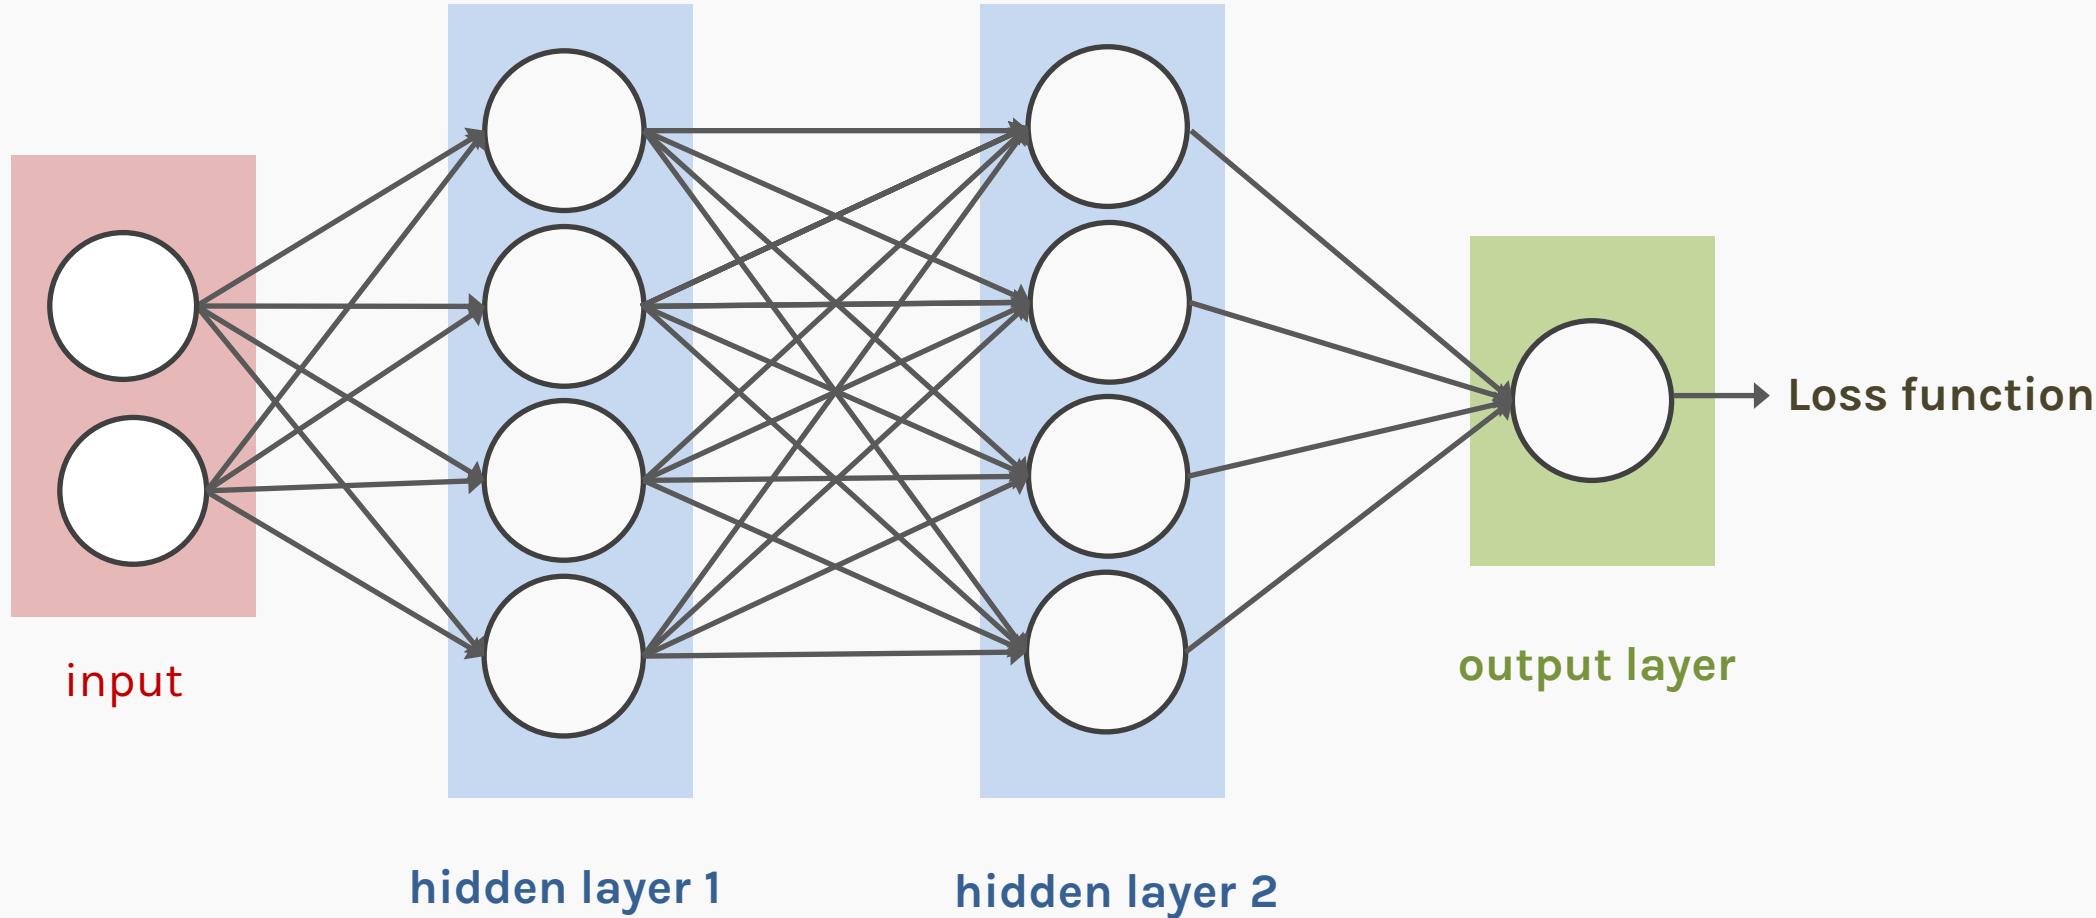
Each edge in this graph represents multiplication by a different weight, w_i .

Quick review of MLPs

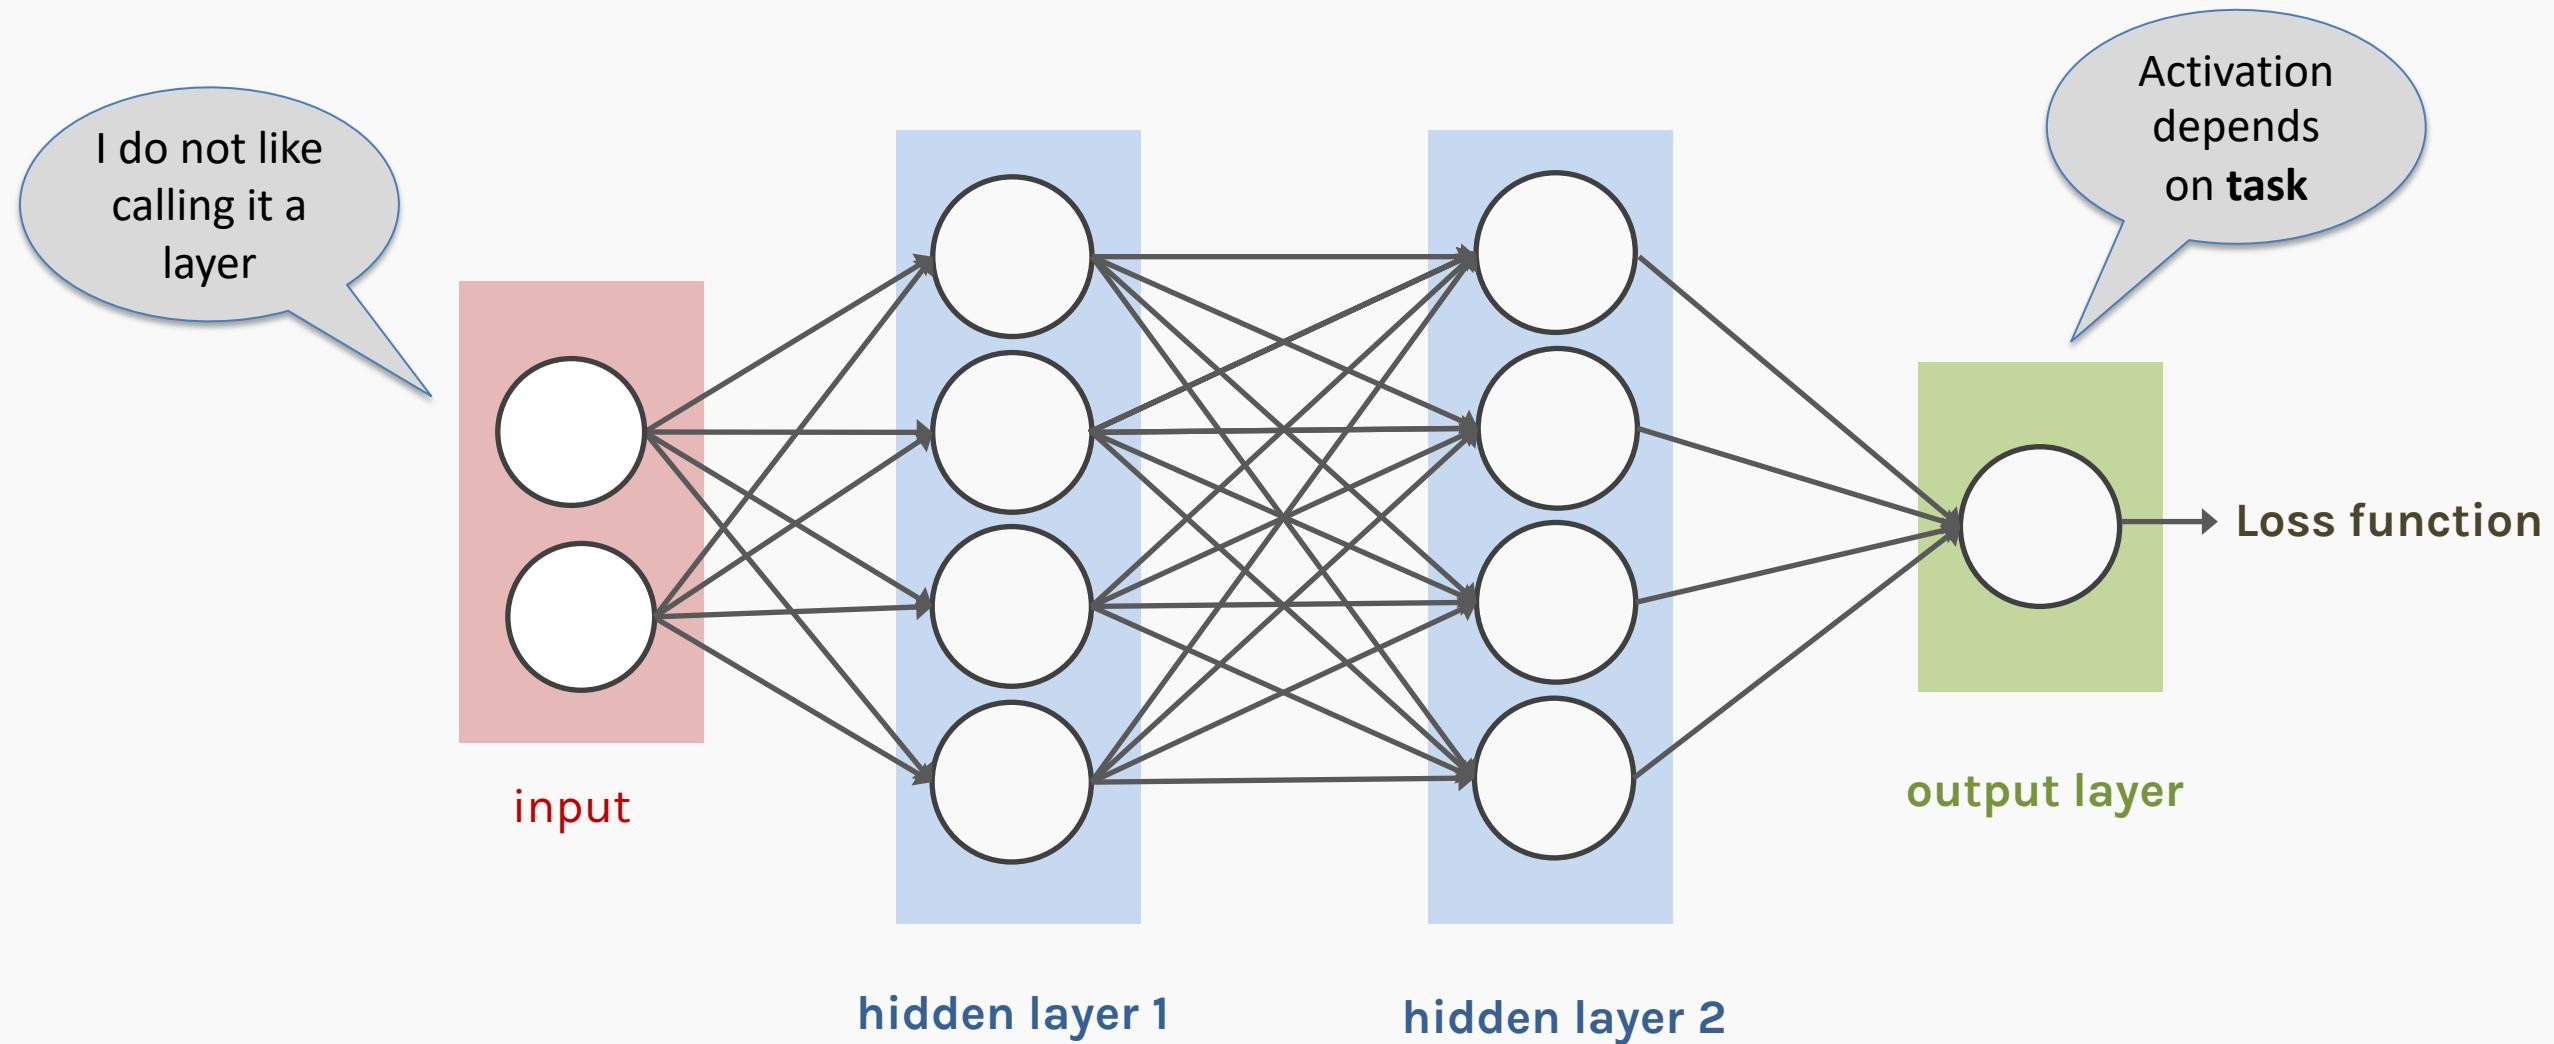


Quick review of MLPs

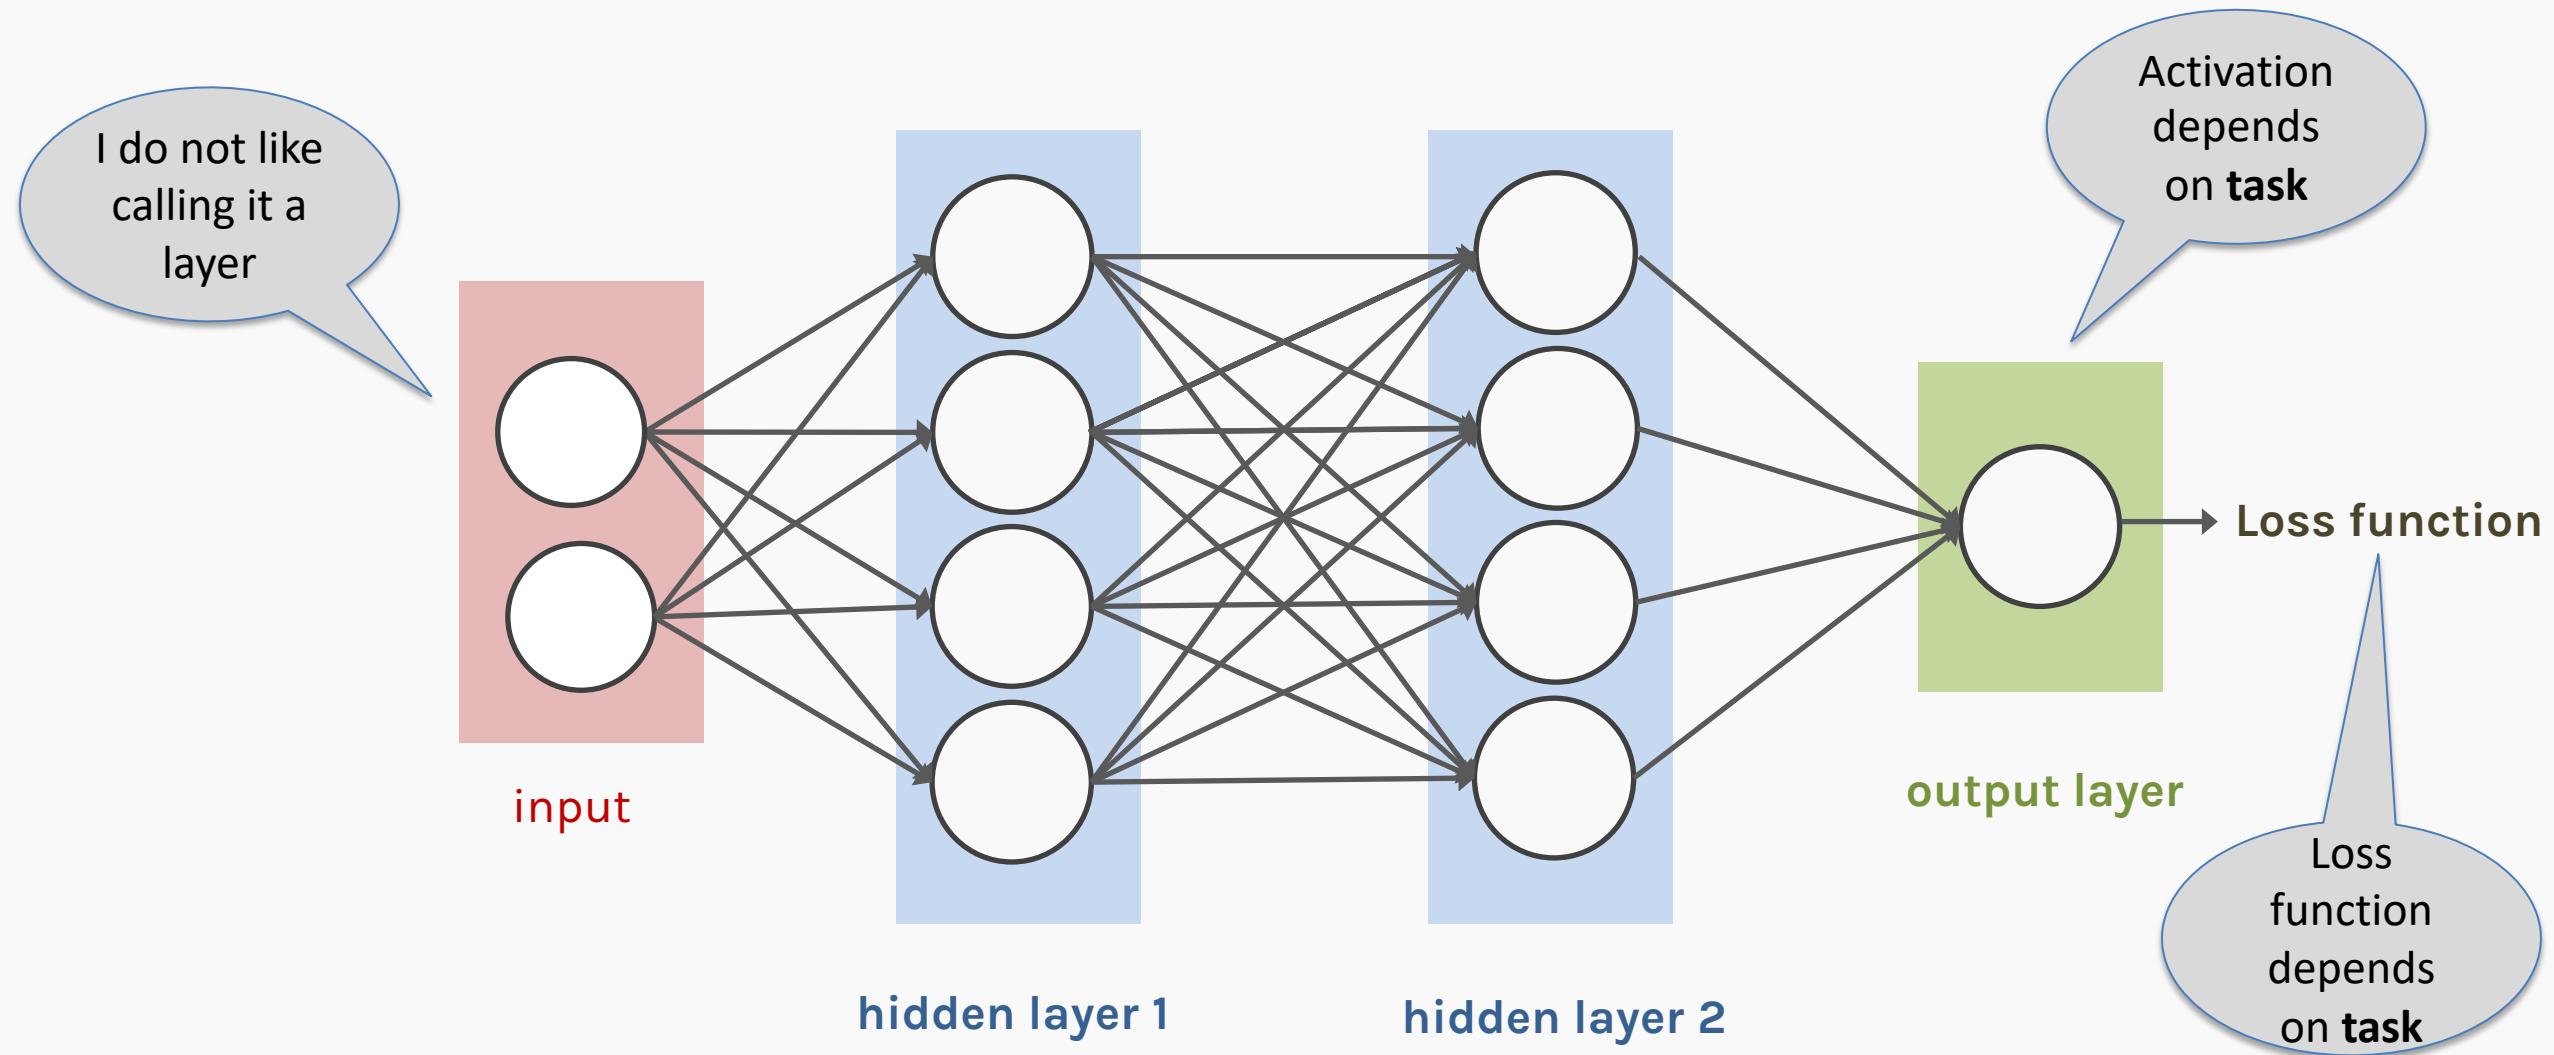
I do not like
calling it a
layer



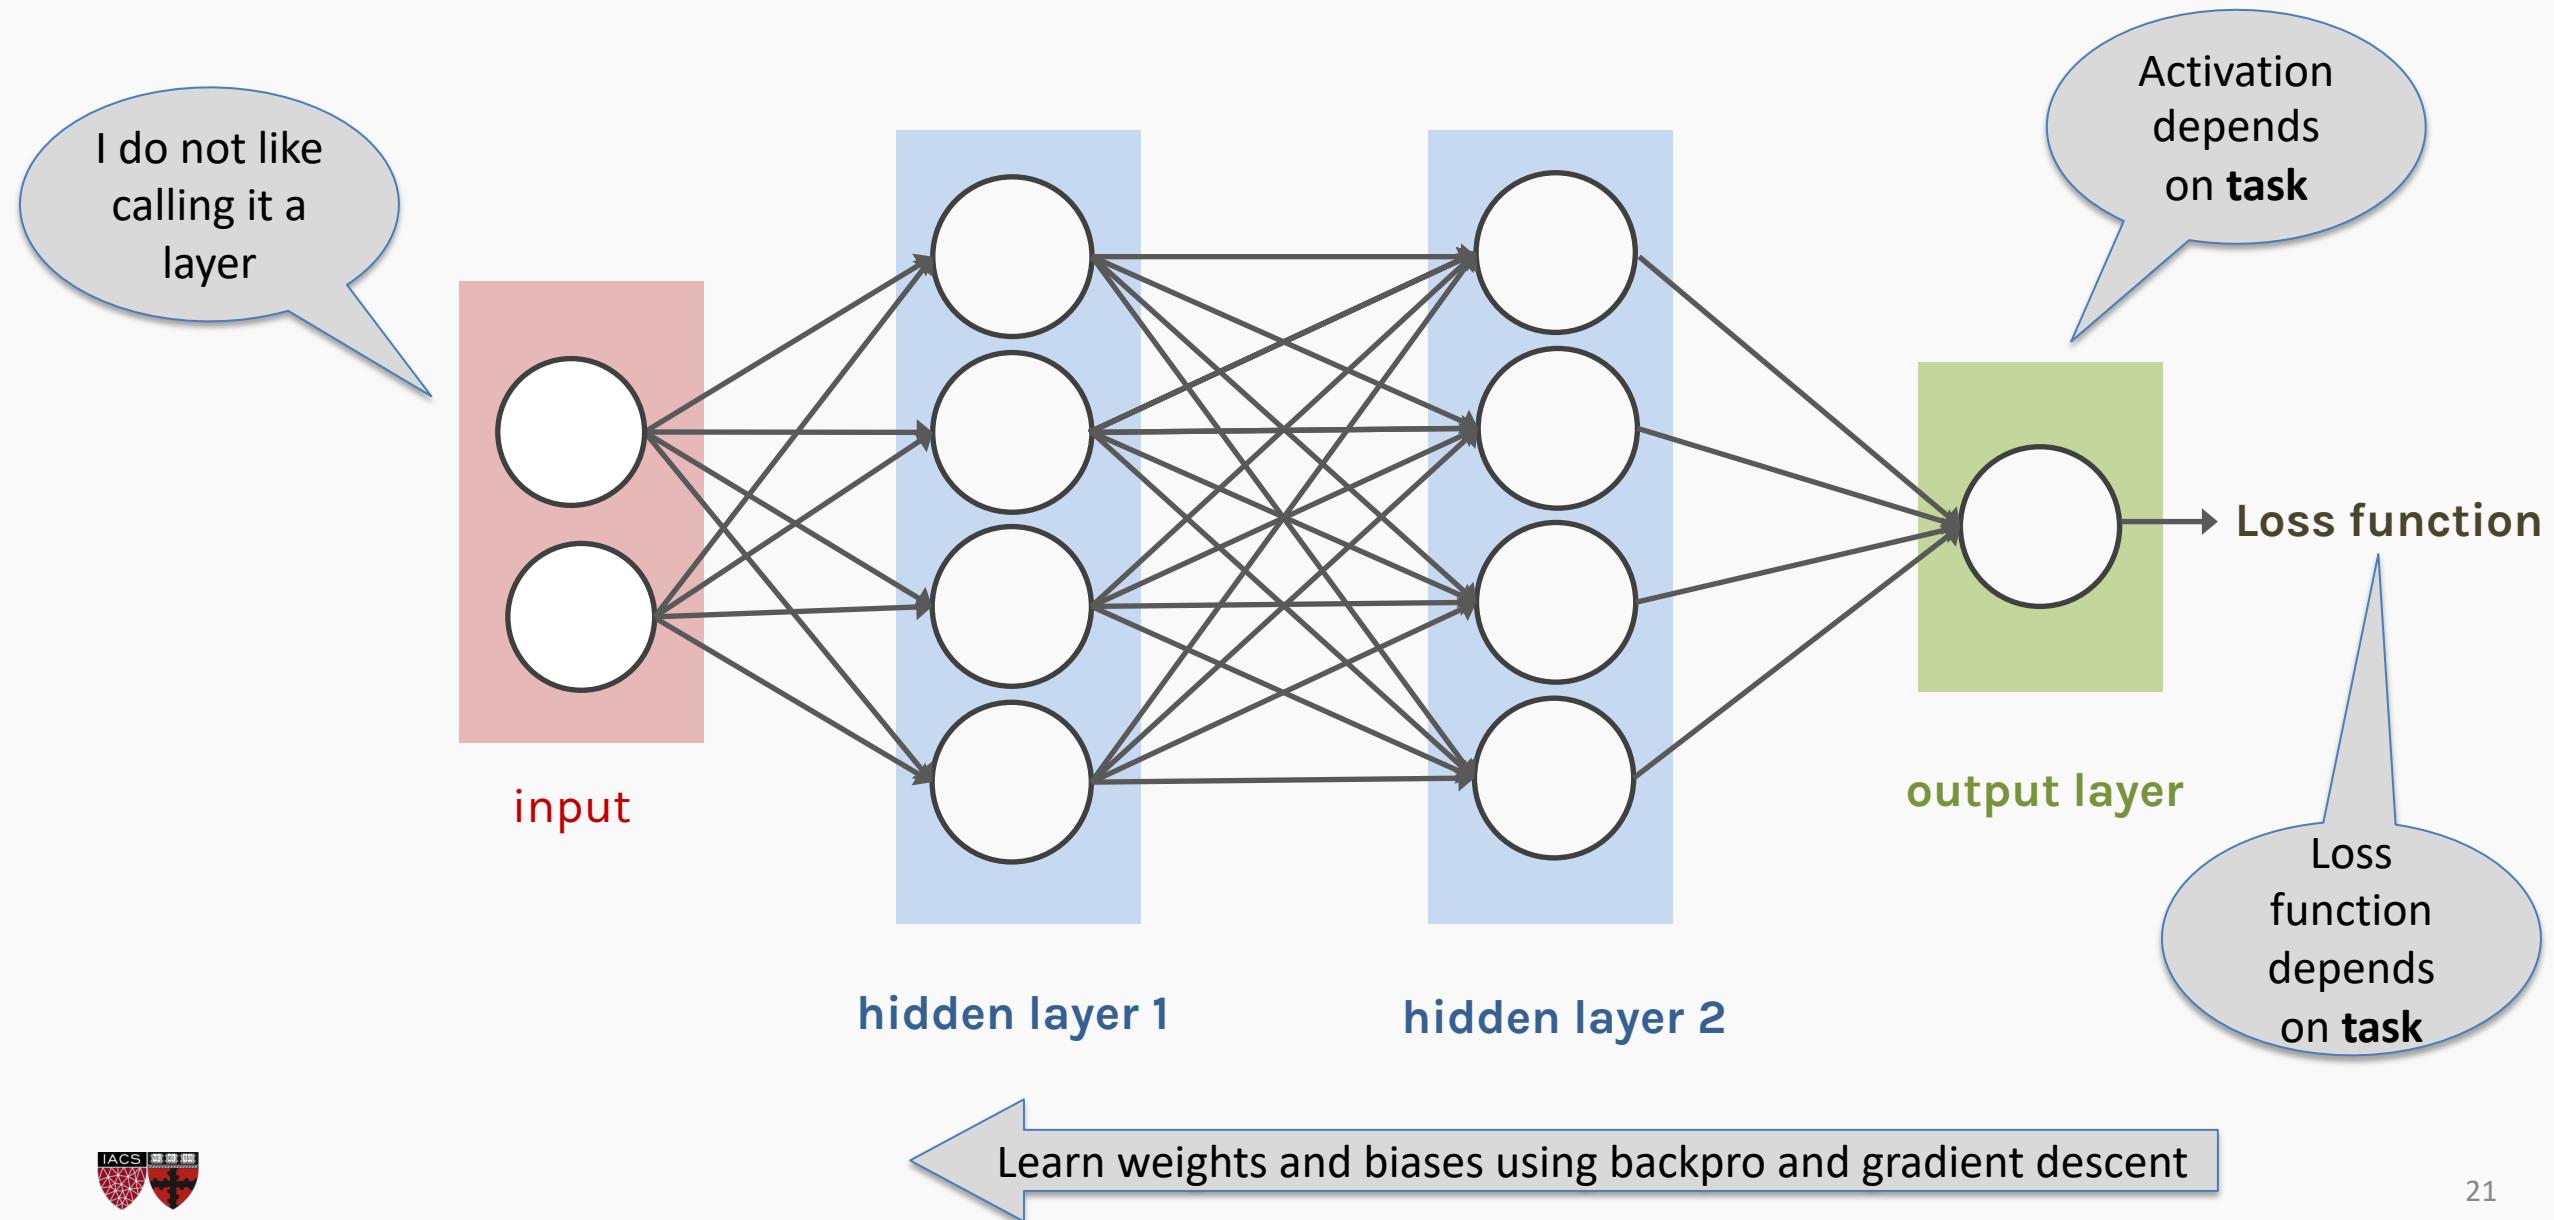
Quick review of MLPs



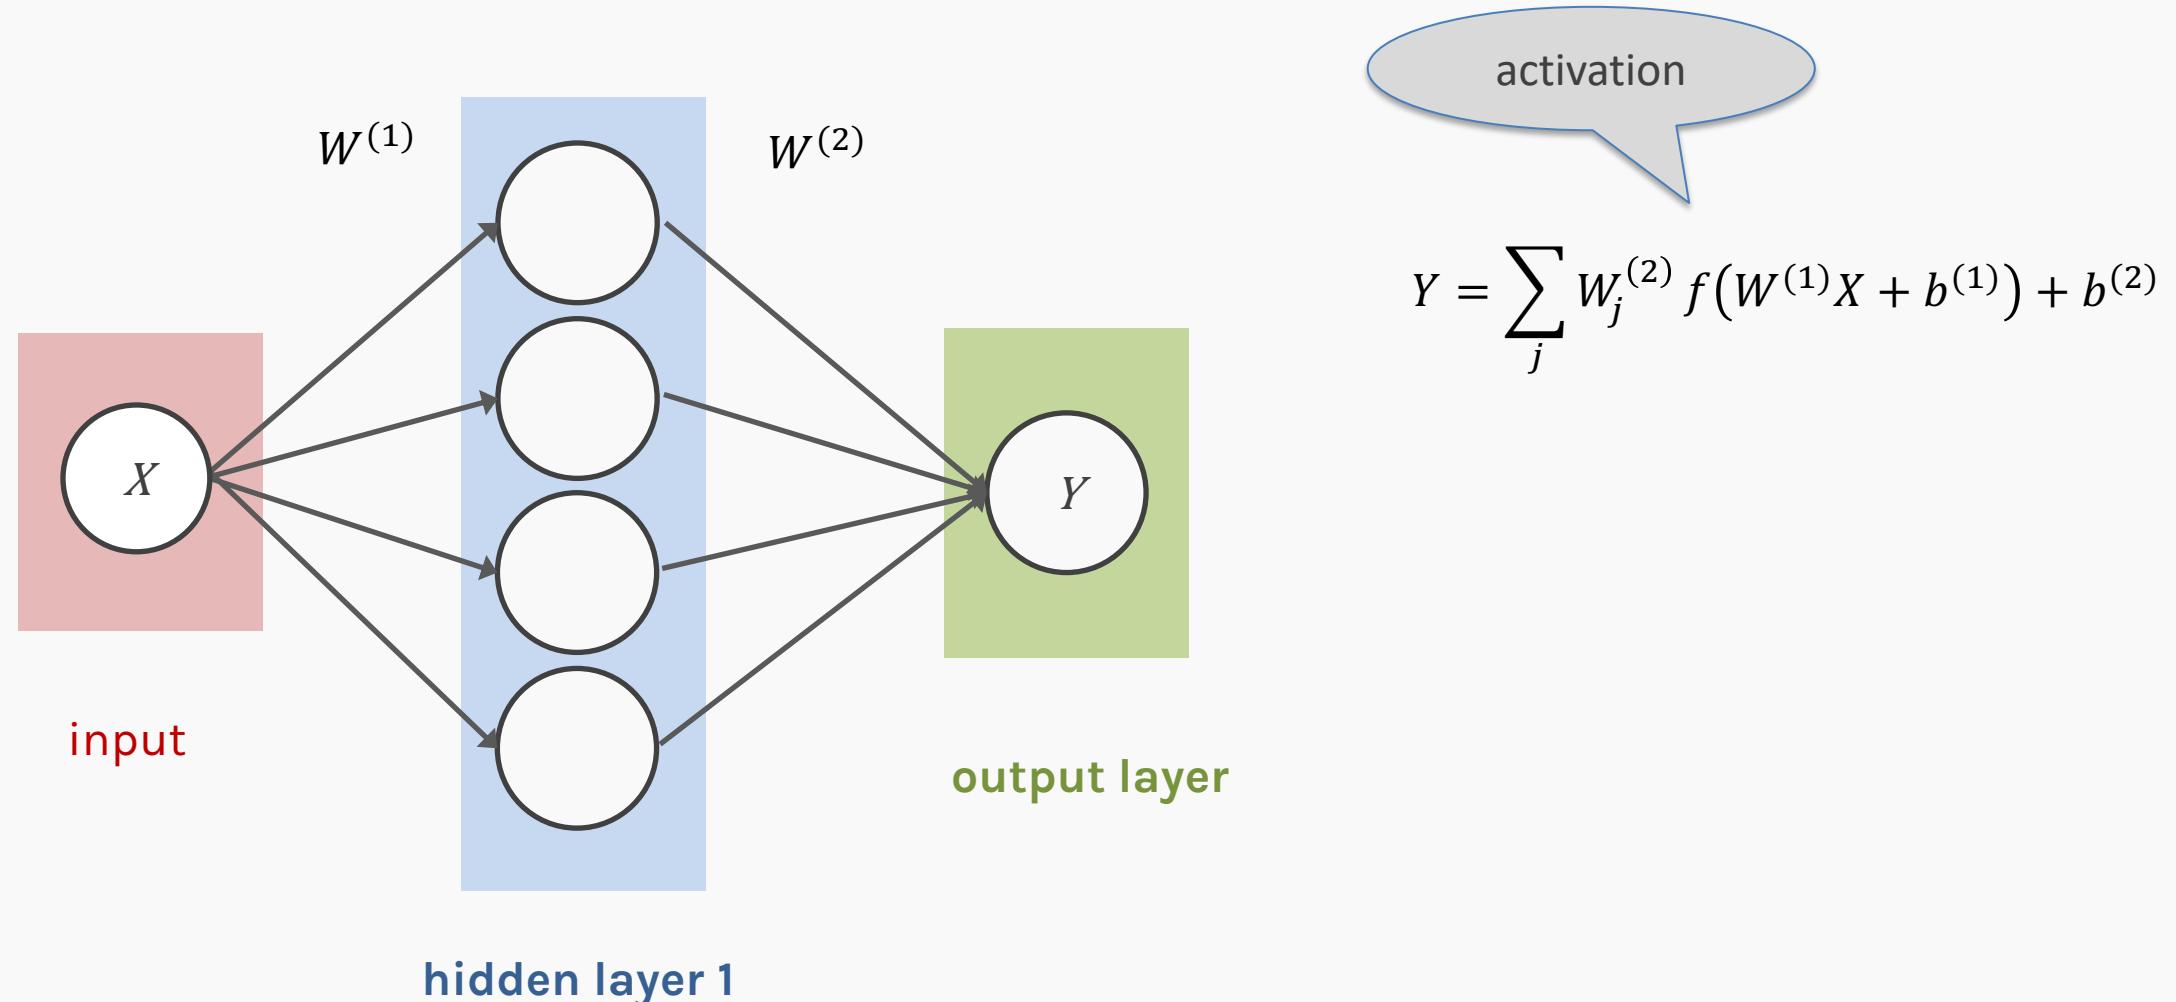
Quick review of MLPs



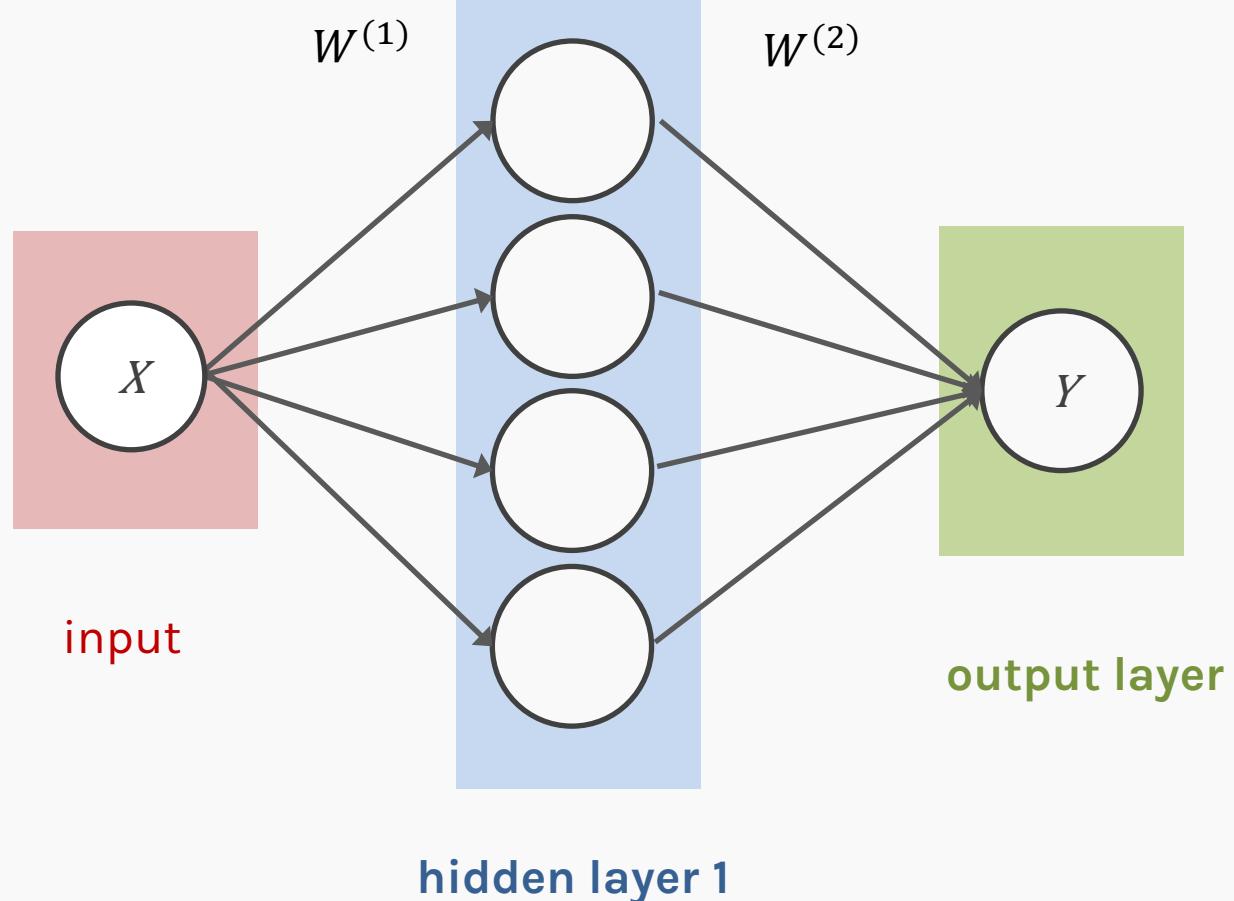
Quick review of MLPs



MLP as an additive model



MLP as an additive model

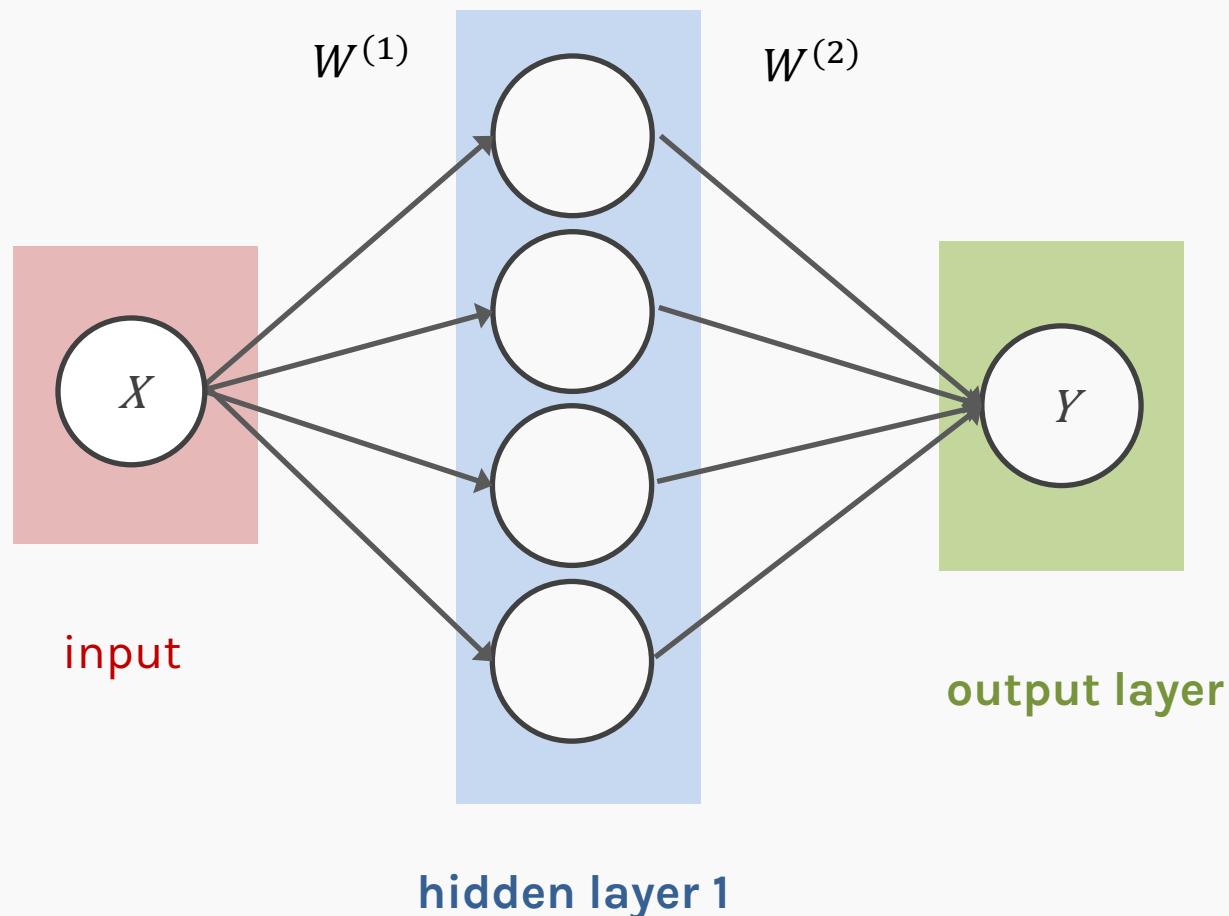


activation

$$Y = \sum_j W_j^{(2)} f(W^{(1)}X + b^{(1)}) + b^{(2)}$$

Basis functions.

MLP as an additive model



activation

$$Y = \sum_j W_j^{(2)} f(W^{(1)}X + b^{(1)}) + b^{(2)}$$

Basis functions.

Y is a linear combination of these basis functions.

We learn the coefficients of the basis functions $W_j^{(2)}$ as well as the parameters of the basis functions ($W_j^{(1)}, \beta_j$)

MLP as an additive model (cont)

From lecture 1:

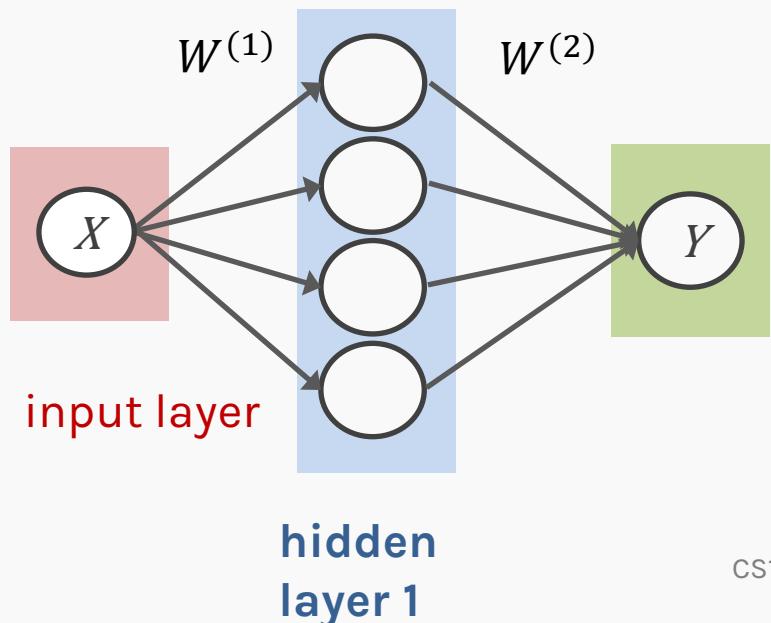
$$E(Y|x) = \alpha_0 + \alpha_1 x + \beta_1(x - \xi_1)_+ + \beta_2(x - \xi_2)_+ + \dots + \beta_k(x - \xi_k)_+$$

Minor modification:

$$E(Y|x) = \alpha_0 + \beta_0(x - \infty)_+ + \beta_1(x - \xi_1)_+ + \beta_2(x - \xi_2)_+ + \dots + \beta_k(x - \xi_k)_+$$

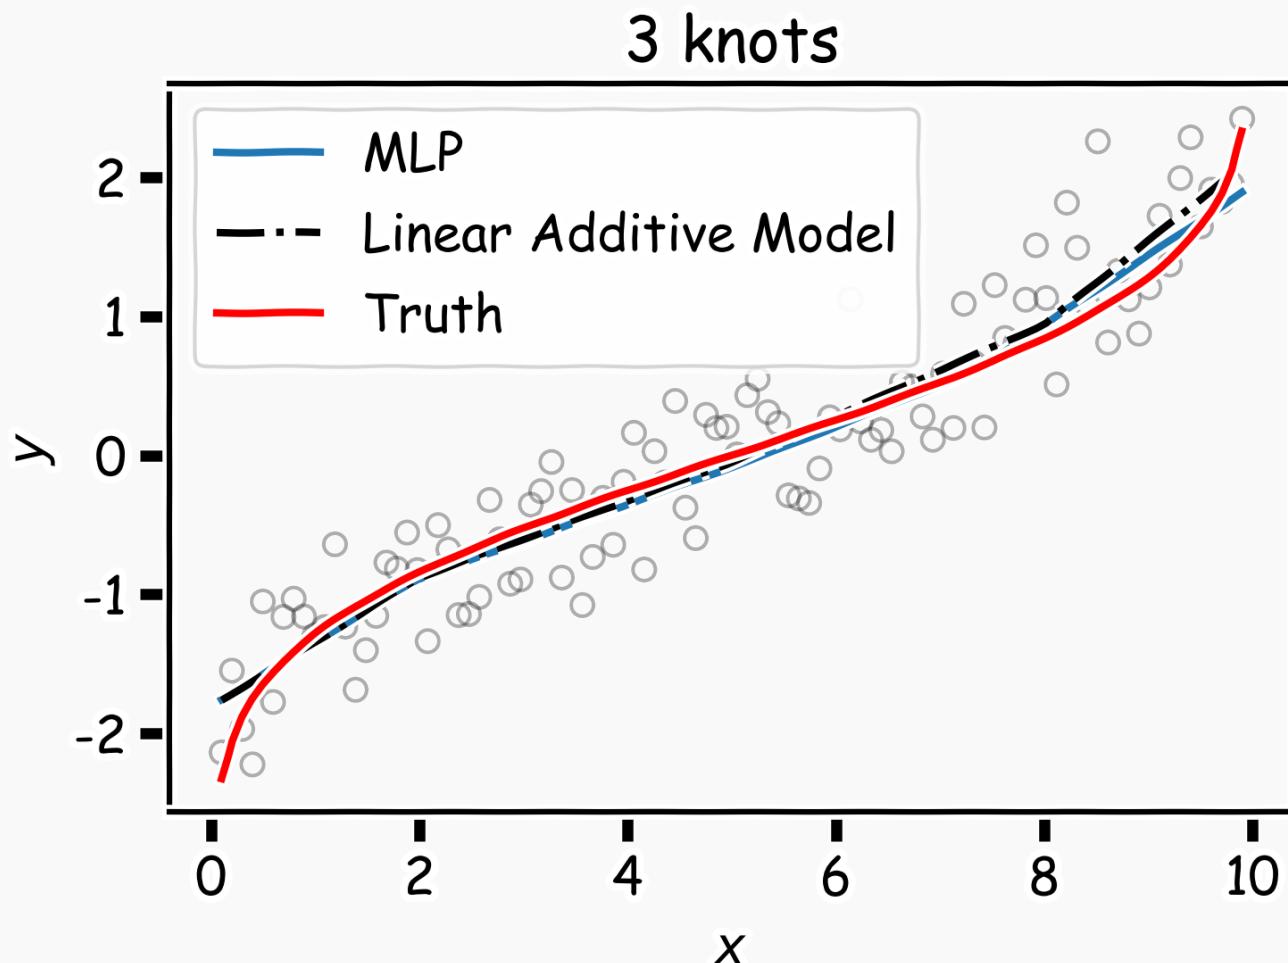


$$\text{ReLU}(Wx + \xi_1) \text{ where } W = 1$$



Location of Knots can be learned
as well as the β 's and α_0

MLP as an additive model (cont)



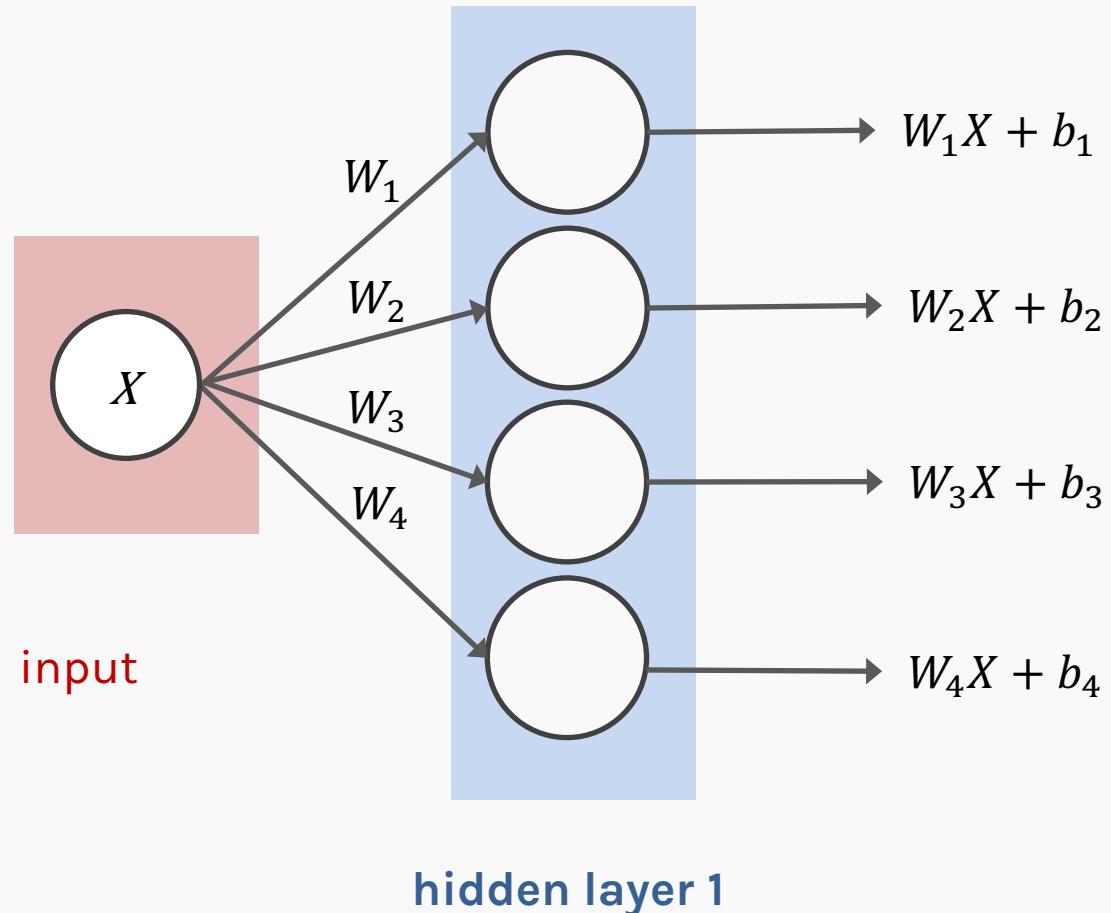
MLP:

$$\begin{aligned}\xi_1 &= 1.98248 \\ \xi_2 &= 5.03615 \\ \xi_3 &= 7.91110\end{aligned}$$

Main drawbacks of MLPs

- MLPs use one node for each input (e.g. pixel in an image, or 3 pixel values in RGB case). The number of weights rapidly becomes unmanageable for large images.
- Training difficulties arise, *overfitting* can appear.
- MLPs react differently to an input (images) and its shifted version – they are not translation invariant.

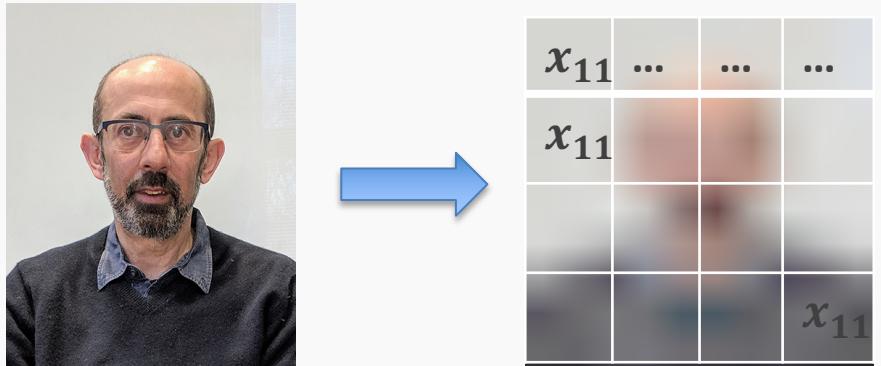
MLP: number of weights



How many weights?

- If $X \in \mathbb{R}$ then $|W| = 1$
- If $X \in \mathbb{R}^m$ then $|W| = m$

MLP: number of weights for images



If we consider each pixel as an independent predictor, then $X \in \mathbb{R}^{4 \times 4}$ or 16 predictors, there are **16 weights for each node** in the first hidden layer.

A strong motivation for performing model selection is to avoid overfitting, which can happen when:

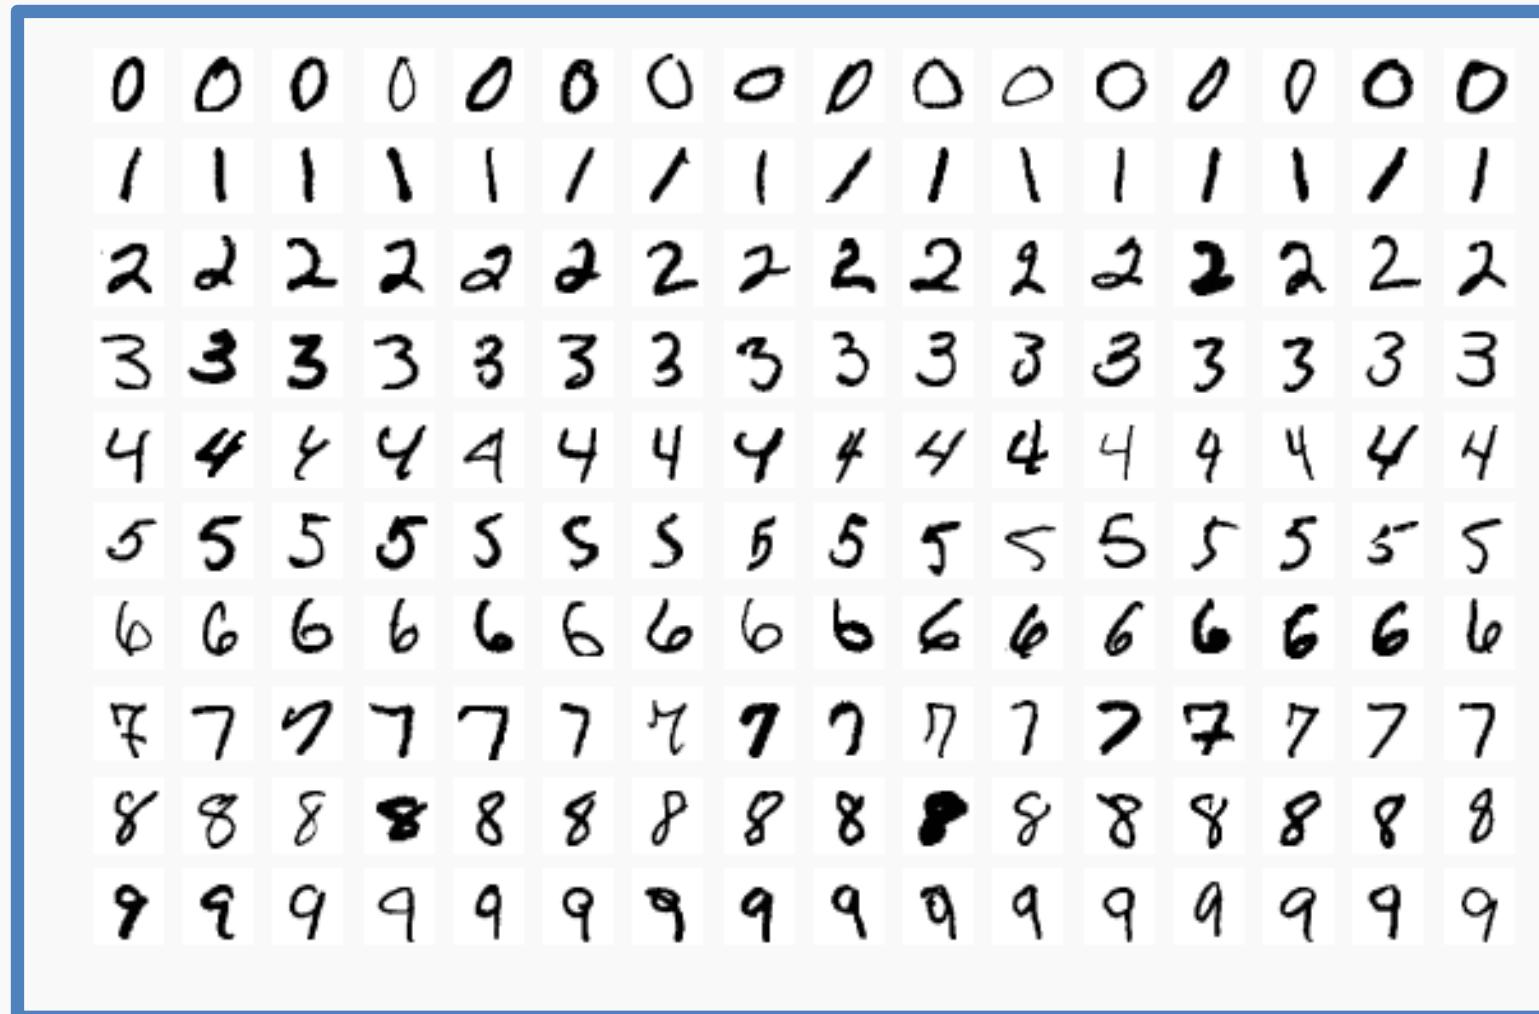
- there are too many predictors
- **the feature space is high dimensional**
- the polynomial degree is too high
- too many cross-terms are considered

Common Dataset: MNIST

MNIST database is a large set of handwritten digits.

It contains 60,000 training images and 10,000 testing images.

Every image 28x28 pixel and anti-aliased, which introduced grayscale levels



MLP: number of weights for images

Example: CIFAR10 is a dataset of images that are commonly used to train machine learning models. It contains 60,000 32x32 color images in 10 different classes.

Each pixel is a feature: an MLP would have $32 \times 32 \times 3 + 1 = 3073$ weights per neuron!

airplane



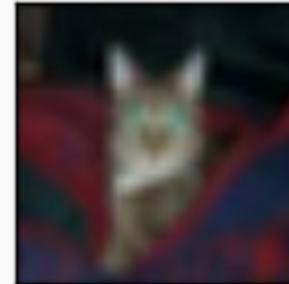
automobile



bird



cat



deer



dog



frog



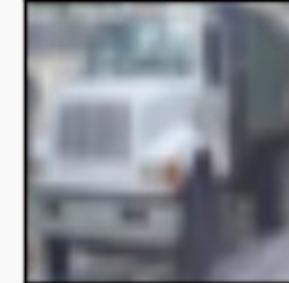
horse



ship



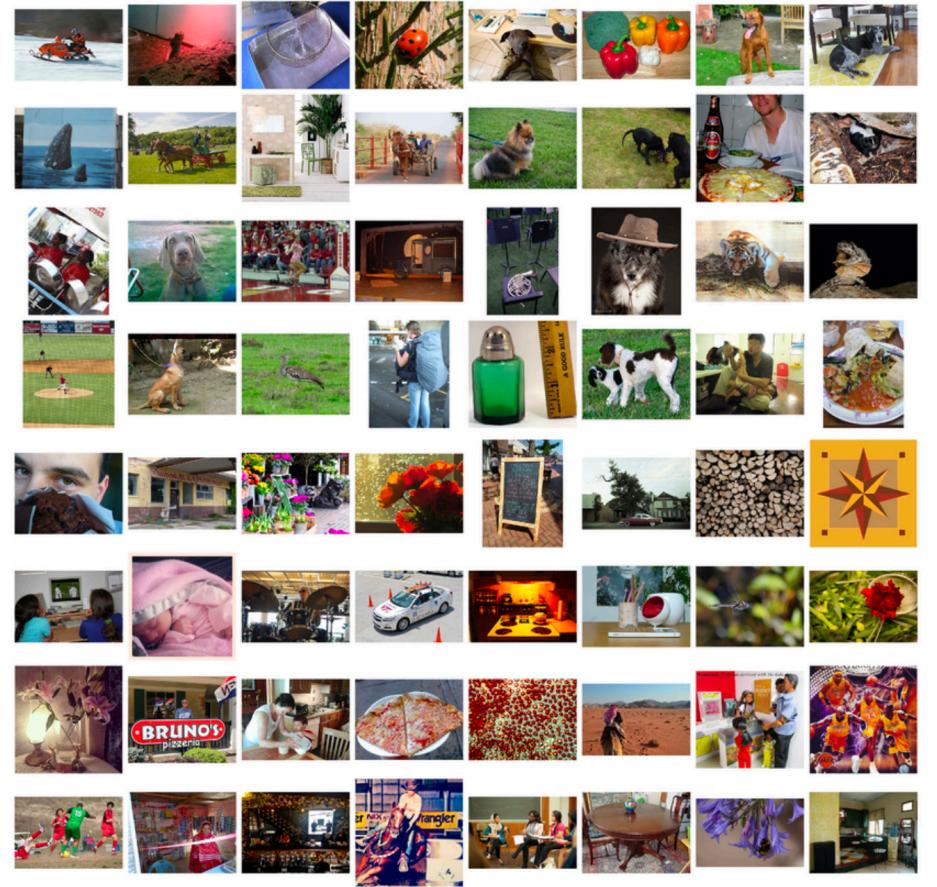
truck



MLP: number of weights for images

Example: [ImageNet](#) is a large visual database designed for use in visual object recognition software research. More than 14 million images have been hand-annotated by the project to indicate what objects are pictured. In at least one million of the images, **bounding boxes** are also provided.

Images are usually 224x224x3: an MLP would have **150129 weights per neuron**. If the first layer of the MLP is around 128 nodes, which is small, this already becomes very heavy to train.



Will my neural network
stop overfitting someday?

AS SURE AS THE
SUN WILL RISE

Model Selection and Dimensionality Reduction

Recall from CS109A that to reduce the number of predictors we can:



Model Selection and Dimensionality Reduction

Recall from CS109A that to reduce the number of predictors we can:

- PCA



Model Selection and Dimensionality Reduction

Recall from CS109A that to reduce the number of predictors we can:

- PCA
- Stepwise Variable Selection

Model Selection and Dimensionality Reduction

Recall from CS109A that to reduce the number of predictors we can:

- PCA
- Stepwise Variable Selection
- Regularization, in particular L1 will produce sparsity



Model Selection and Dimensionality Reduction

Recall from CS109A that to reduce the number of predictors we can:

- PCA
- Stepwise Variable Selection
- Regularization, in particular L1 will produce sparsity
- Drop predictors that are highly correlated

Model Selection and Dimensionality Reduction

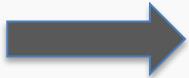
Recall from CS109A that to reduce the number of predictors we can:

- PCA
- Stepwise Variable Selection
- Regularization, in particular L1 will produce sparsity
- Drop predictors that are highly correlated
- Summarize input (image) with high level features => feature extraction or representation learning

Feature extraction



x



Features:

1. Bald
2. Grey hair
3. Oval shape head
4. Glasses

Feature extraction



x



Features:

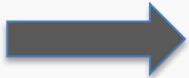
1. Bald
2. Grey hair
3. Oval shape head
4. Glasses

WAIT FOR IT

Feature extraction



x



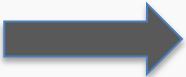
Features:

1. Bald
2. Grey hair
3. Oval shape head
4. Glasses
5. Awesome

Feature extraction

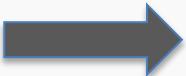


x



Features:

1. Bald
2. Grey hair
3. Oval shape head
4. Glasses
5. Awesome



Features:

1. Bald
2. Grey hair
3. Oval shape head
4. No Glasses
5. Awesome

Feature extraction

x



Features:

- 1. Bald
- 2. Grey hair
- 3. Oval shape head
- 4. Glasses
- 5. Awesome



Features:

- 1. Bald
- 2. Grey hair
- 3. Oval shape head
- 4. No Glasses
- 5. Awesome



$x' \rightarrow$

$\hat{y} : \text{PAVLOS or NOT PAVLOS}$

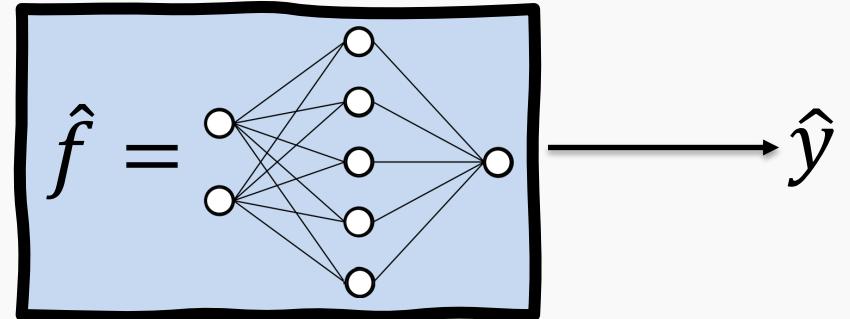


Image analysis

Imagine that we want to recognize swans in an image:



Image analysis

Imagine that we want to recognize swans in an image:

Oval-shaped white blob (body)



Image analysis

Imagine that we want to recognize swans in an image:



Image analysis

Imagine that we want to recognize swans in an image:



Cases can be a bit more complex...



Cases can be a bit more complex...

Round, elongated
head with orange
or black beak



Cases can be a bit more complex...

Round, elongated head with orange or black beak

Long white neck, square shape



Cases can be a bit more complex...

Round, elongated head with orange or black beak

Long white neck, square shape

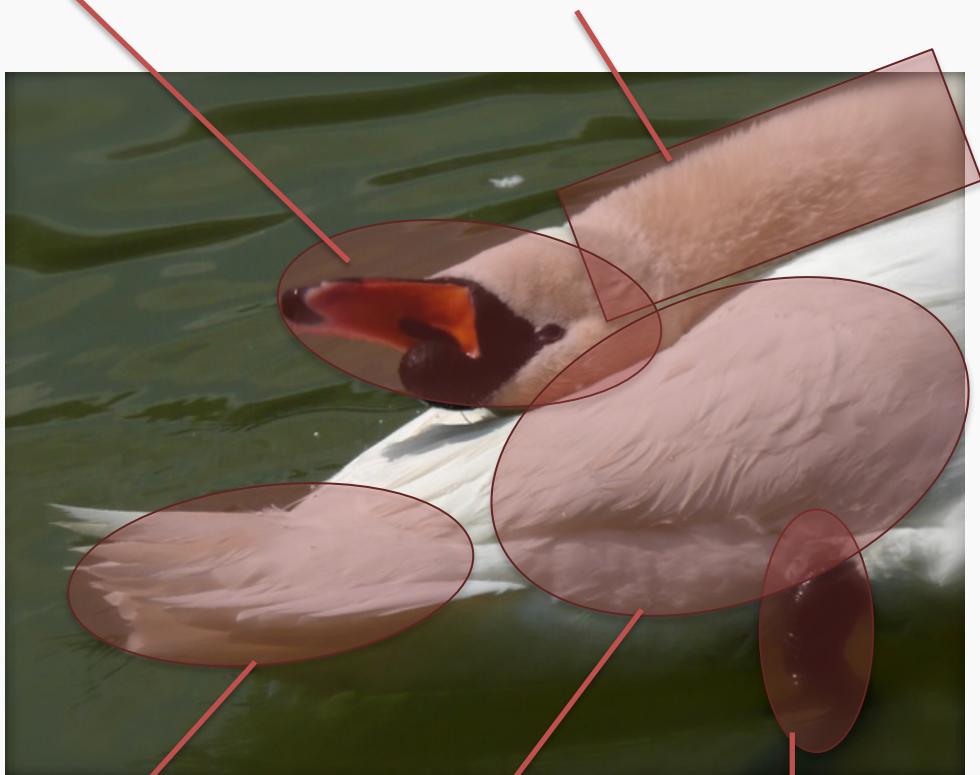


Oval-shaped white body with or without large white symmetric blobs (wings)

Now what?

Round, elongated head with orange or black beak, can be turned backwards

Long white neck, can bend around, not necessarily straight



White tail, generally far from the head, looks feathery

White, oval shaped body, with or without wings visible

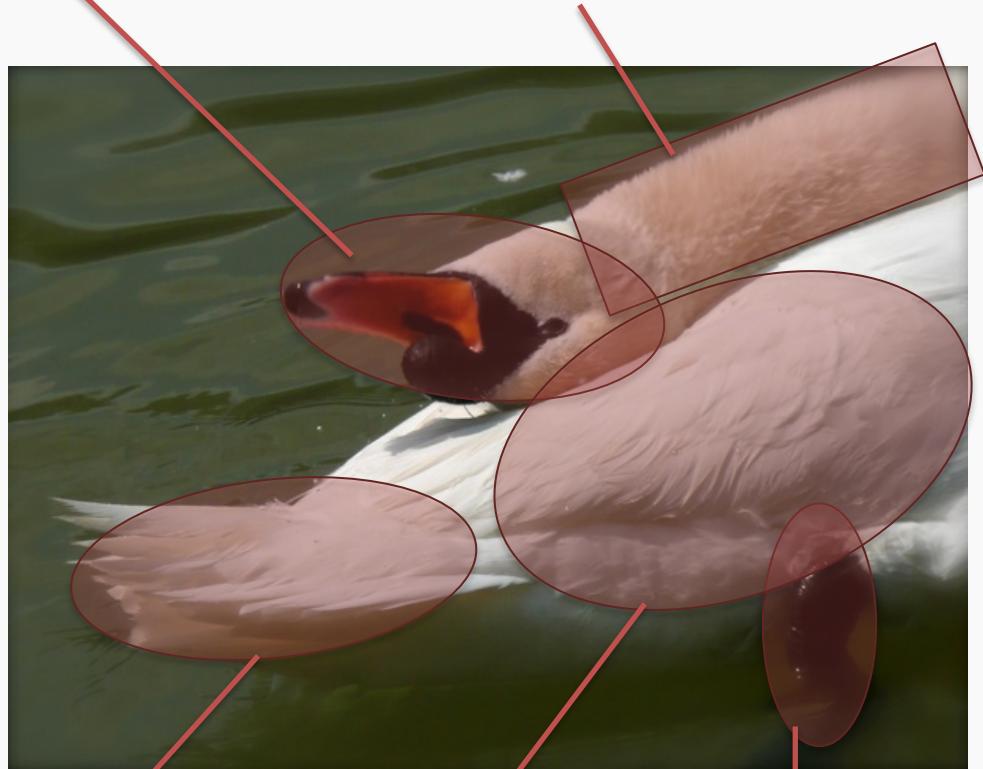
Black feet, under body, can have different shapes



Now what?

Round, elongated head with orange or black beak, can be turned backwards

Long white neck, can bend around, not necessarily straight



White tail, generally far from the head, looks feathery

White, oval shaped body, with or without wings visible

Black feet, under body, can have different shapes

Small black circles, can be facing the camera, sometimes can see both

Black triangular shaped form, on the head, can have different sizes



White elongated piece, can be squared or more triangular, can be obstructed sometimes

Luckily, the color is consistent...

We need to be able to deal with these cases



And these



And these



And these



And these



Image features

- We've been basically talking about **detecting features in images** in a very naïve way.
- Researchers built multiple computer vision techniques to deal with these issues: **SIFT, FAST, SURF, BRIEF, etc.**
- However, similar problems arose: **the detectors were either too general or too over-engineered.** Humans were **designing these feature detectors**, and that made them either too simple or hard to generalize.

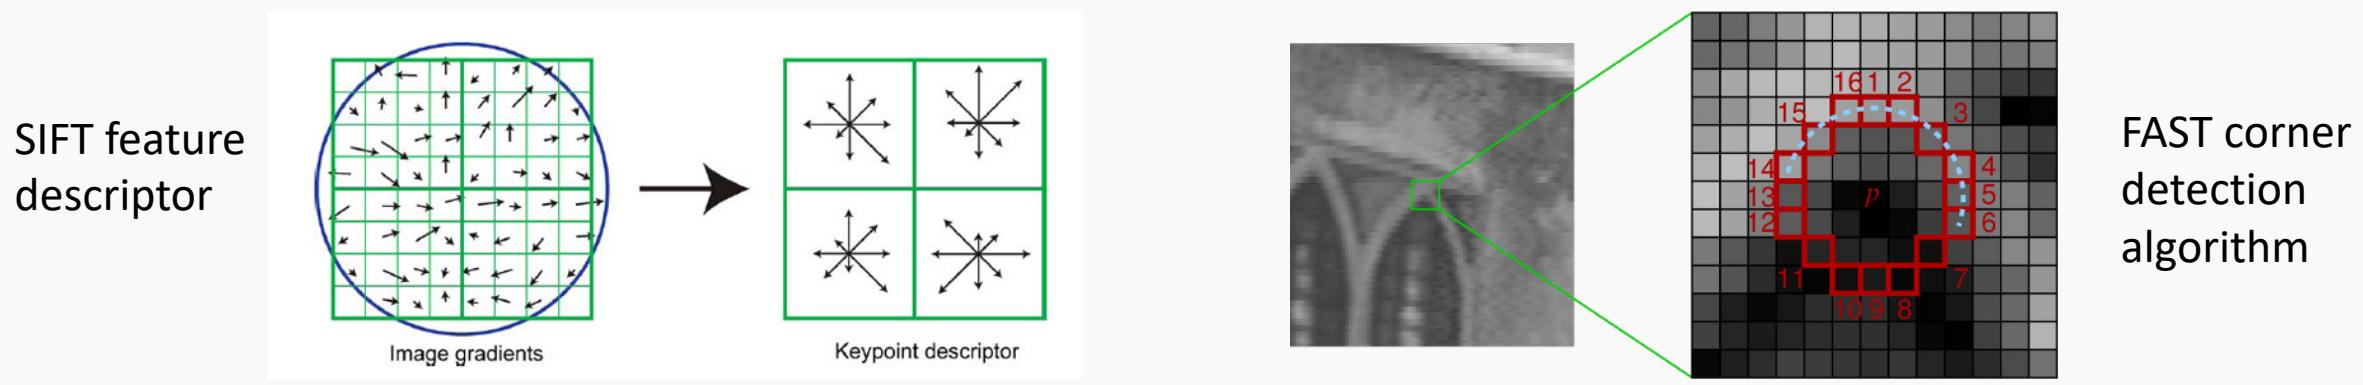


Image features (cont)

- What if we learned the features?
- We need a system that can do *Representation Learning* or *Feature Learning*.

Image features (cont)

- What if we learned the features?
- We need a system that can do *Representation Learning* or *Feature Learning*.

Representation Learning: technique that allows a system to automatically find relevant features for a given task. Replaces manual feature engineering.

Image features (cont)

- What if we learned the features?
- We need a system that can do *Representation Learning* or *Feature Learning*.

Representation Learning: technique that allows a system to automatically find relevant features for a given task. Replaces manual feature engineering.

Multiple techniques for this:

- Unsupervised (K-means, PCA, ...).
- Supervised Dictionary learning
- Neural Networks!

Some things to consider



- Nearby Pixels are more strongly related than distant ones
- Objects are built up out of smaller parts
- Images are Local and Hierarchical

Images are Invariant

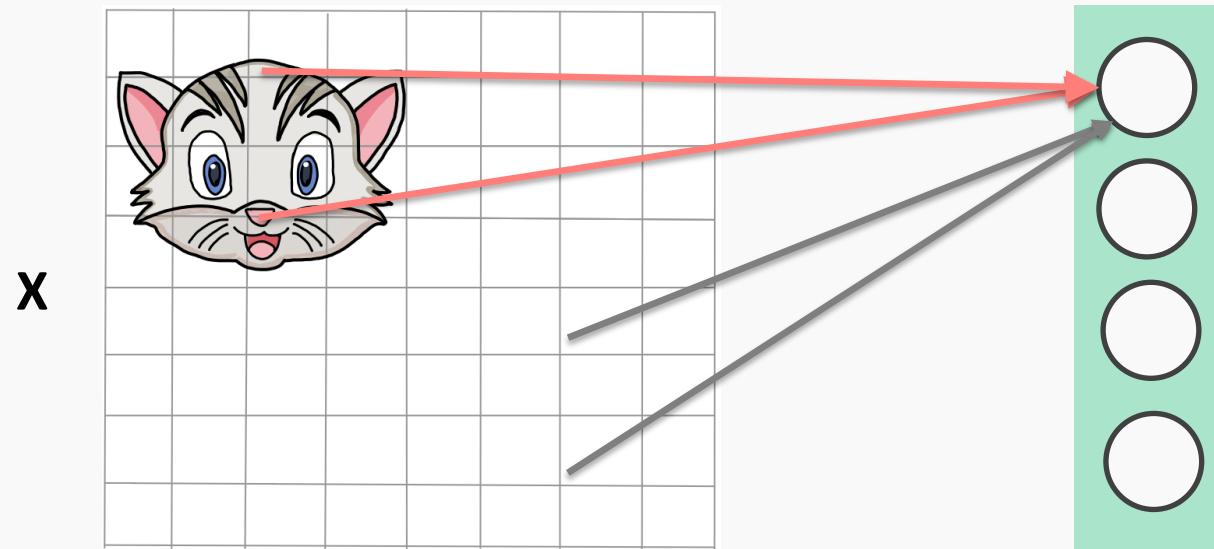


Outline

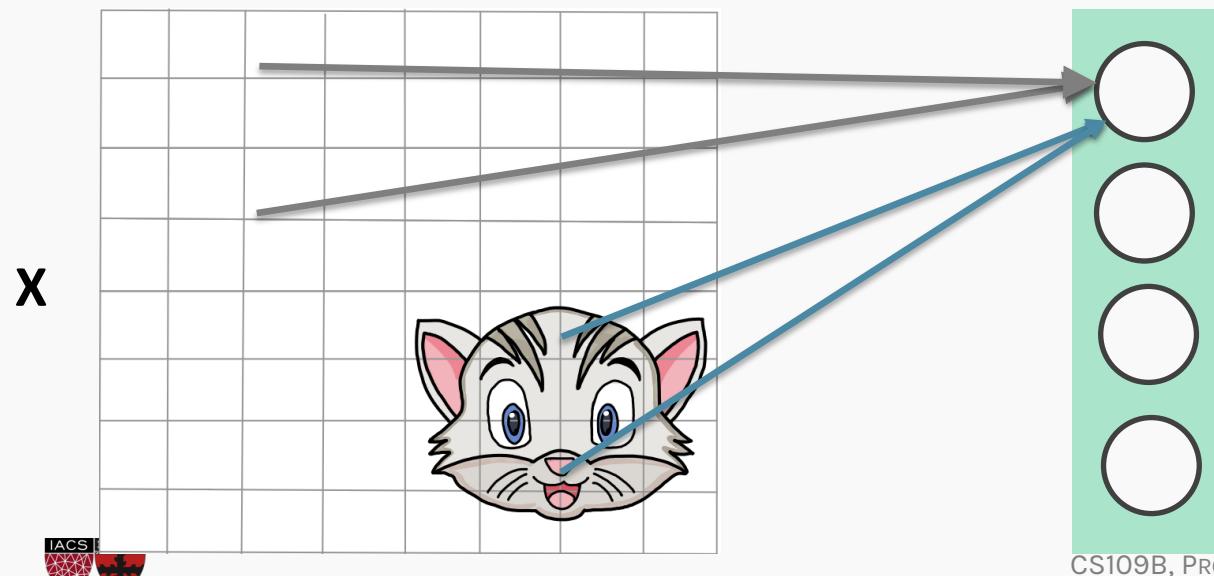
1. Motivation
2. CNN basic ideas
3. Building a CNN



Each neuron from the first layer has one weight per pixel. Recall that the importance of the predictors (here pixels) is given by the value of the coefficient (there the weight w)

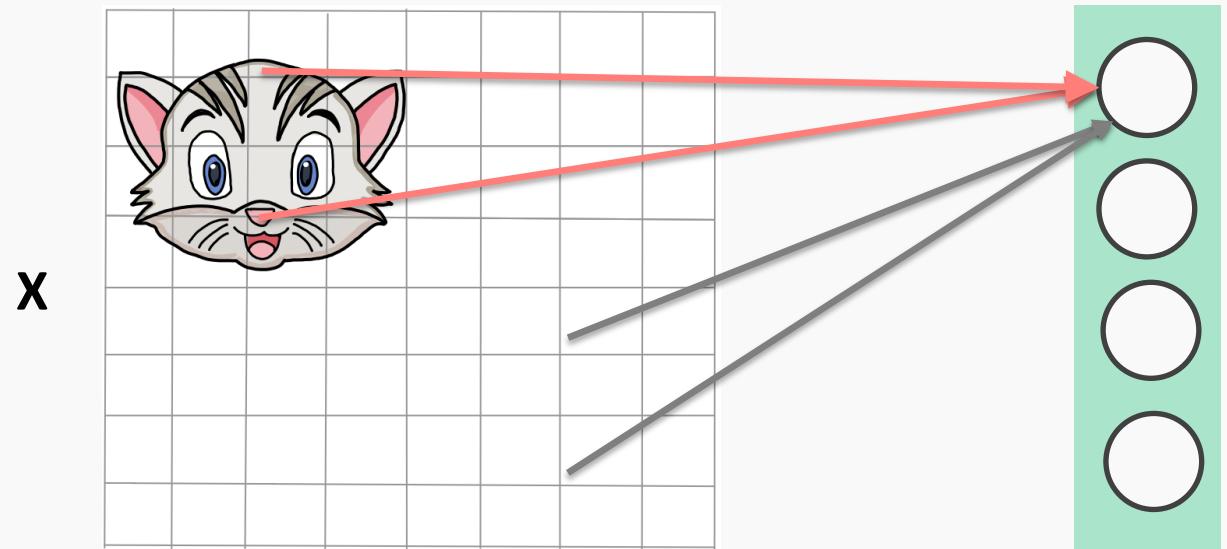


In this case, the **red weights** will be larger to better recognize "cat".

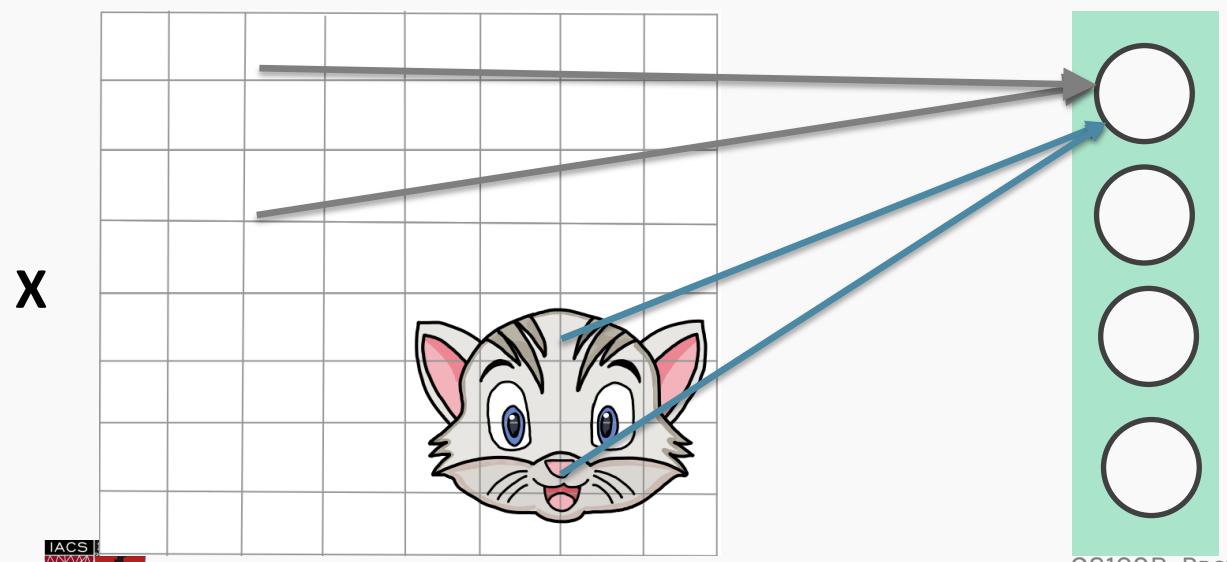


In this case, the **blue weights** will be larger.

Each neuron from the first layer has one weight per pixel. Recall that the importance of the predictors (here pixels) is given by the value of the coefficient (there the weight w)

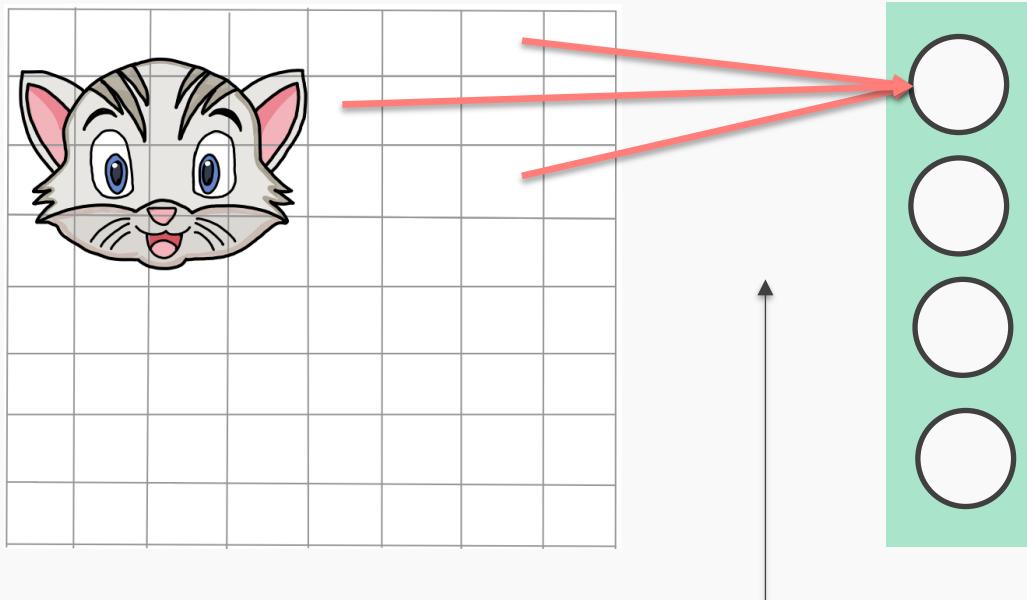


We are learning **redundant** features.
Approach is not robust, as cats could appear in yet another position.



Solution: Cut the image into smaller pieces.

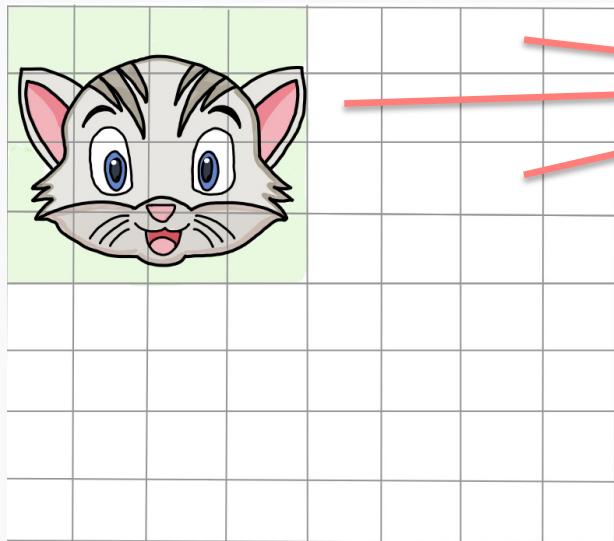
$X: \mathbb{R}^{8 \times 8}$



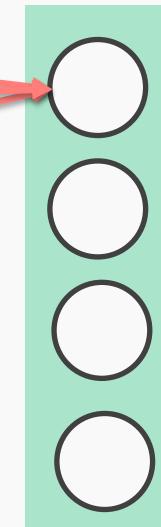
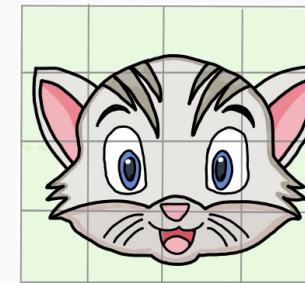
64 weights per neuron

Solution: Cut the image to smaller pieces.

$X: \mathbb{R}^{8 \times 8}$



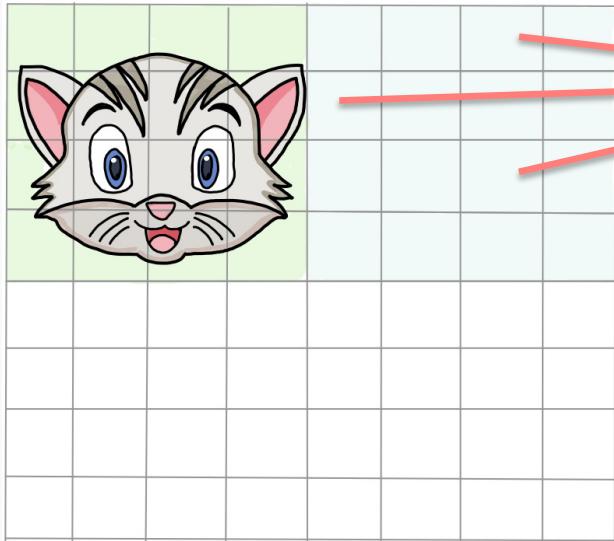
$X: \mathbb{R}^{4 \times 4}$



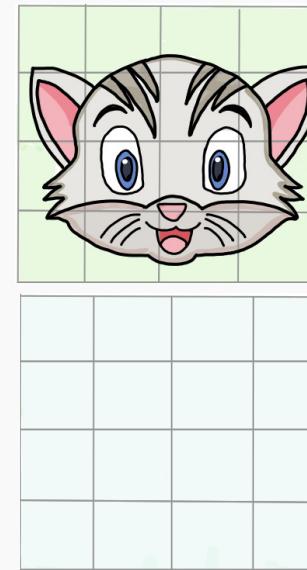
64 weights per neuron

Solution: Cut the image to smaller pieces.

$$X: \mathbb{R}^{8 \times 8}$$



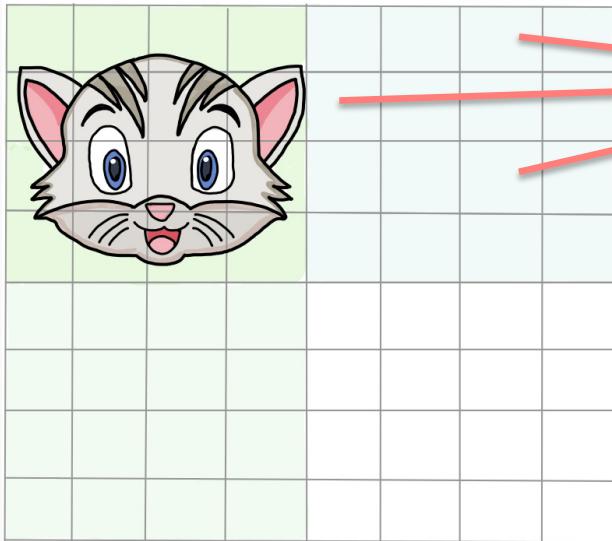
$$X: \mathbb{R}^{4 \times 4}$$



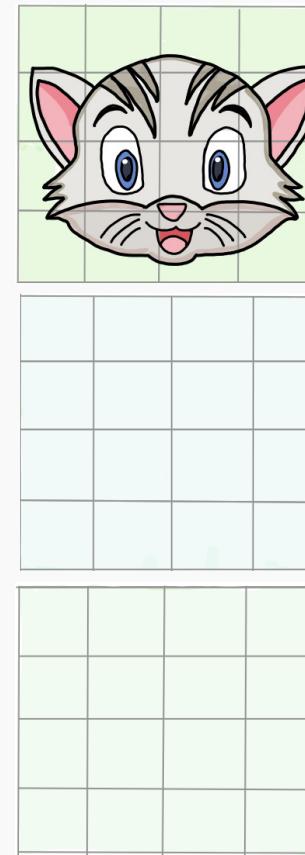
64 weights per neuron

Solution: Cut the image to smaller pieces.

$$X: \mathbb{R}^{8 \times 8}$$



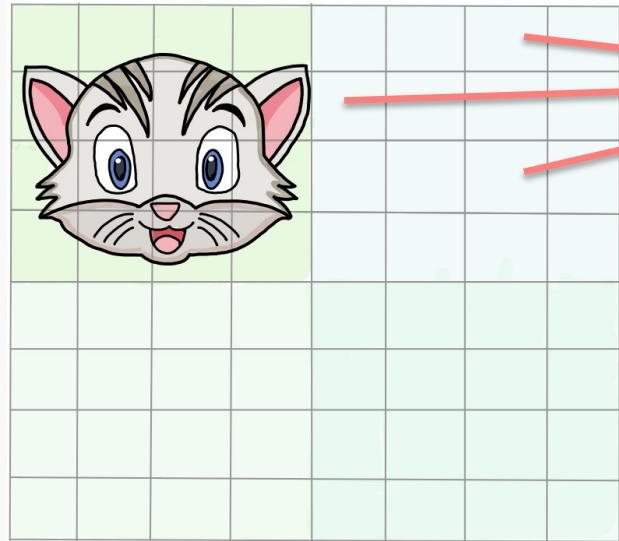
$$X: \mathbb{R}^{4 \times 4}$$



64 weights per neuron

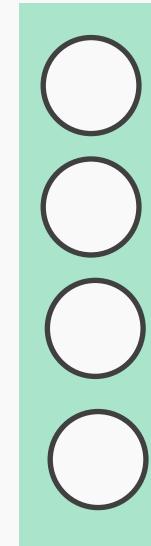
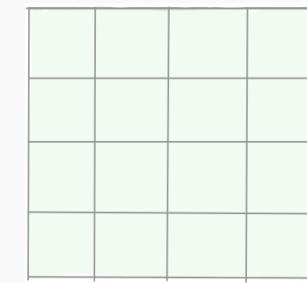
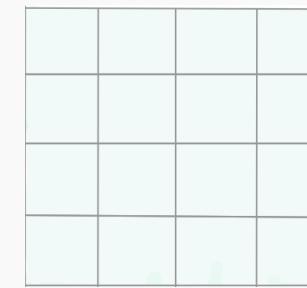
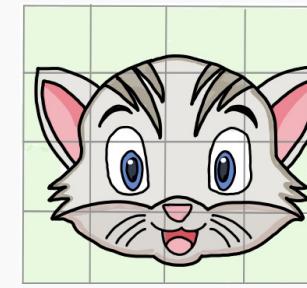
Solution: Cut the image to smaller pieces.

$X: \mathbb{R}^{8 \times 8}$



64 weights per neuron

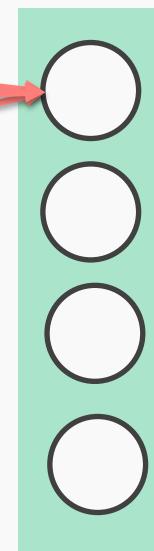
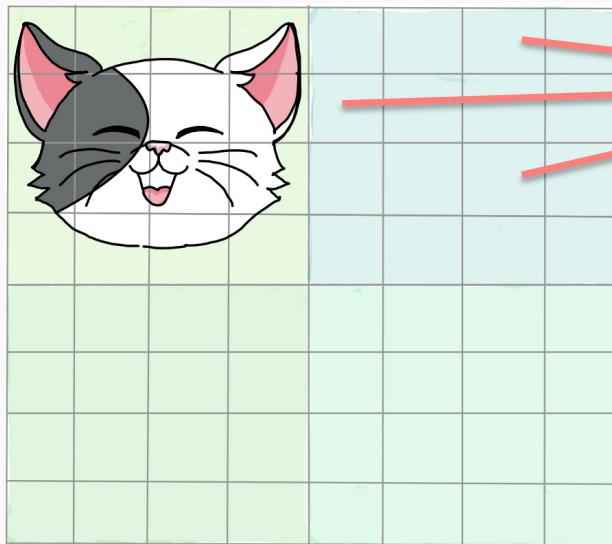
$X: \mathbb{R}^{4 \times 4}$



16 weights per neuron **but**
4 times more training
examples.

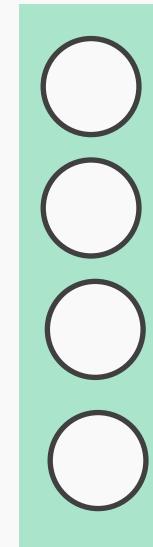
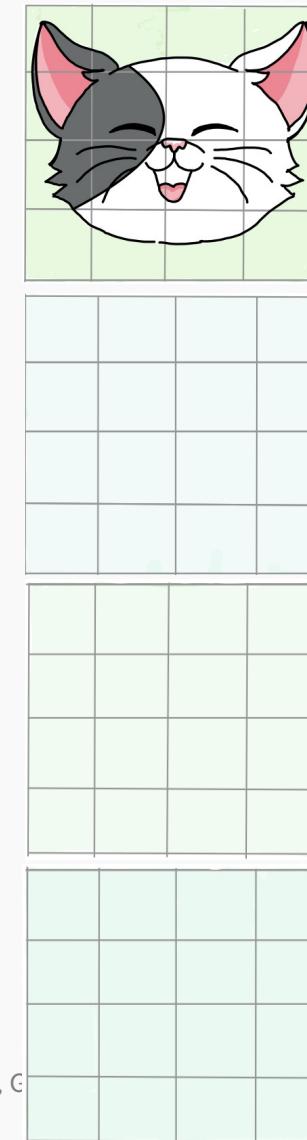
Do the same for all images

$X: \mathbb{R}^{8 \times 8}$



64 weights per neuron

$X: \mathbb{R}^{4 \times 4}$

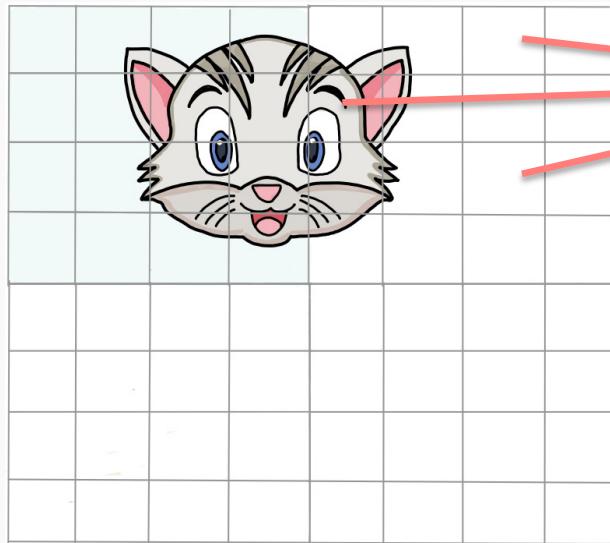


16 weights per neuron but
4 times more training
examples.

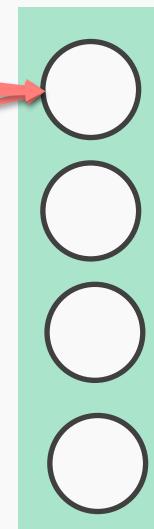
What if the cat is not entirely in one of the 4 boxes?

What if the cat is not entirely in one of the 4 boxes?

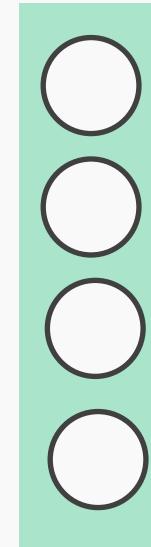
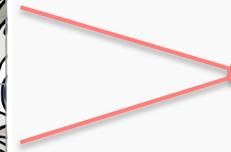
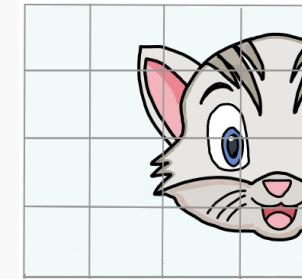
$X: \mathbb{R}^{8 \times 8}$



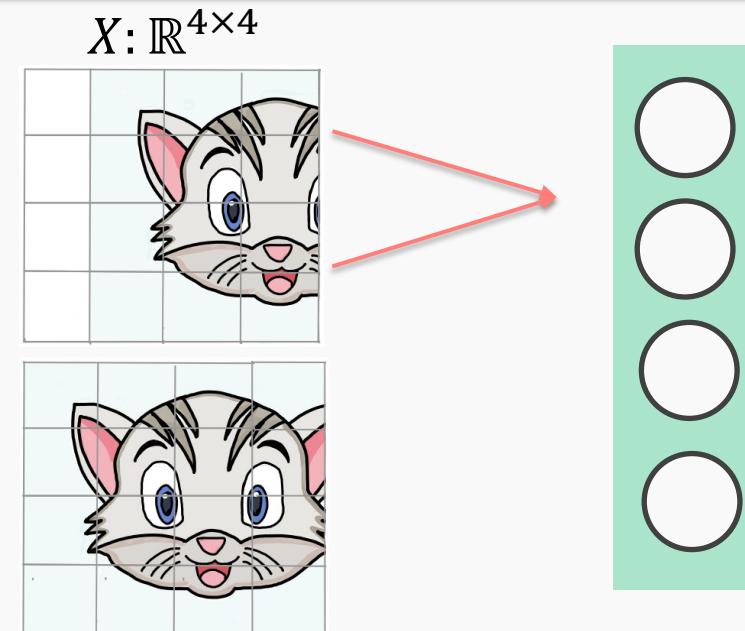
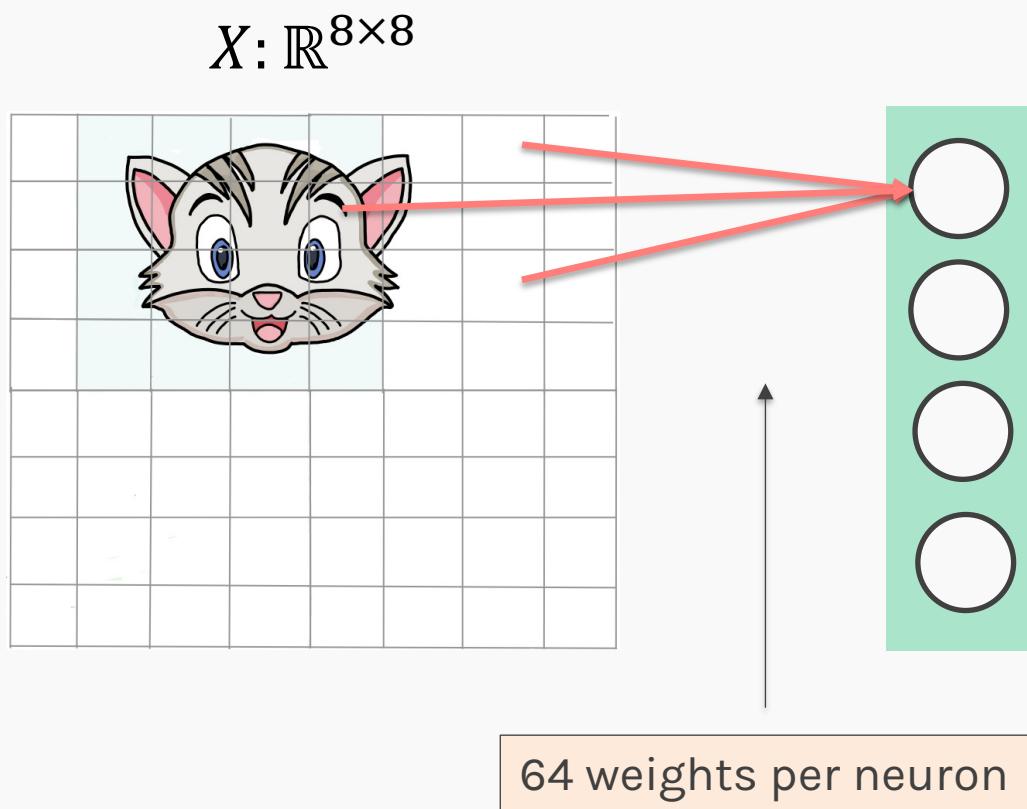
64 weights per neuron



$X: \mathbb{R}^{4 \times 4}$

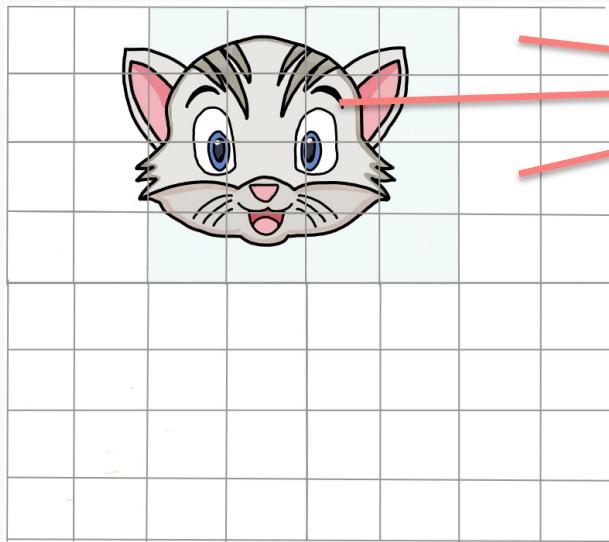


What if the cat is not entirely in one of the 4 boxes?



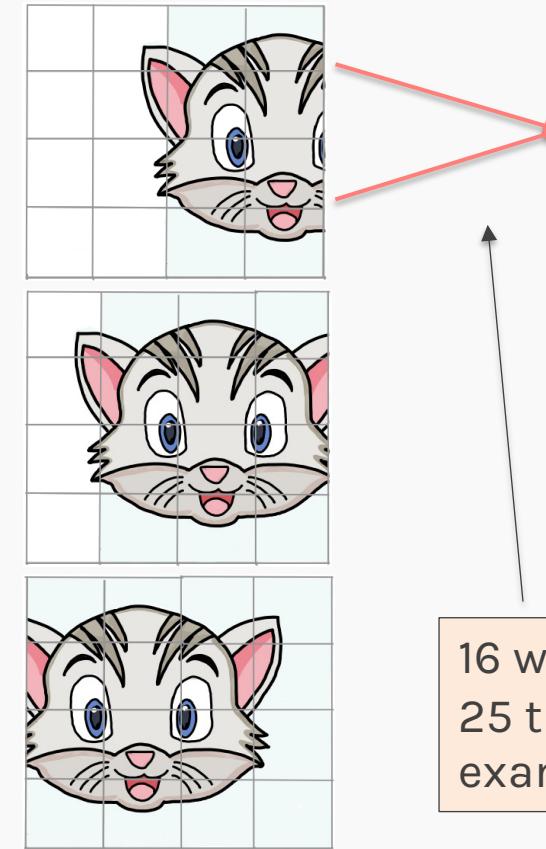
What if the cat is not entirely in one of the 4 boxes?

$X: \mathbb{R}^{8 \times 8}$



64 weights per neuron

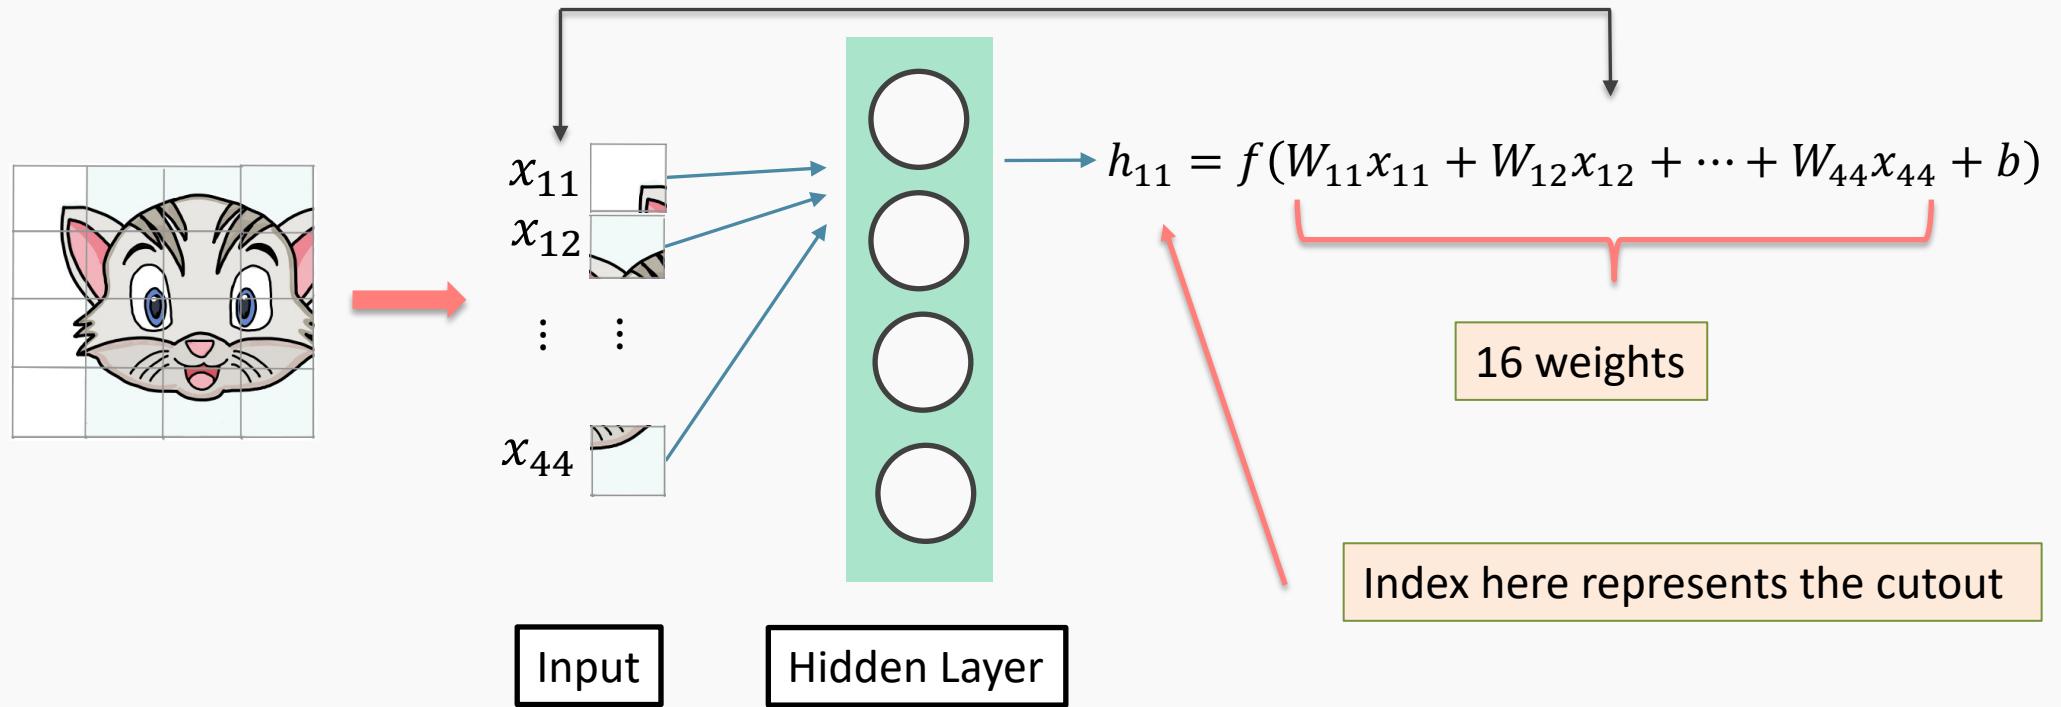
$X: \mathbb{R}^{4 \times 4}$



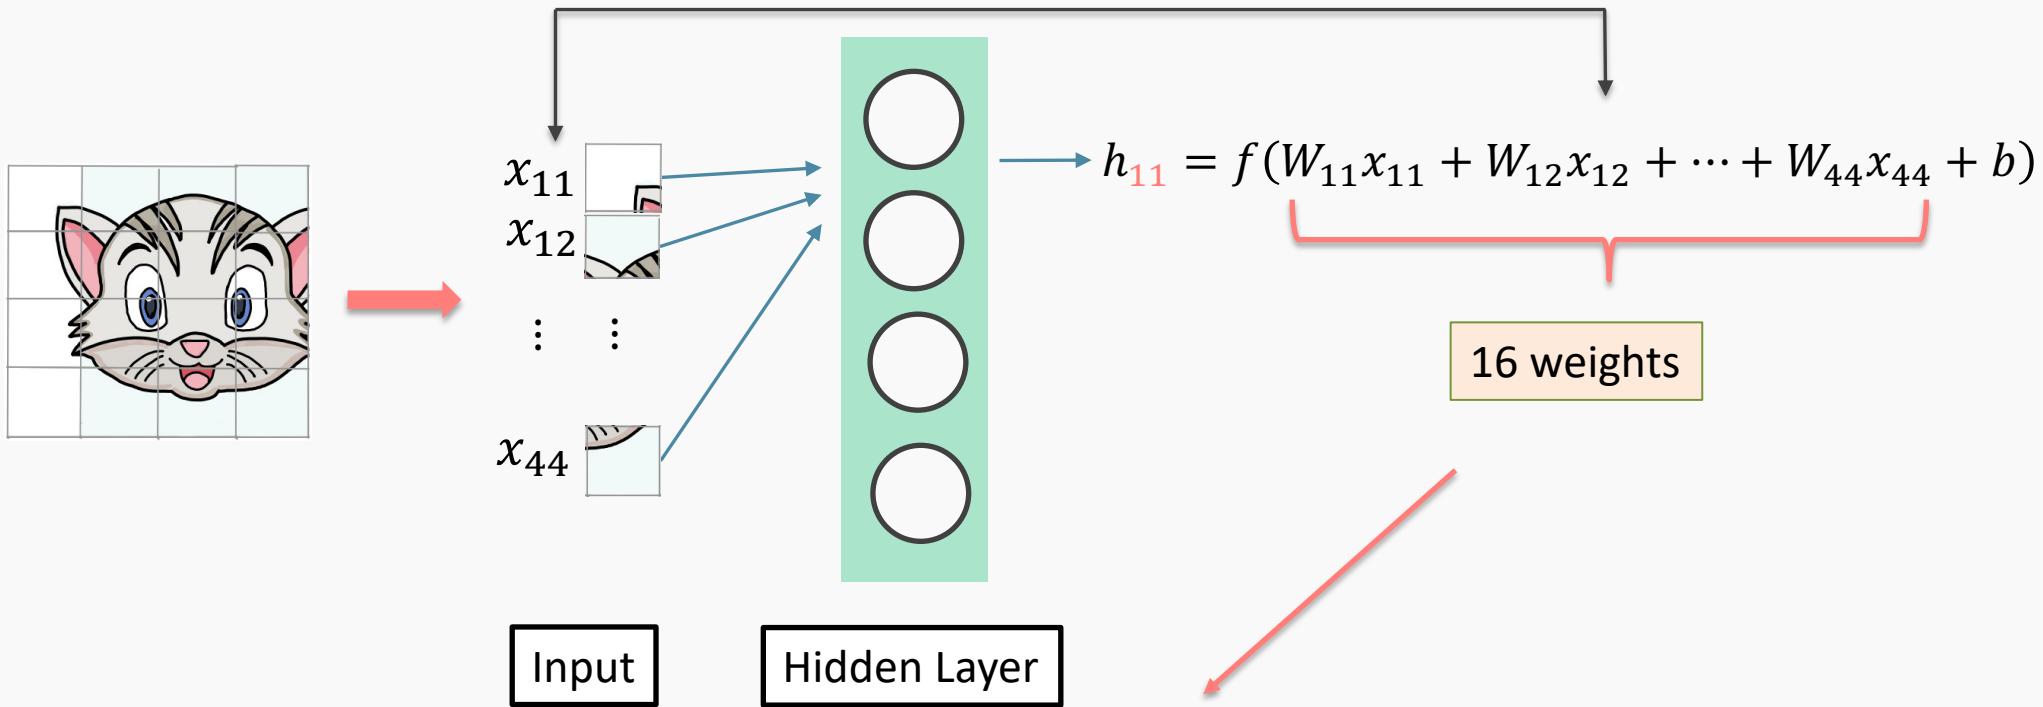
16 weights per neuron but
25 times more training
examples.

Sliding Window

Convolution



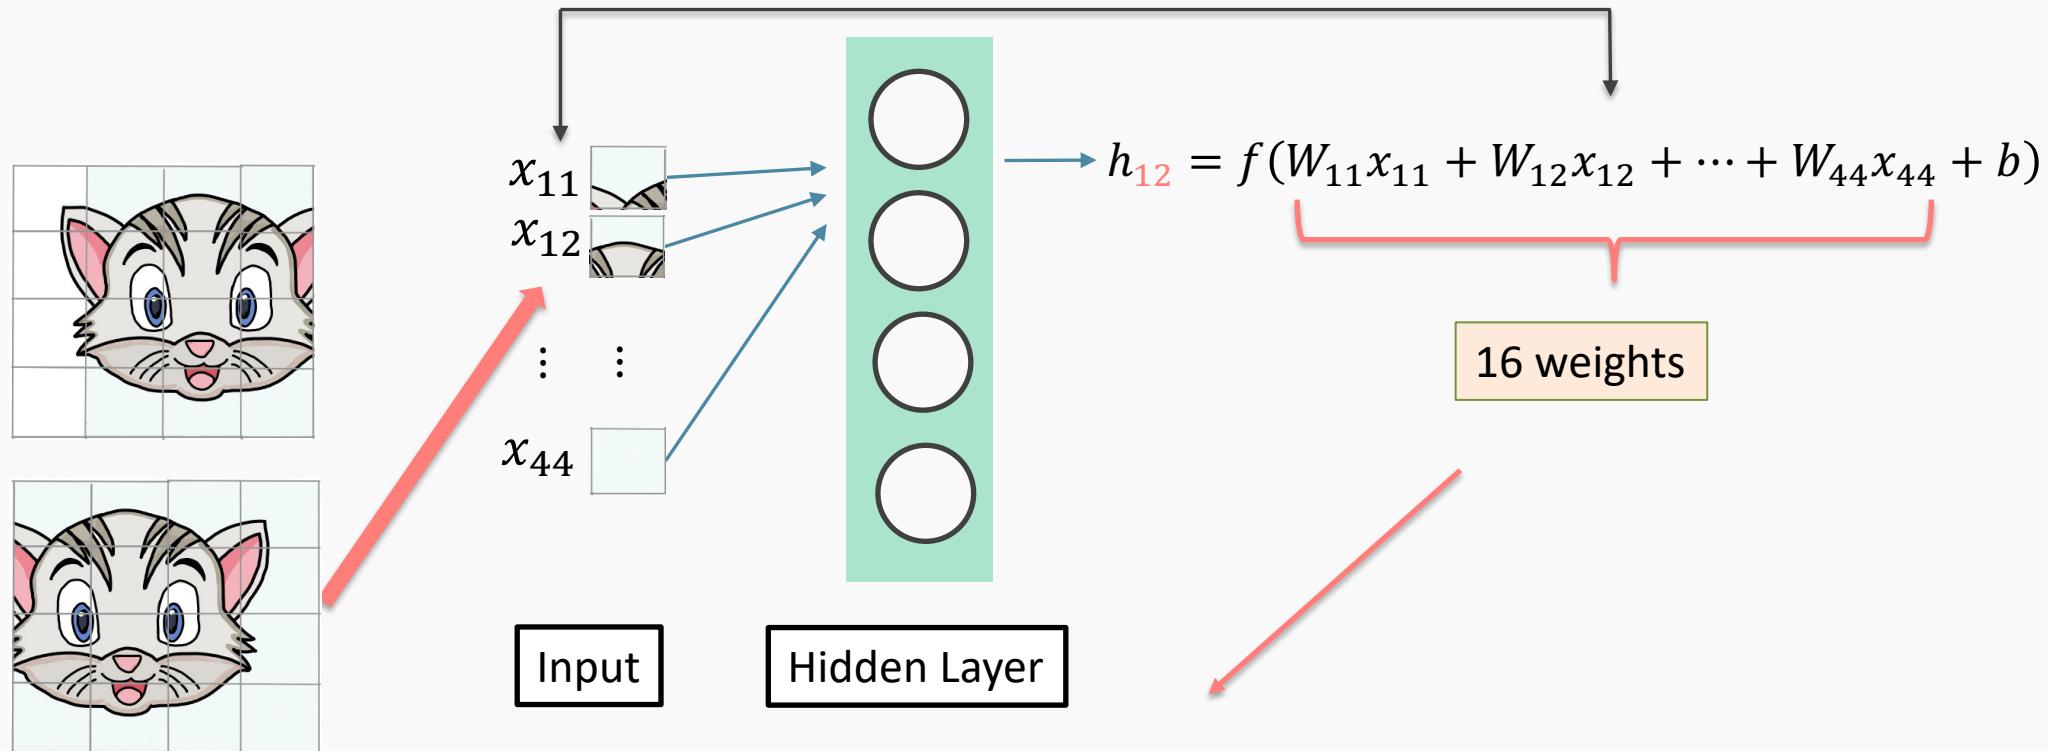
Convolution



$$h_{11} = \sum \begin{array}{|c|c|c|c|} \hline W_{11} & W_{12} & W_{13} & W_{14} \\ \hline & & & W_{44} \\ \hline \end{array} \star \begin{array}{|c|c|c|c|} \hline x_{11} & x_{12} & x_{13} & x_{14} \\ \hline & & & x_{44} \\ \hline \end{array}$$

Element wise multiplication and addition of all products

Convolution



$$h_{12} = \sum \begin{array}{|c|c|c|c|} \hline W_{11} & W_{12} & W_{13} & W_{14} \\ \hline & & & W_{44} \\ \hline \end{array} \star \begin{array}{|c|c|c|c|} \hline x_{11} & x_{12} & x_{13} & x_{14} \\ \hline & & & x_{44} \\ \hline \end{array}$$

Element wise multiplication and addition of all products

Convolution

$$h_{ij} = \sum$$

W_{11}	W_{12}	W_{13}	W_{14}
			W_{44}

★

x_{11}	x_{12}	x_{13}	x_{14}
			x_{44}

Convolution

$$h_{ij} = \sum$$

W_{11}	W_{12}	W_{13}	W_{14}
			W_{44}
			★

x_{11}	x_{12}	x_{13}	x_{14}
			x_{44}
			★

KERNEL, K

X is the cutout of image center at $\{i, j\}$

Convolution

$$h_{ij} = \sum$$



W_{11}	W_{12}	W_{13}	W_{14}
			W_{44}

★

x_{11}	x_{12}	x_{13}	x_{14}
			x_{44}

Index here represents the output
from this operation

KERNEL, K

X is the cutout of image center at $\{i, j\}$

Element wise multiplication and addition of all products = CONVOLUTION

$$H = K * X$$

Convolution and cross-correlation

- A **convolution** of f and g , $(f * g)$, is defined as the integral of the product, having one of the functions inverted and shifted:

$$(f * g)(t) = \int_a f(a)g(t - a)da$$

Function is
inverted and
shifted left by t

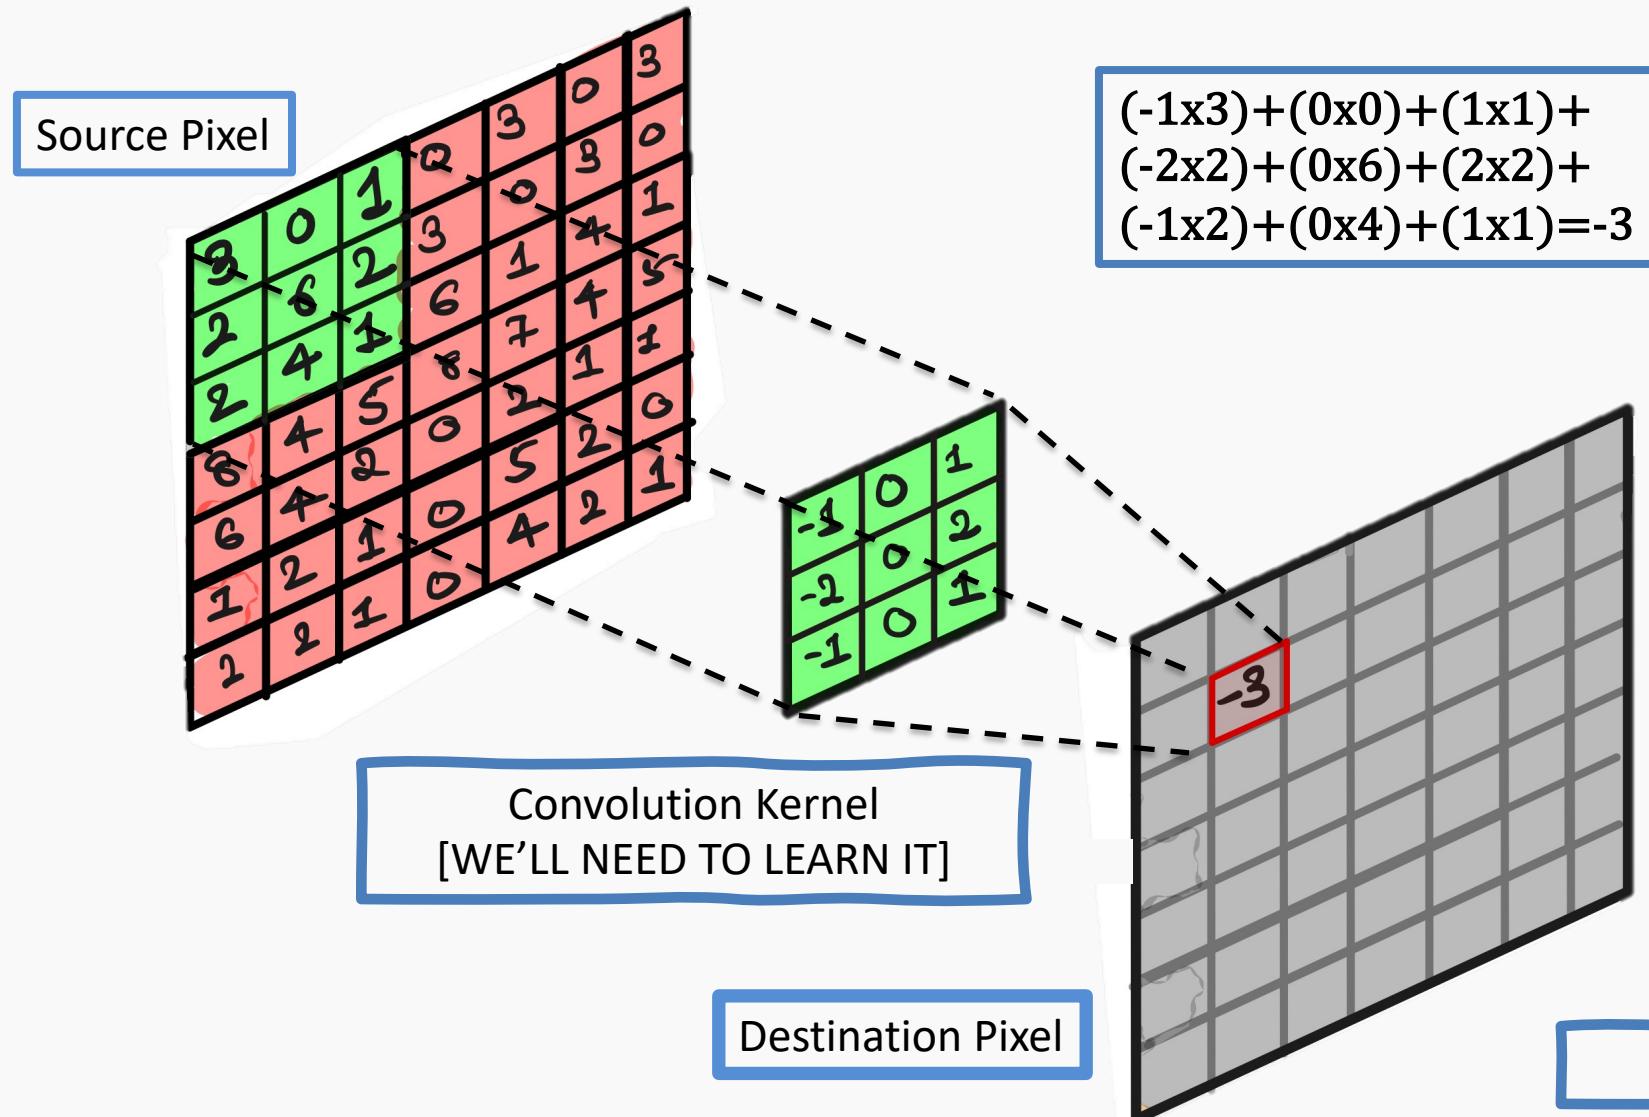
- Discrete convolution:

$$(f * g)(t) = \sum_{a=-\infty}^{\infty} f(a)g(t - a)$$

- Discrete cross-correlation:

$$(f \star g)(t) = \sum_{a=-\infty}^{\infty} f(a)g(t + a)$$

“Convolution” Operation



“Convolution” Operation in action

What does convolving an image with a Kernel do?



“Convolution” Operation in action

What does convolving an image with a Kernel do?



Edge detection

$$* \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$

Kernel

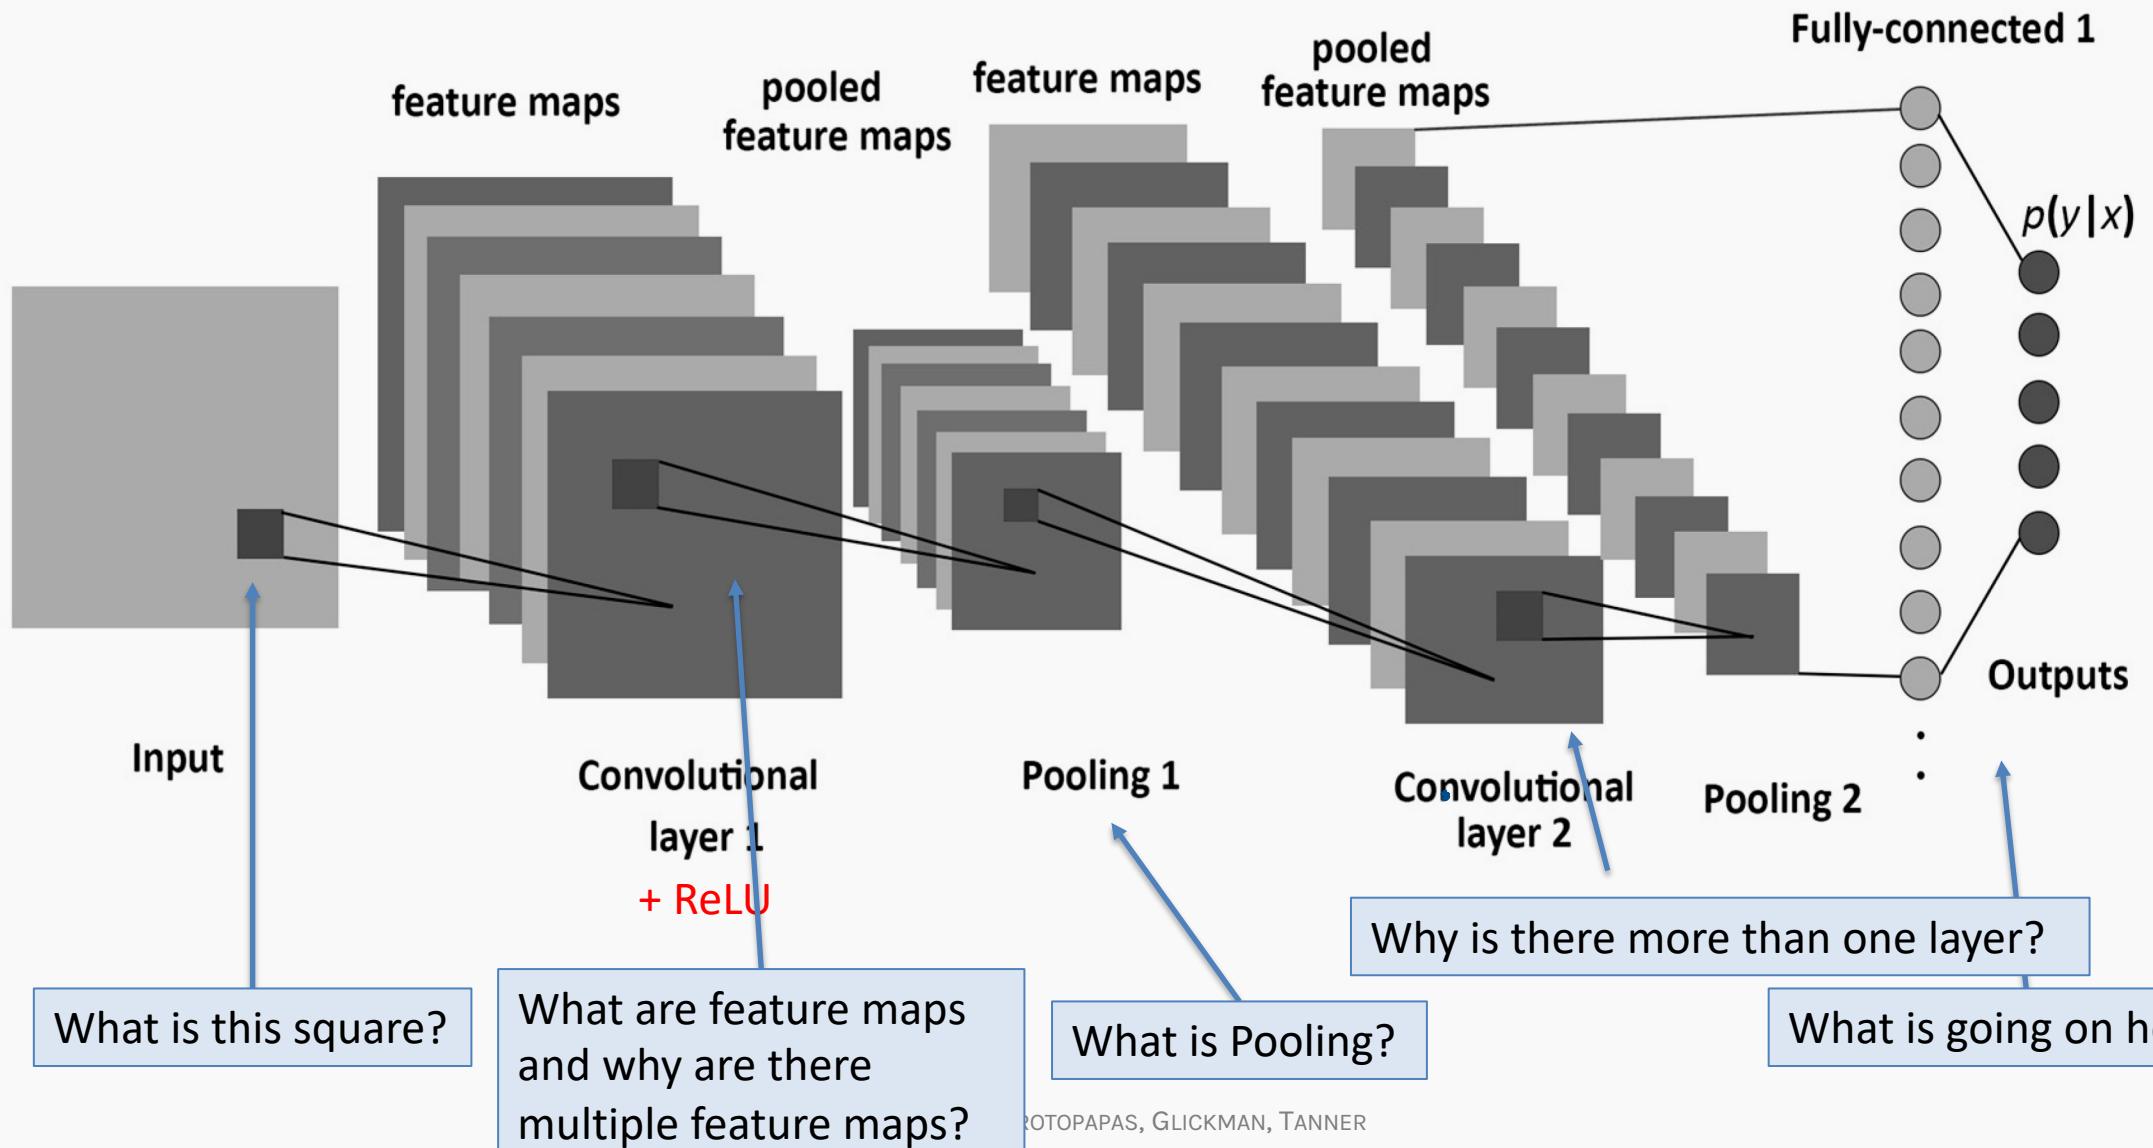
A binary image showing the edges of the man's face highlighted in white against a black background.



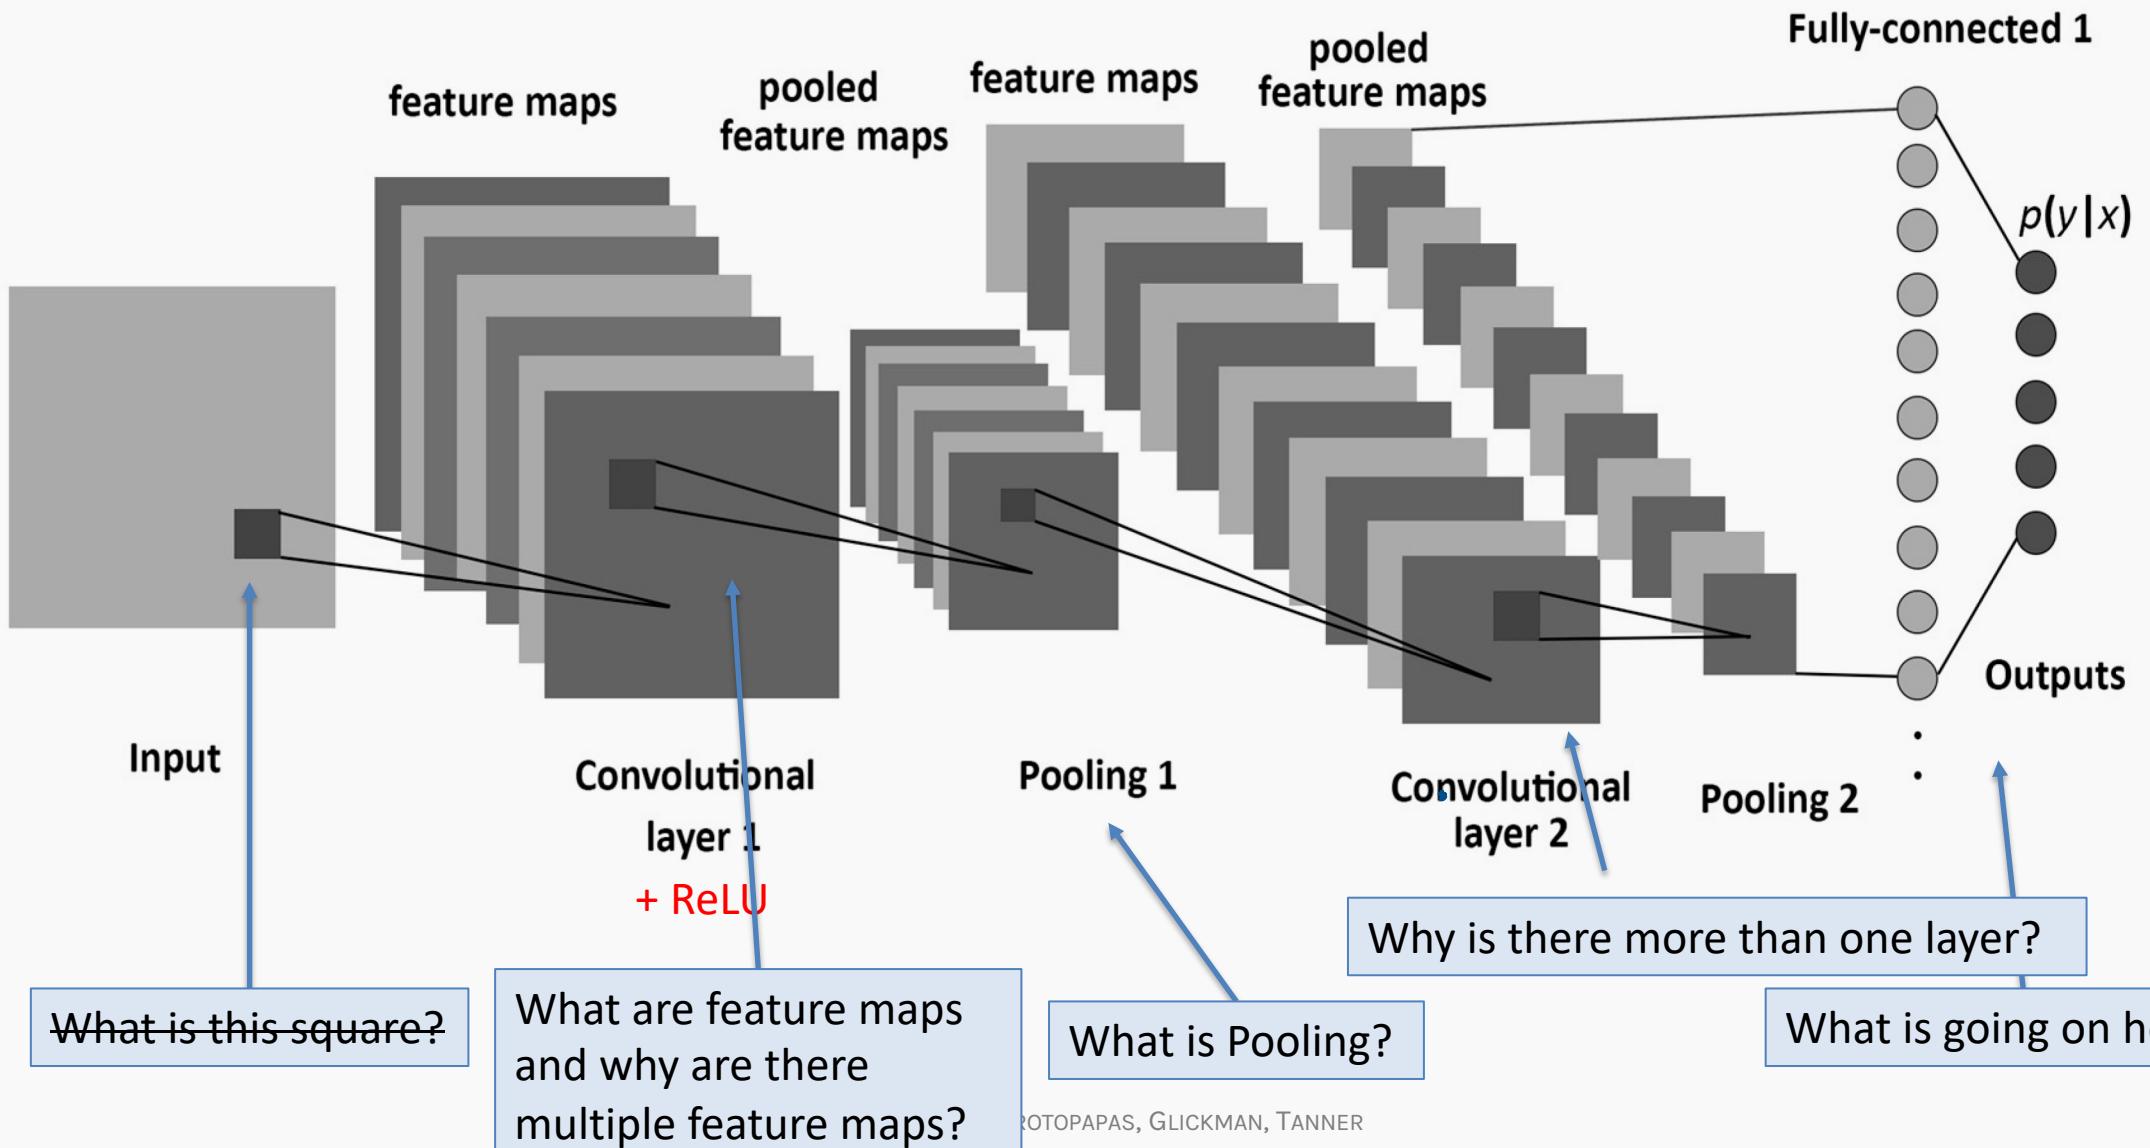
Sharpen

$$* \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix} =$$
A photograph of the man with glasses and a beard, where the edges appear more pronounced and the overall contrast is higher.

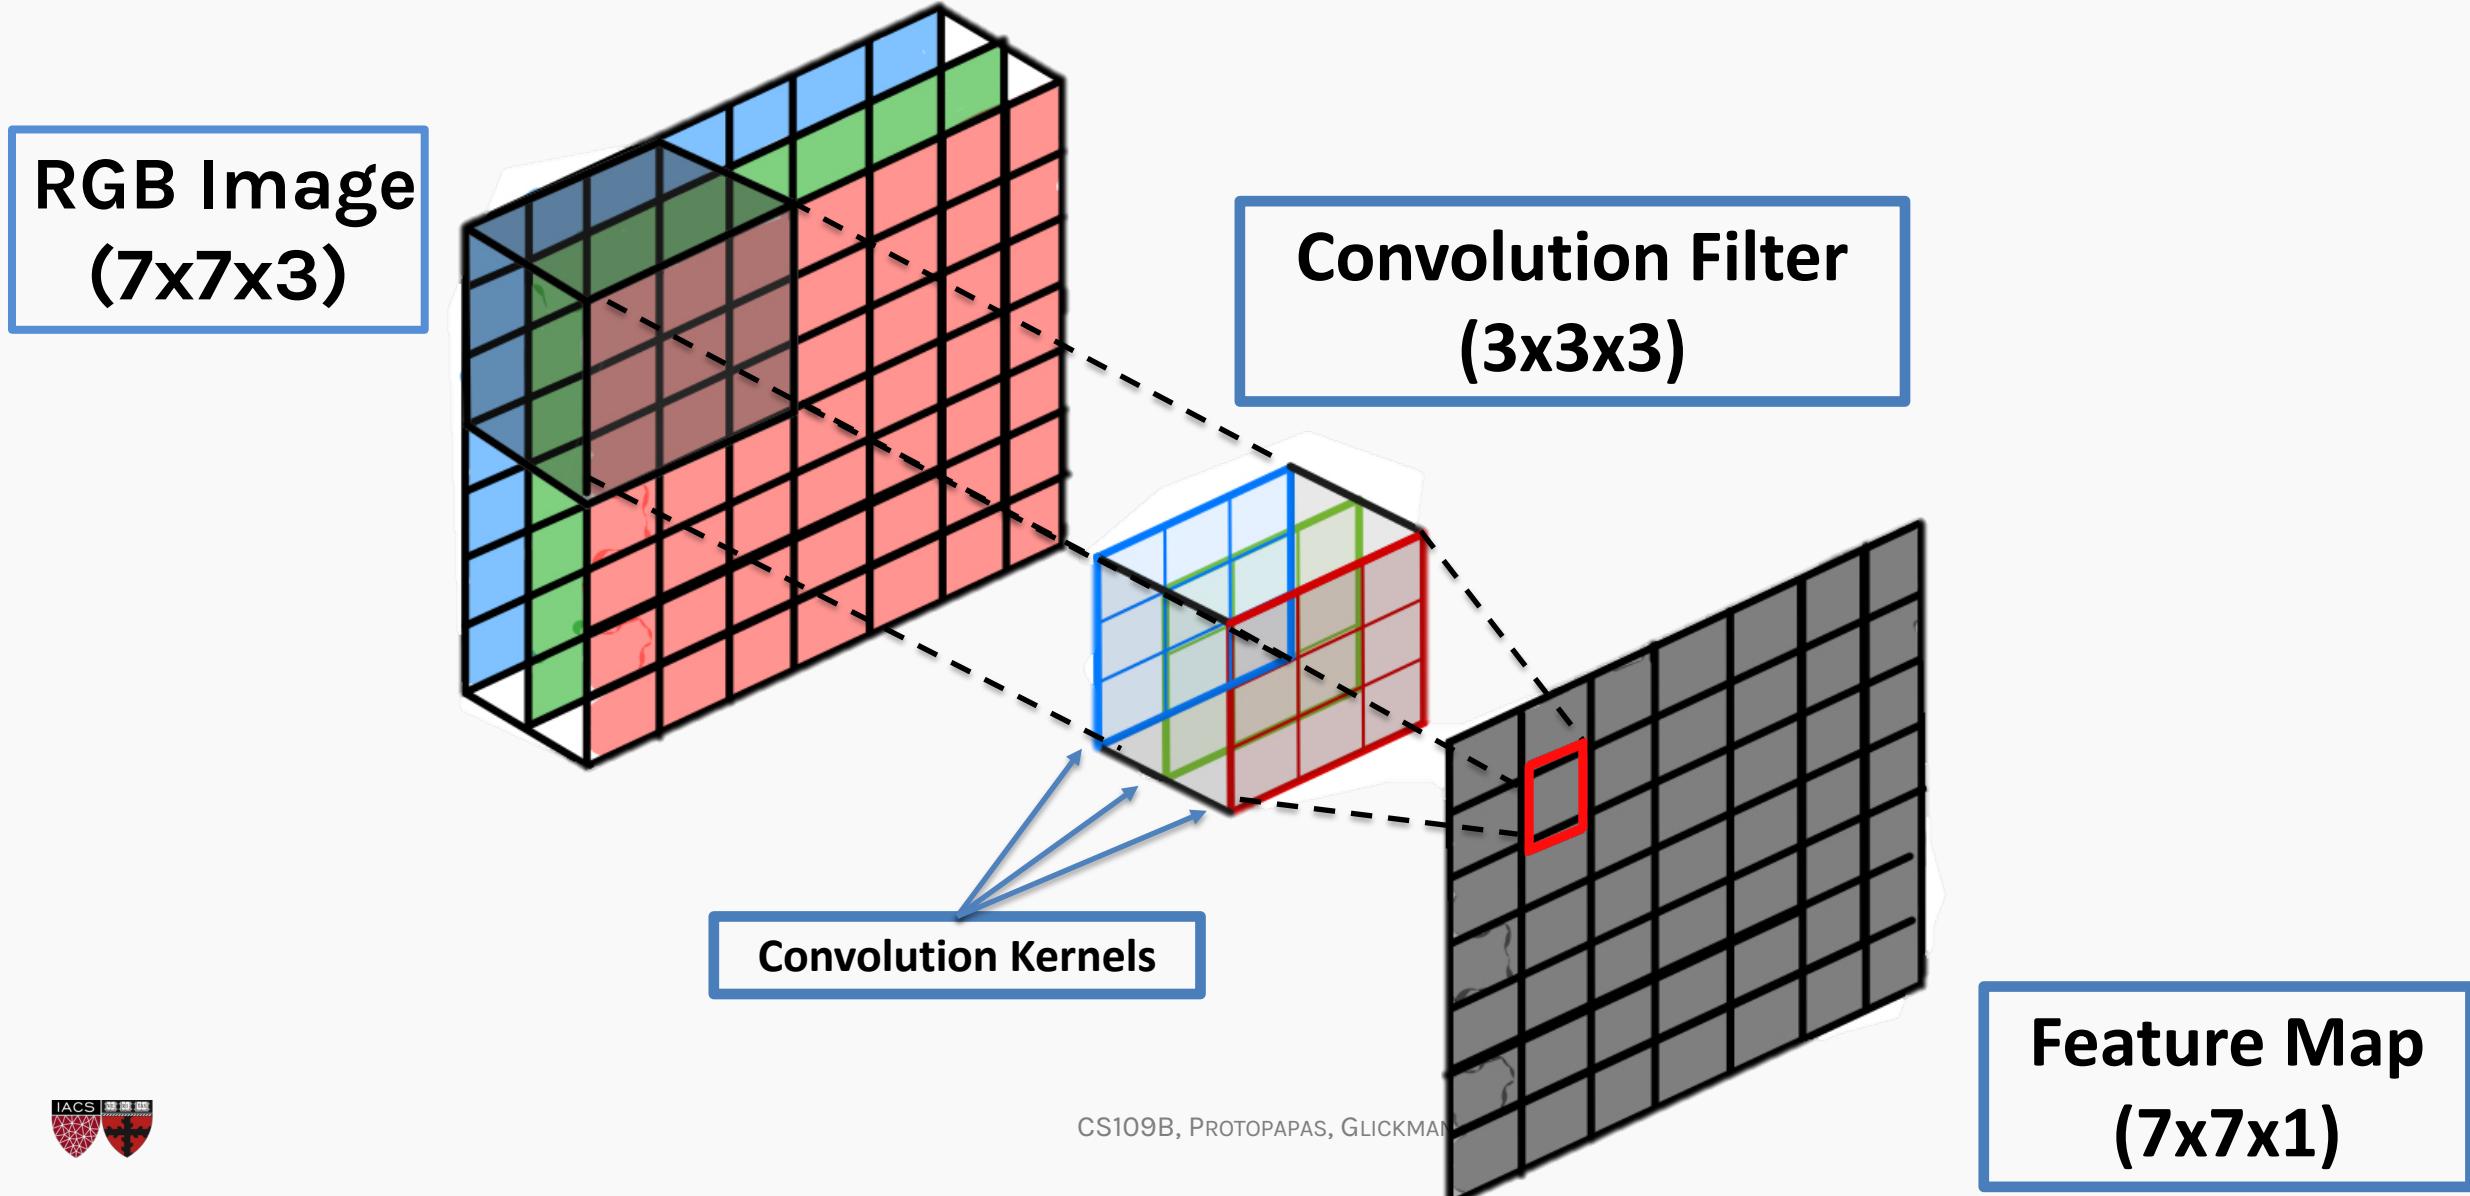
A Convolutional Network



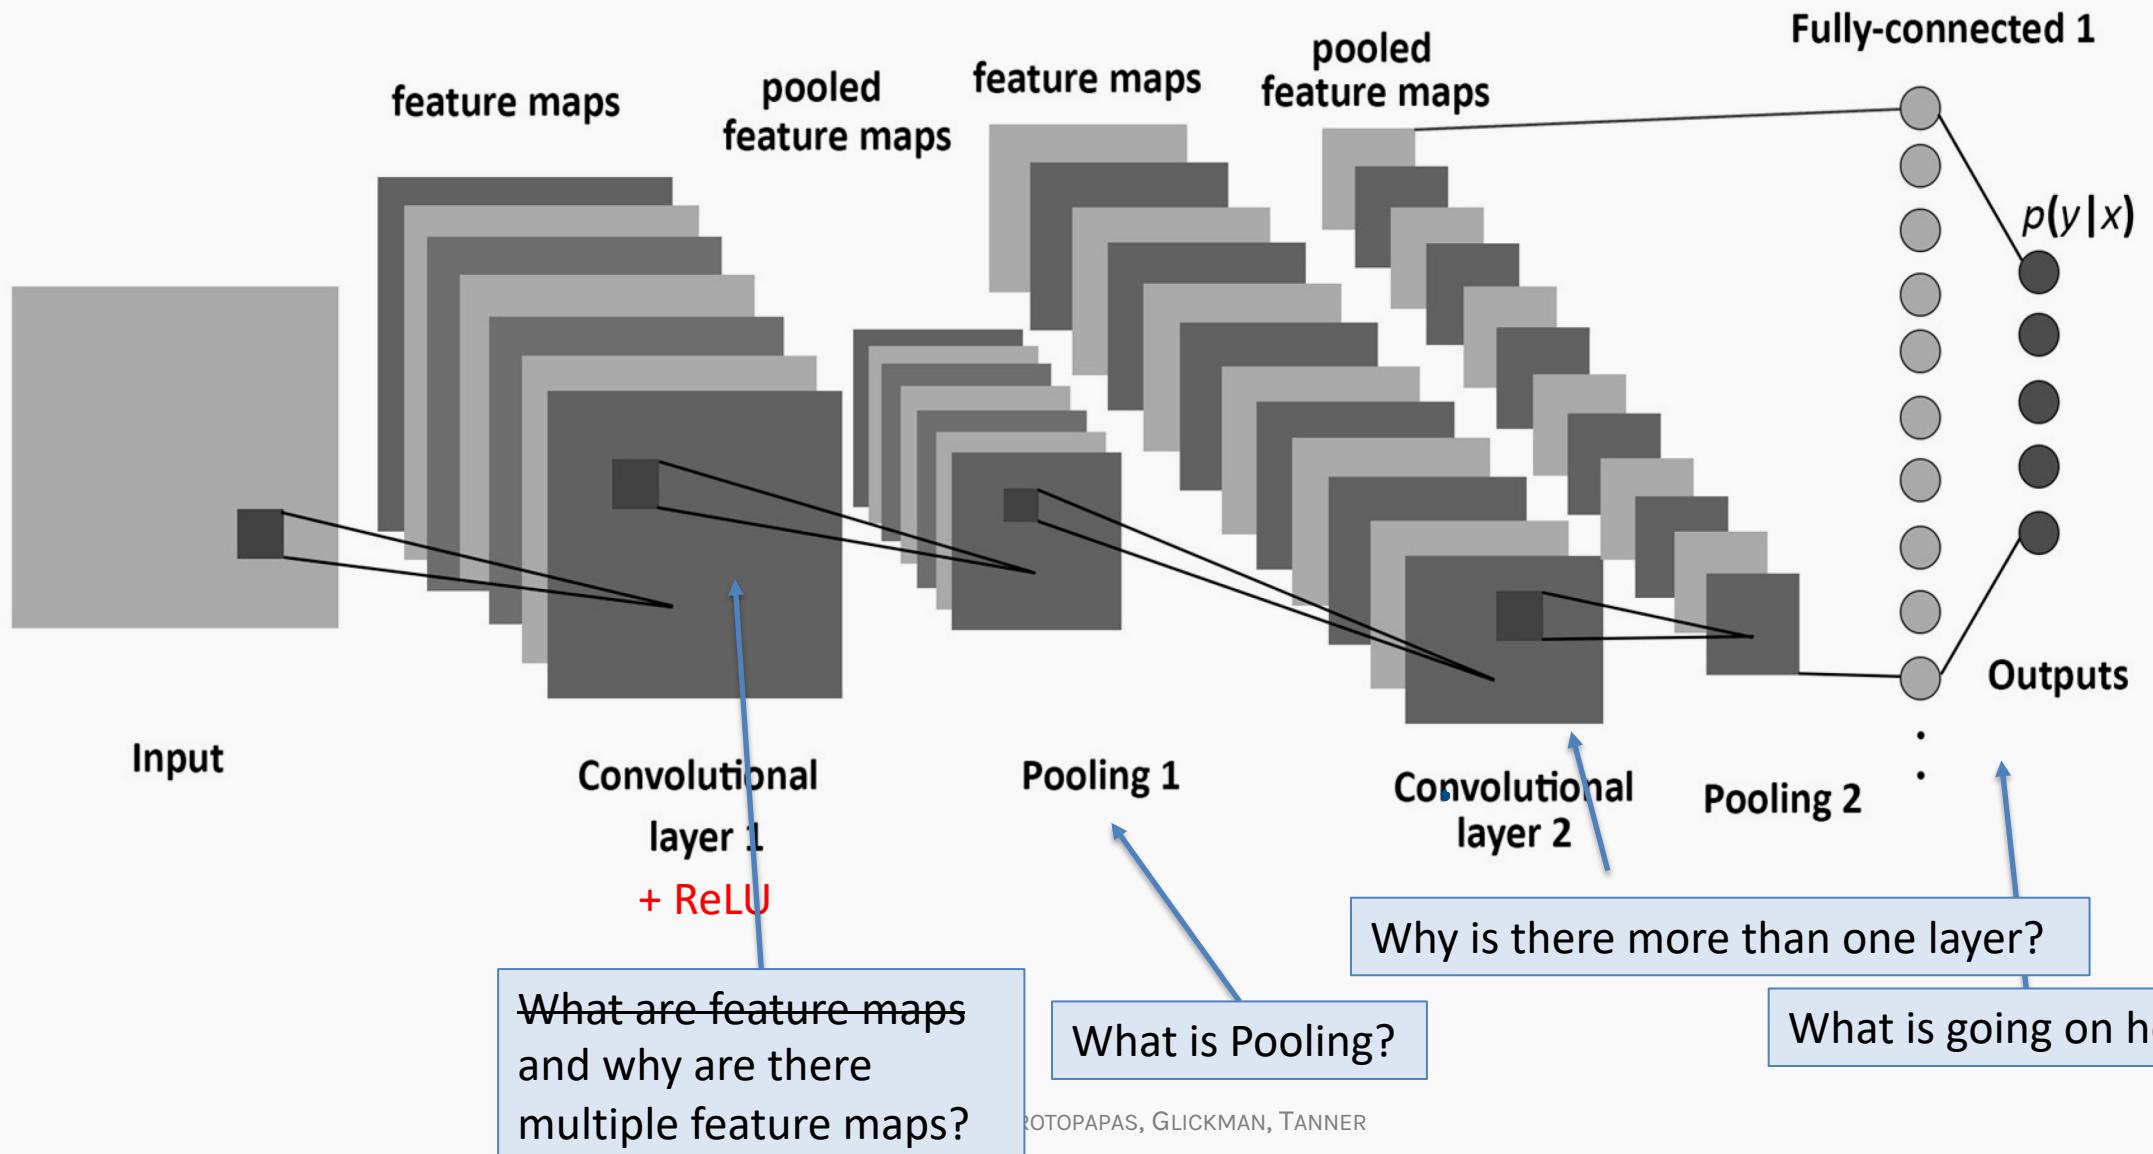
A Convolutional Network



“Convolution” Operation



A Convolutional Network



Why more than one feature map?

LAYER 1:



Why more than one feature map?



LAYER 1:

Filter 1: Horizontal Lines

Why more than one feature map?



LAYER 1:

Filter 1: Horizontal Lines

Filter 2: Vertical Lines

Why more than one feature map?



LAYER 1:

Filter 1: Horizontal Lines

Filter 2: Vertical Lines

Filter 3: Orange bulb

Why more than one feature map?



LAYER 1:

Filter 1: Horizontal Lines

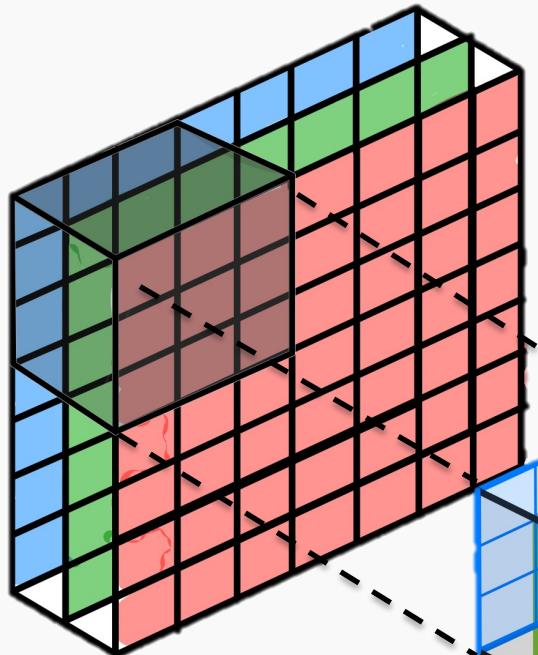
Filter 2: Vertical Lines

Filter 3: Orange bulb

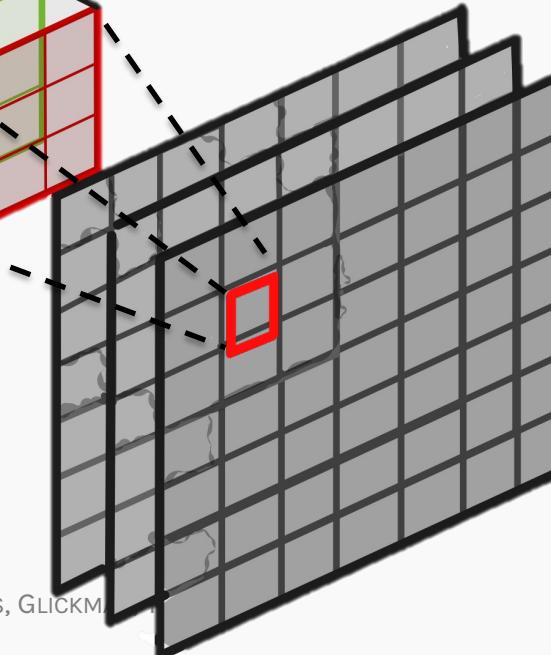
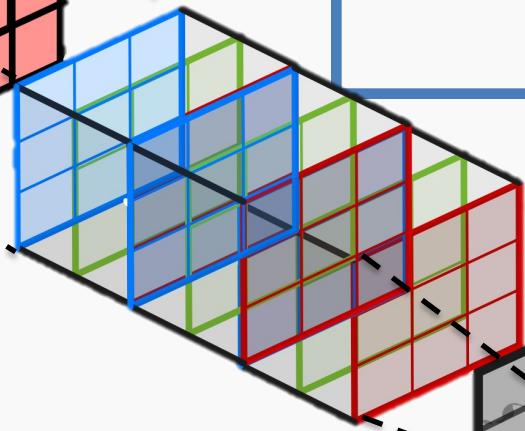
Different filters identify different features.

“Convolution” Operation

RGB Image
(7x7x3)



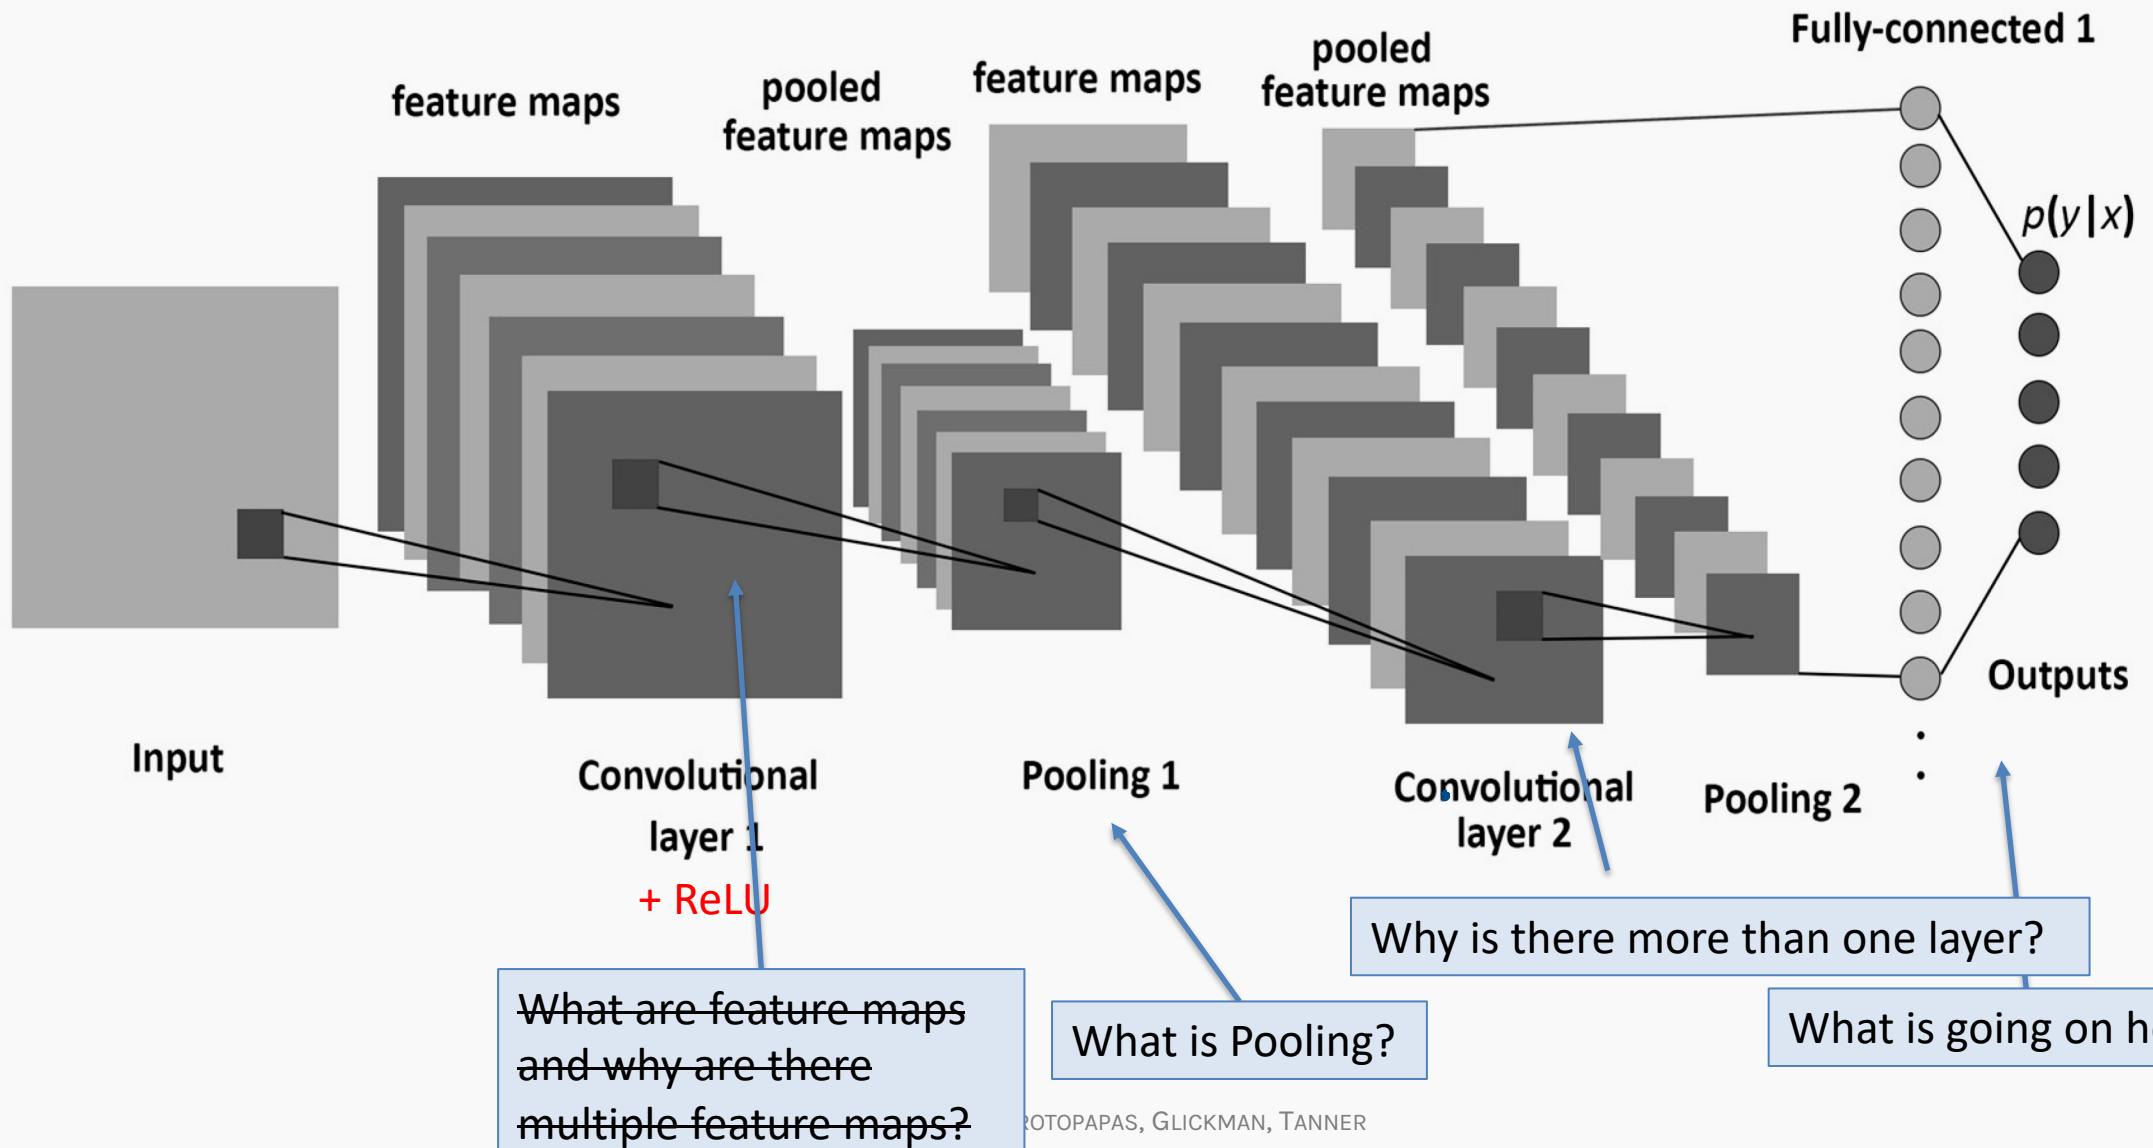
3 Convolution Filters
(3x3x3)



3 Feature Maps
(7x7x3)



A Convolutional Network



Why more than one layer?



Why more than one layer?



Layer 2, Filter 1: Combines
horizontal and vertical lines from
Layer 1 produce diagonal lines.

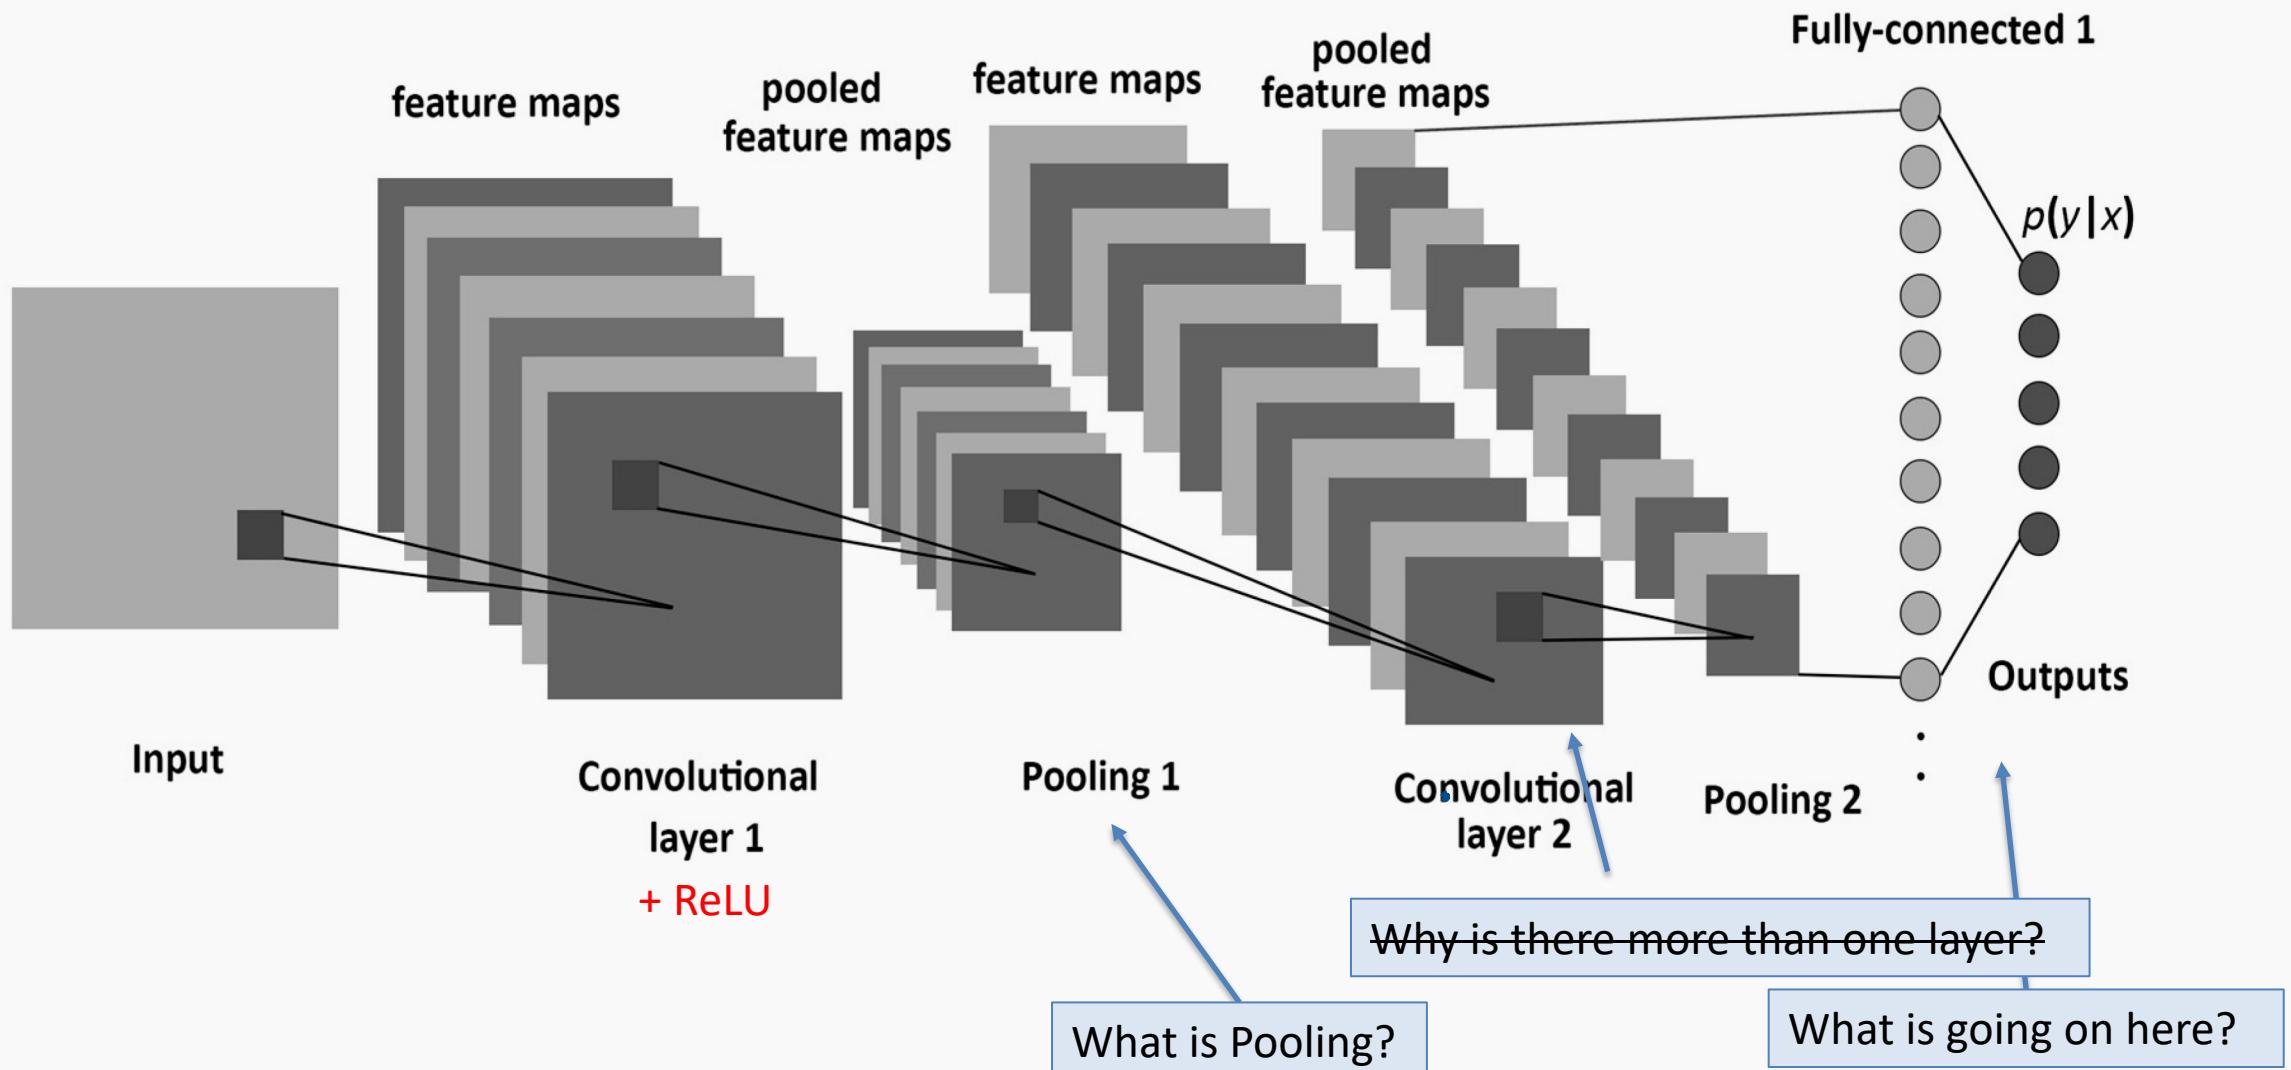
Why more than one layer?



Layer 2, Filter 1: Combines horizontal and vertical lines from Layer 1 produce diagonal lines.

Layer 3, Filter 1: Combines diagonal lines to identify shapes

A Convolutional Network



So far:

We know that MLPs:

- Do not scale well for images
- Ignore the information brought by **pixel position and correlation with neighbors**
- Cannot handle **translations**



So far:

We know that MLPs:

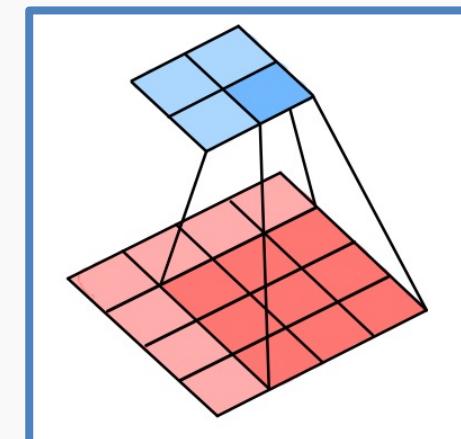
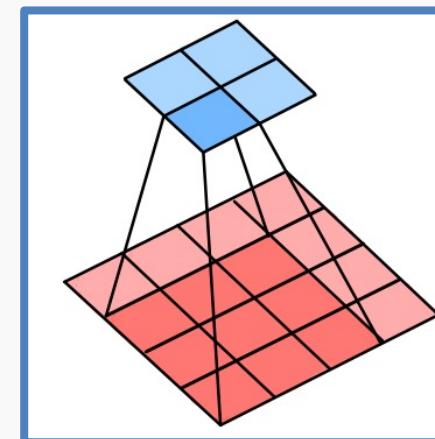
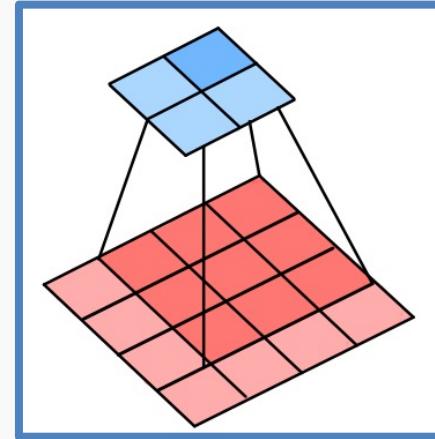
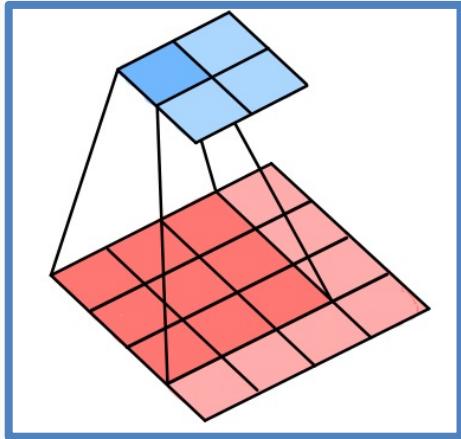
- Do not scale well for images
- Ignore the information brought by pixel position and correlation with neighbors
- Cannot handle translations

The general idea of CNNs is to intelligently adapt to properties of images:

- Pixel position and neighborhood have semantic meanings.
- Elements of interest can appear anywhere in the image.

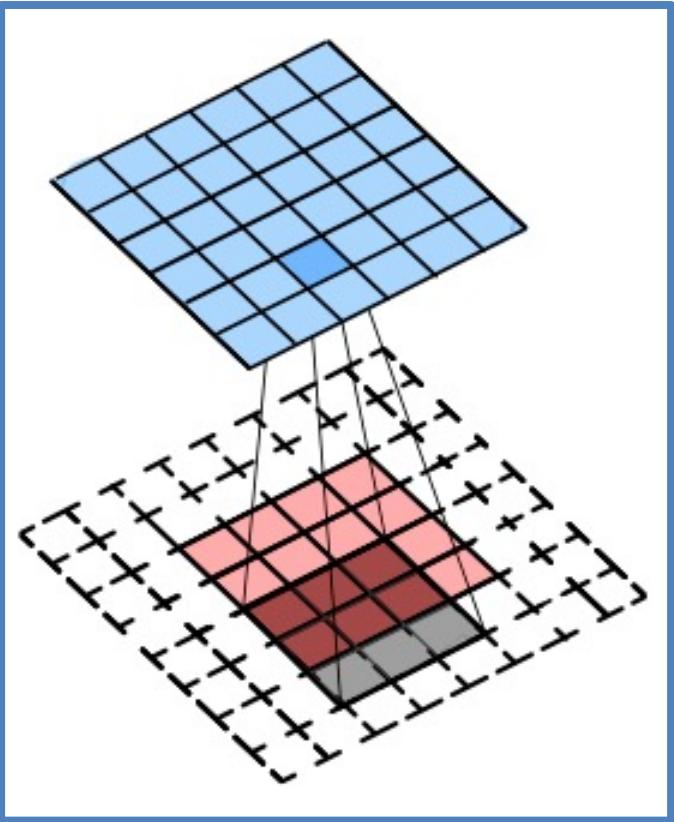
Convolutions – what happens at the edges?

If we apply convolutions on a normal image, the result will be down-sampled by an amount depending on the size of the filter.

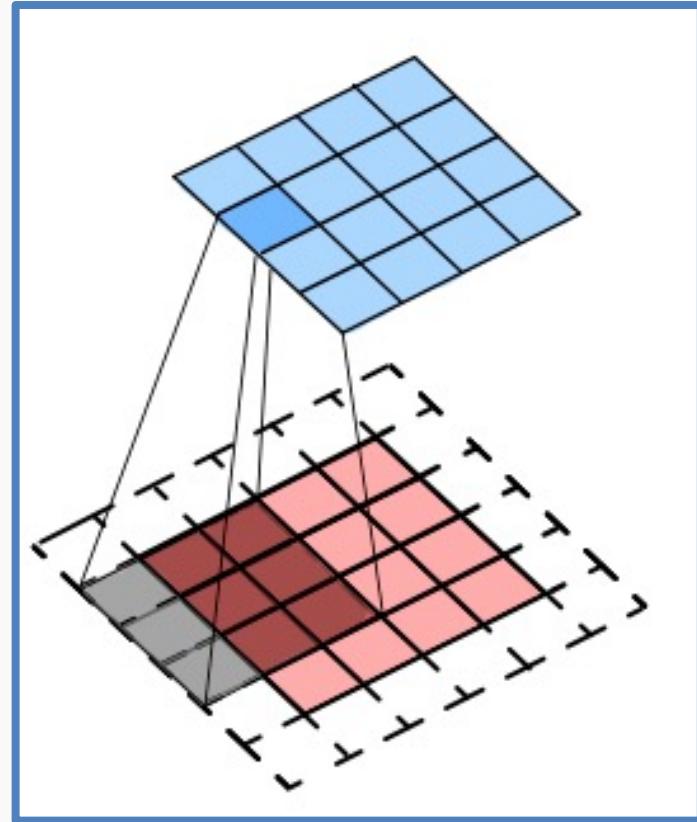


We can avoid this by padding the edges in different ways.

Padding



Full padding. Introduces zeros such that all pixels are visited the same number of times by the filter. Increases size of output.



Same padding. Ensures that the output has the same size as the input.

Stride

Stride controls how the filter convolves around the input volume.

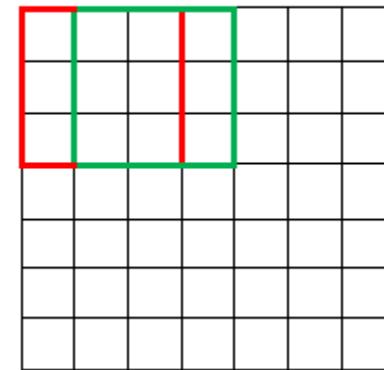
The formula for calculating the output size is:

$$O = \frac{W - K + 2P}{S} + 1$$

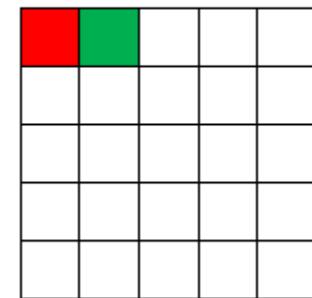
Where O is output dim, W is the input dim, K is the filter size, P is padding and S the stride

Stride = 1

7 x 7 Input Volume

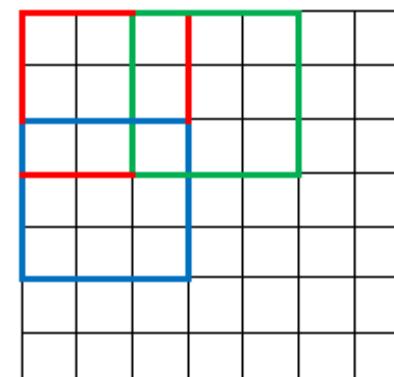


5 x 5 Output Volume

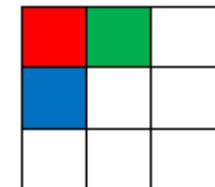


Stride = 2

7 x 7 Input Volume



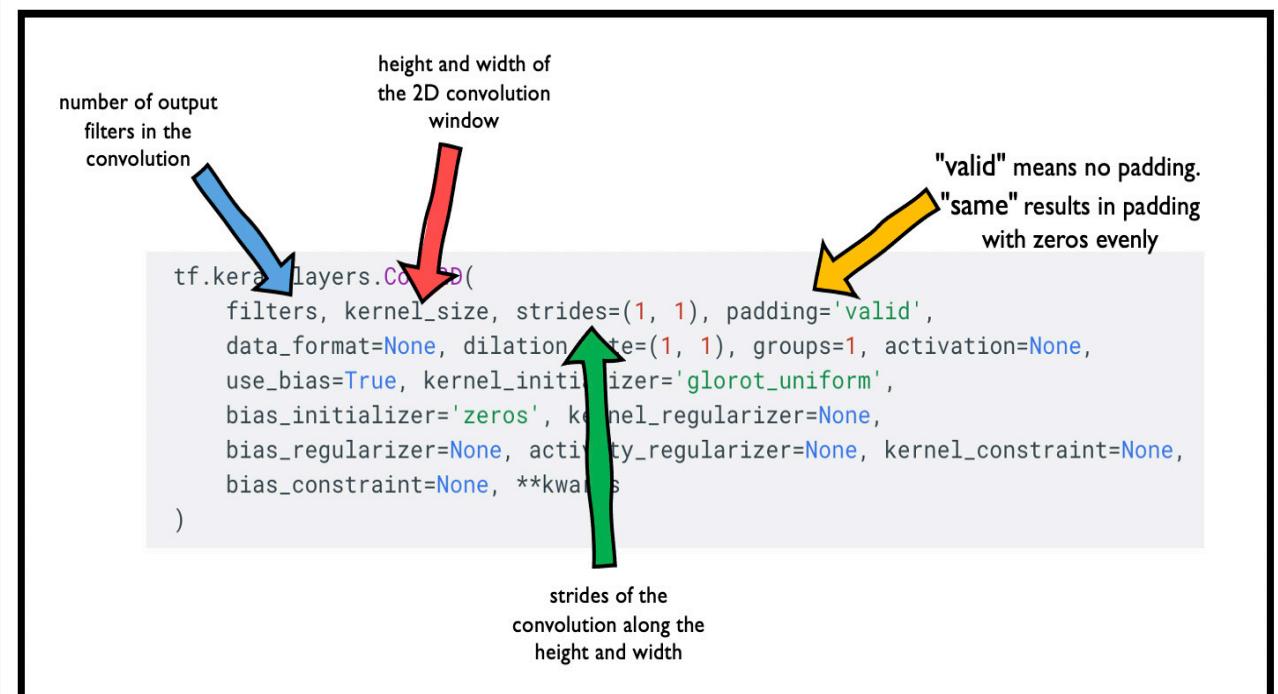
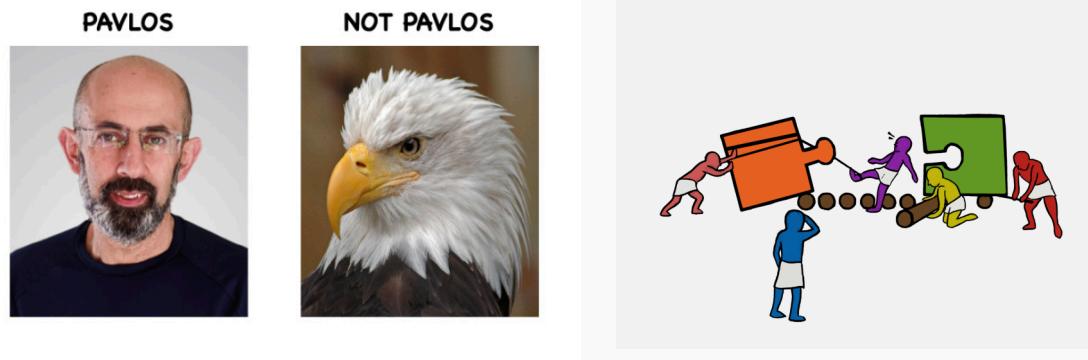
3 x 3 Output Volume



Exercise: Pavlos vs Not Pavlos

The aim of this exercise is to train a dense neural network and a CNN to compare the parameters between them

- Augment the dataset since we only have one image of Pavlos and the eagle
- Build a simple feed-forward network and train it
- Use the convolution layer to build a simple CNN and train it like the network before
- Compare performance and parameters



A Convolutional Network

