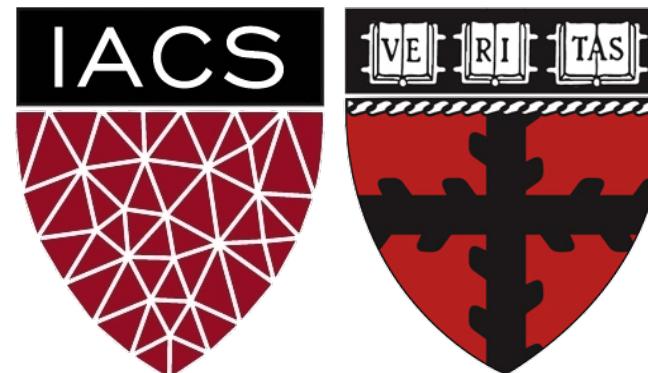


# Transfer Learning

CS109B Data Science 2

Pavlos Protopapas, Mark Glickman, and Chris Tanner



# TRANSFER LEARNING



WHAT SOCIETY THINKS I DO

WHAT MY FRIENDS THINK I DO

WHAT INVESTORS THINK I DO



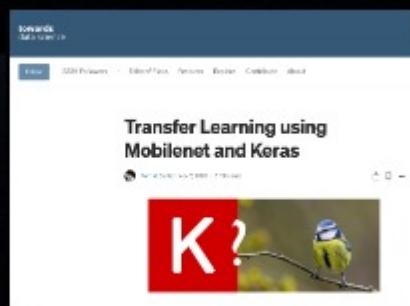
WHAT MY MOM THINKS I DO

imgflip.com



I'm going to learn Jujitsu?

WHAT I THOUGHT I'LL DO



WHAT I ACTUALLY DO

# Outline

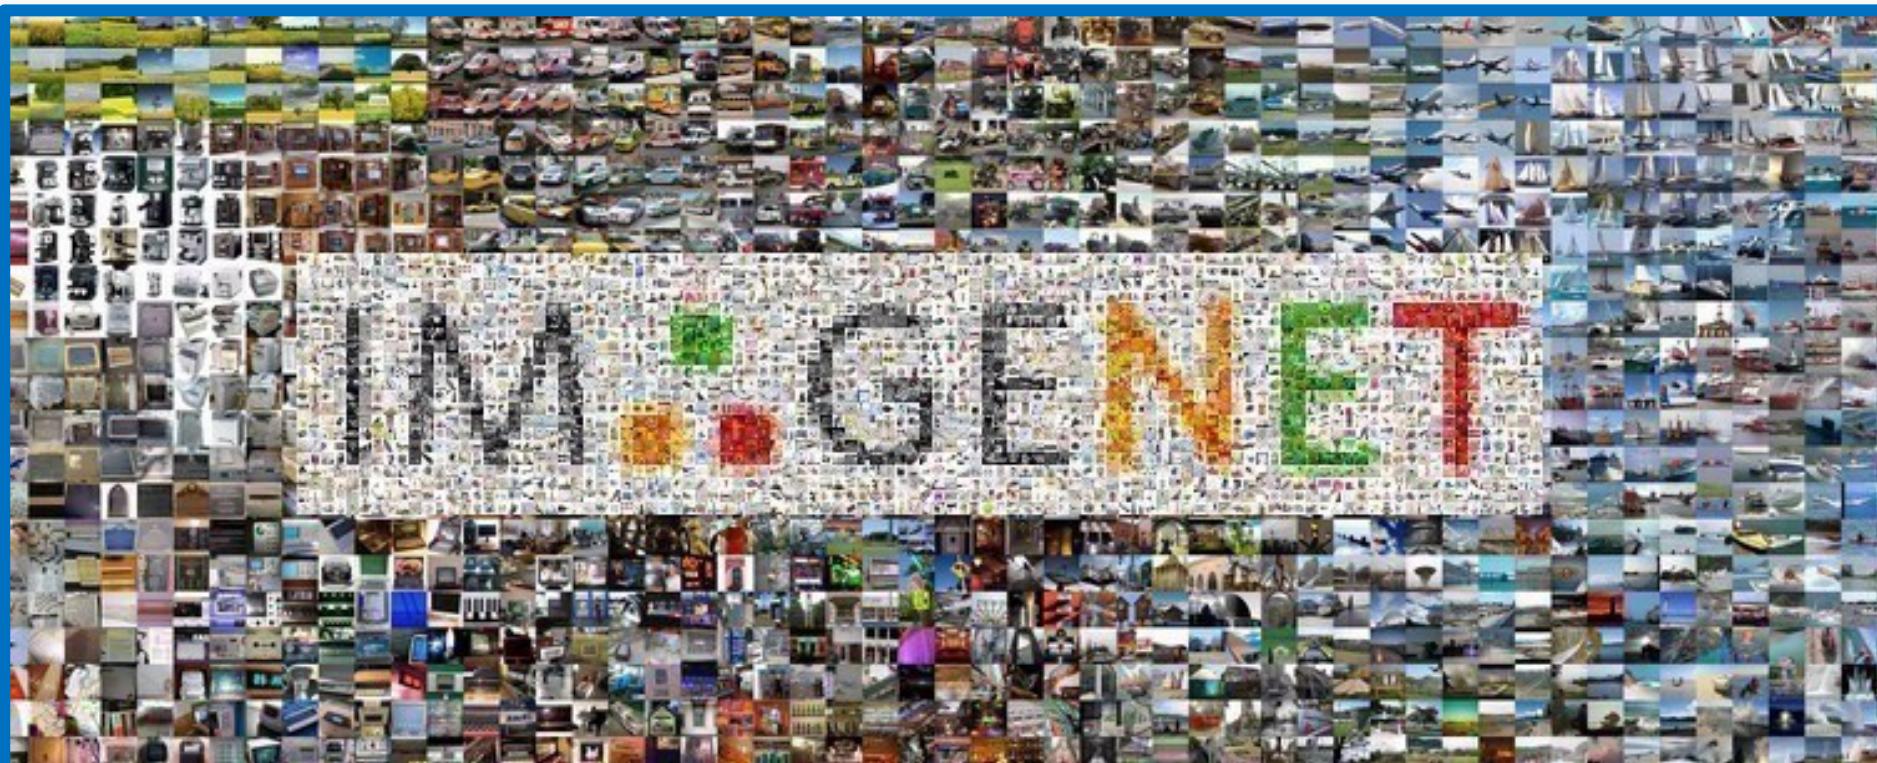
---

1. Review/Questions
2. Motivation
3. The Basics idea for Transfer Learning
4. Representation Learning
5. Transfer Learning Strategies
6. Transfer Learning for Deep Learning

# CNNs: Story so far

## IMAGENET challenge:

- A large visual database designed for use in visual object recognition software research
- More than 14 million images have been hand-annotated by the project to indicate what objects are pictured and in at least one million of the images, bounding boxes are also provided

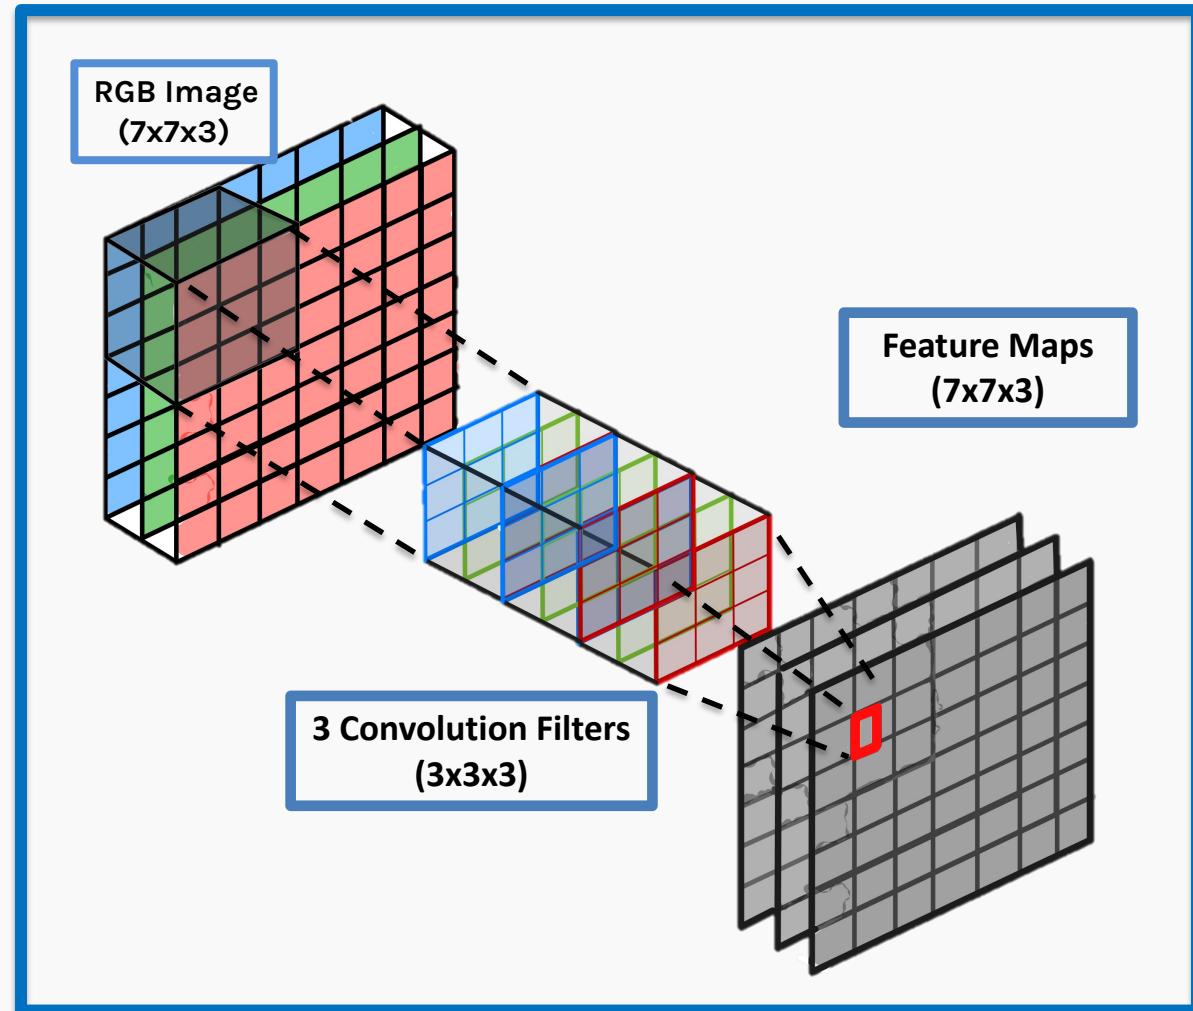


# CNNs: “Convolution” Operation

A **convolutional neural network** typically consists of feature extracting layers and condensing layers.

The feature extracting layers are called **convolutional layers** & each node in these layers uses a small fixed set of weights to transform the image in the way below.

This set of fixed weights for each node in the convolutional layer is often called a **kernel**.

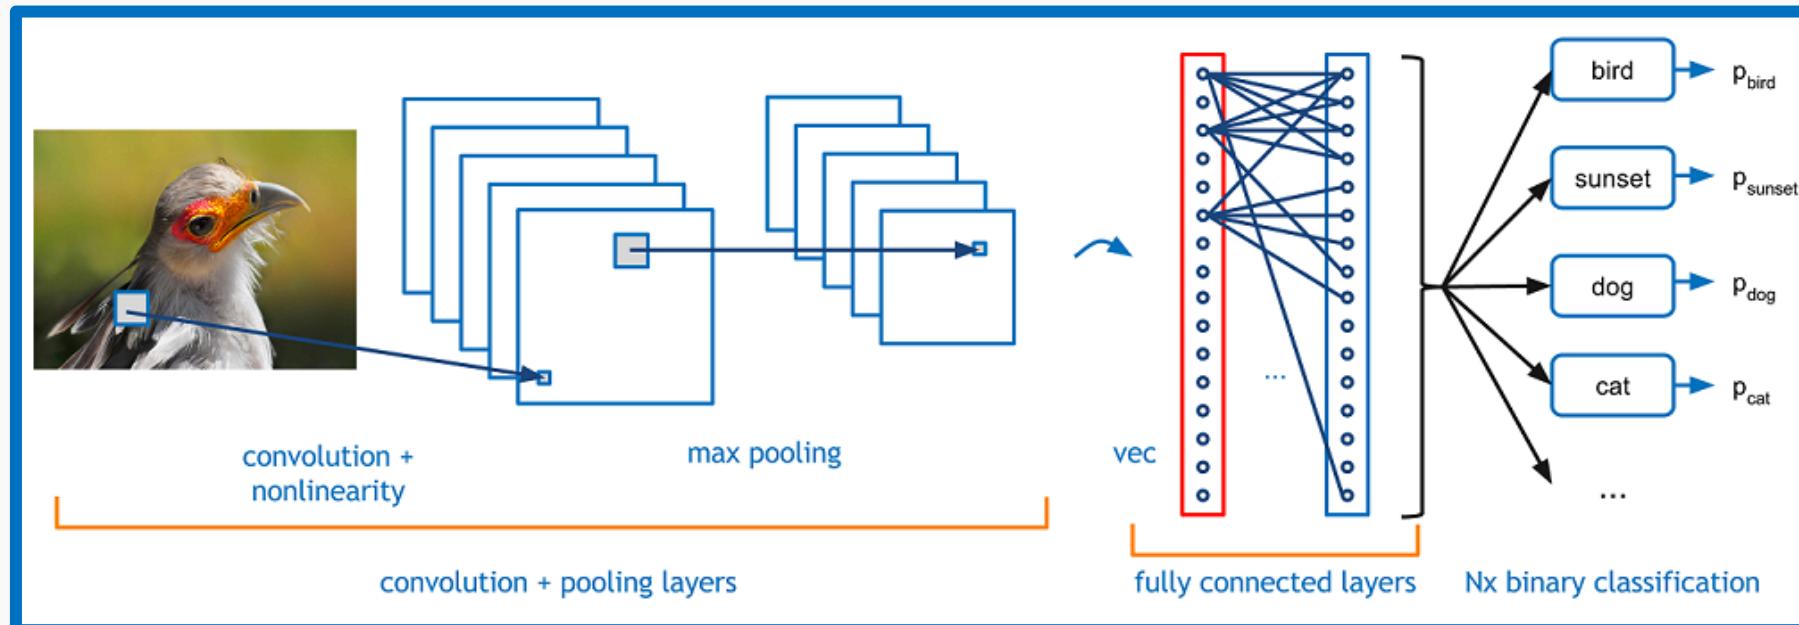


# CNNs: Feature Extraction

Rather than processing image data with a pre-determined set of filters, we want to learn the filters of a CNN for feature extraction. Our goal is to extract features that best helps us to perform our downstream task (e.g. classification).

## Idea:

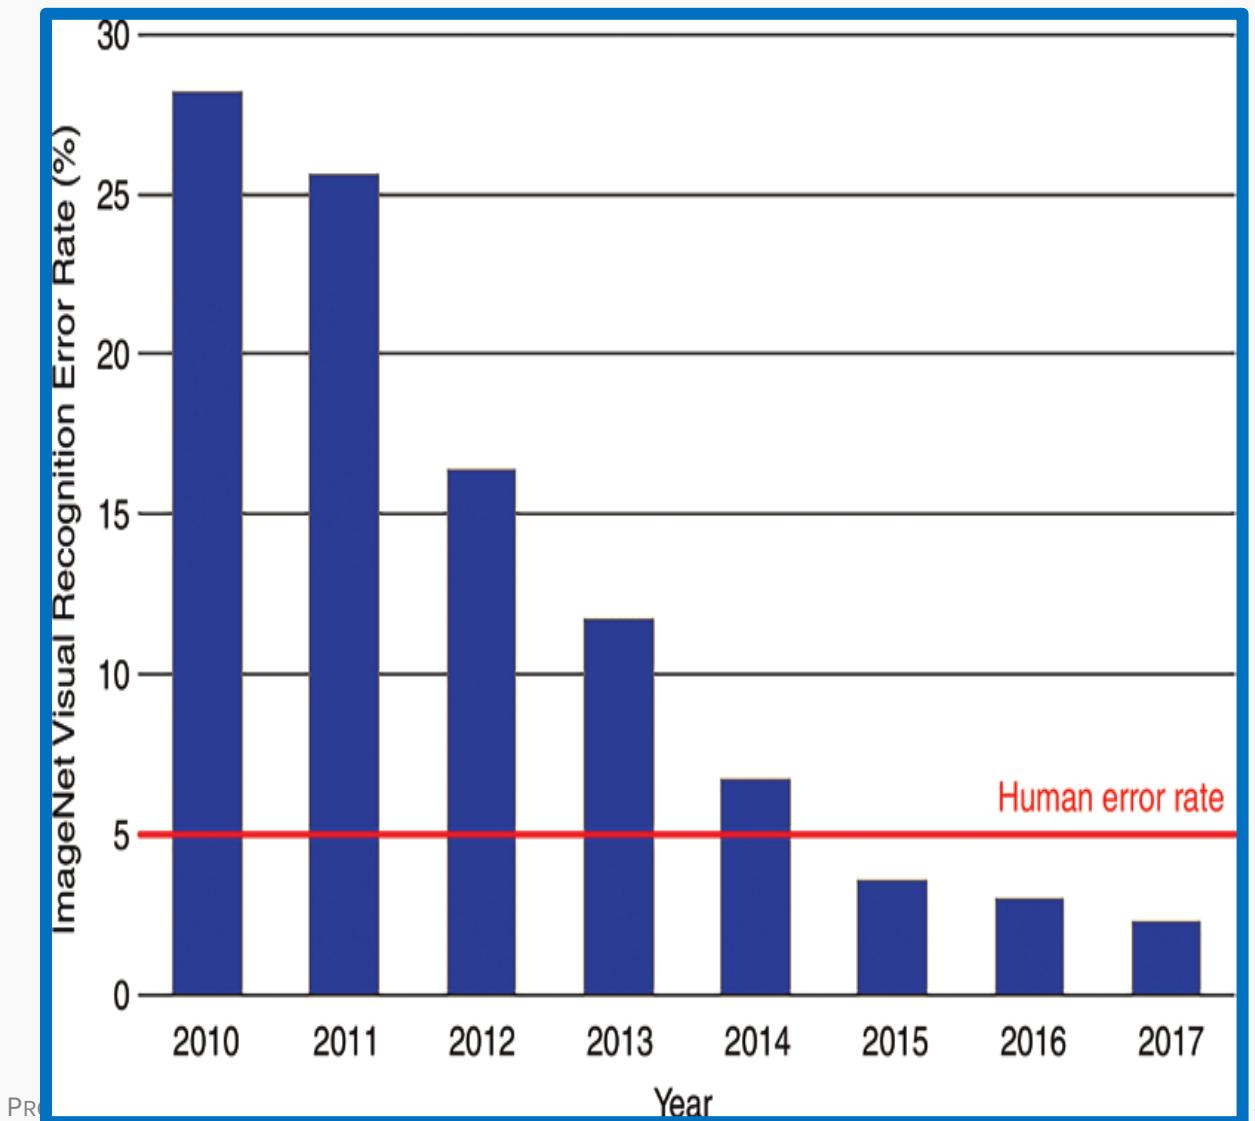
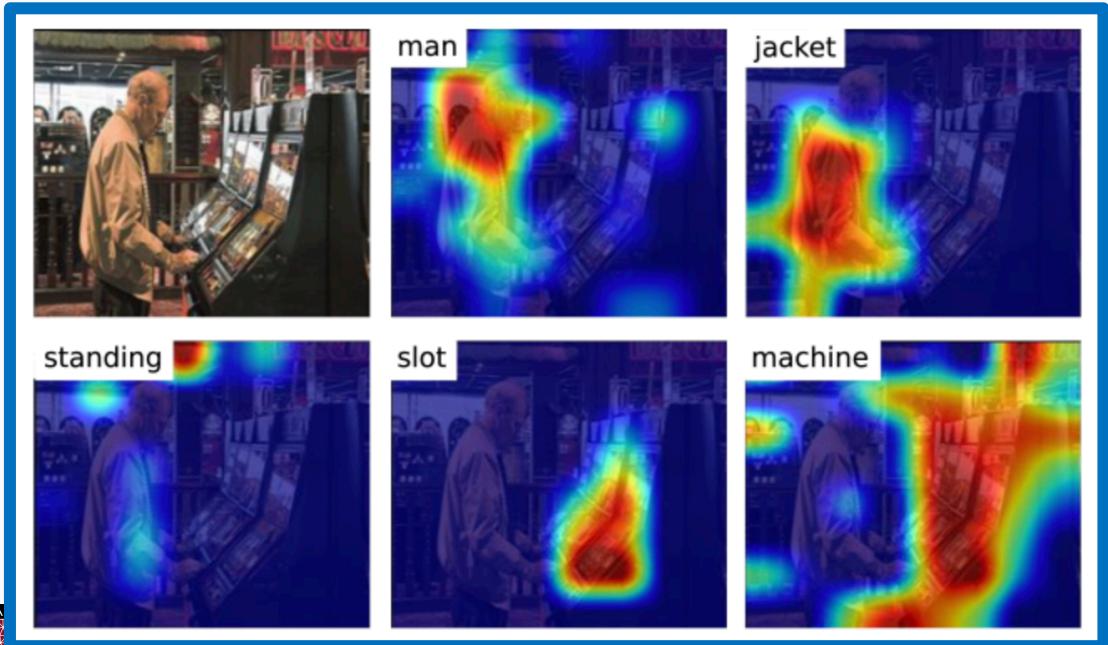
We train a CNN for feature extraction and a model (e.g. MLP, decision tree, logistic regression) for classification, simultaneously and end-to-end.



# CNNs: Successful object detection

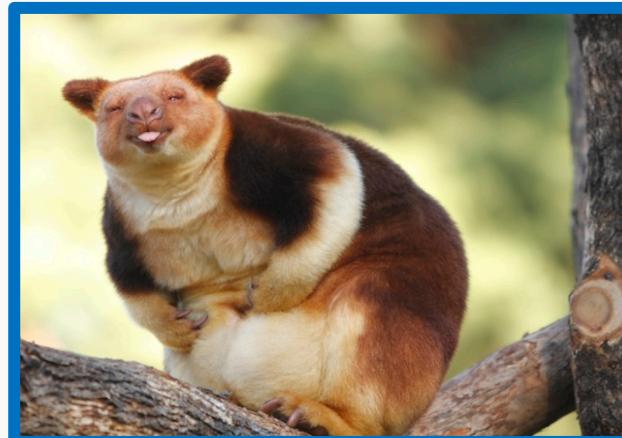
## IMAGENET challenge:

- New model architectures consistently outperform even human error rate
- Ablation studies and saliency maps and confirm models are not overfitting



# CNNs: So what is the problem?

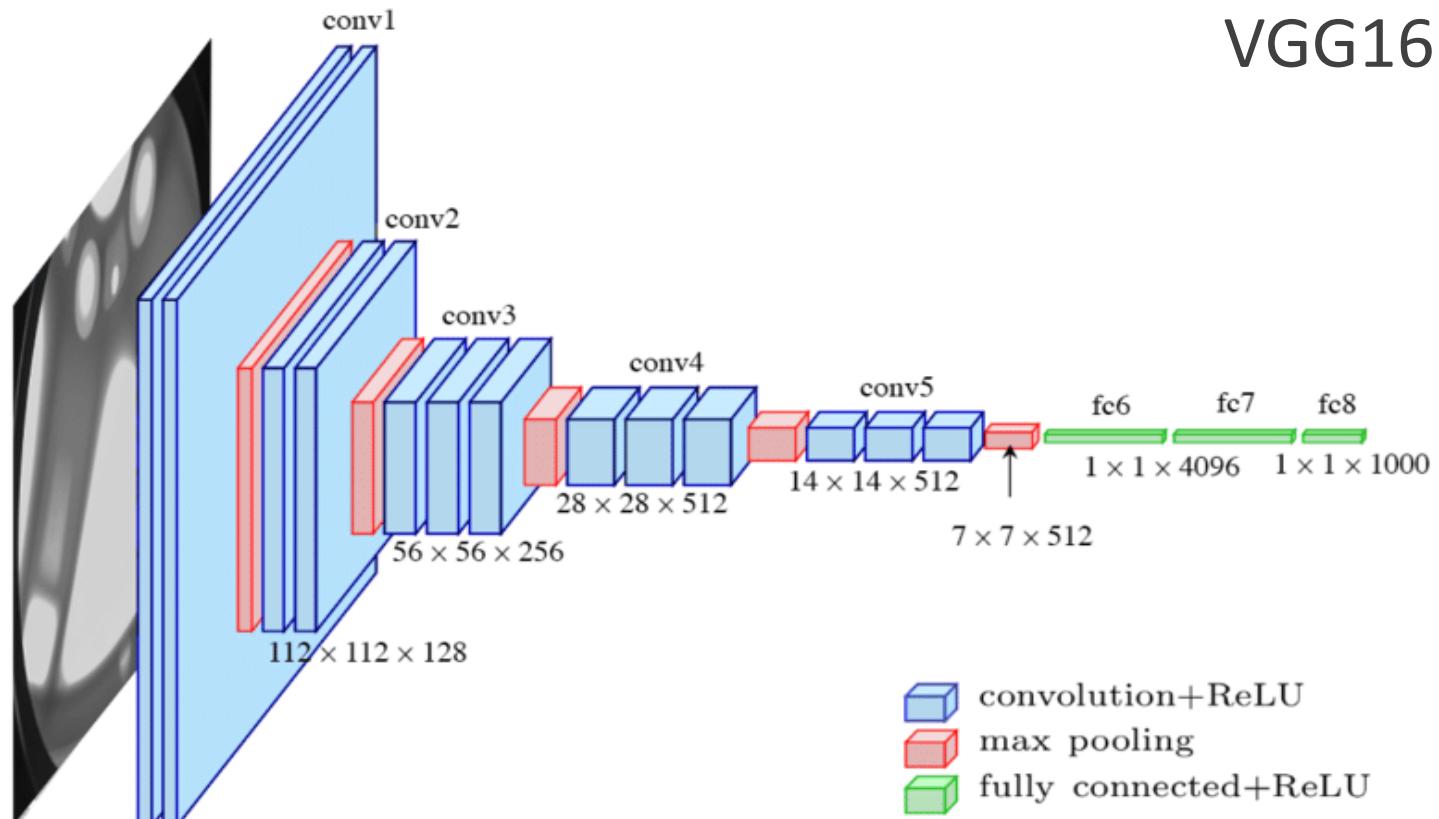
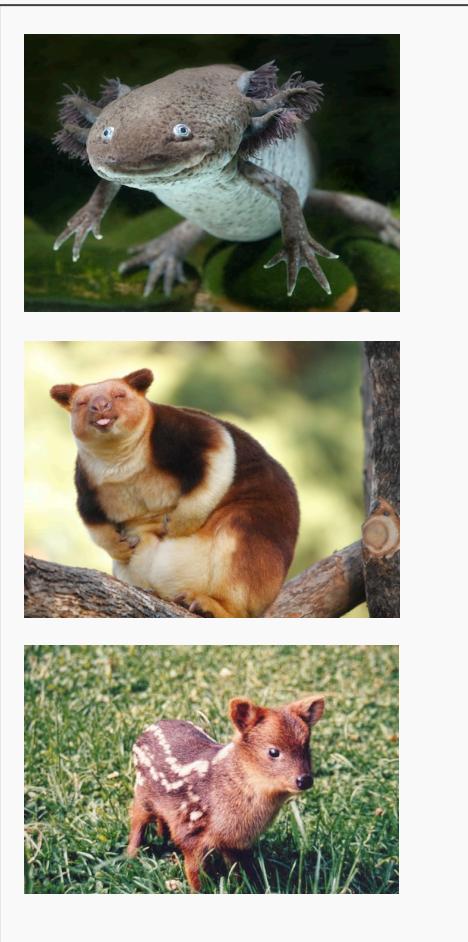
- IMAGENET has more than a 14 million label images and more than 1000 categories.
- However, the imagenet challenge is only a very tiny subset of all possible categories for which we may not have a lot of training data.
- Eg. Can you guess the animals in the images below?



# Classify Rarest Animals



# Classify Rarest Animals

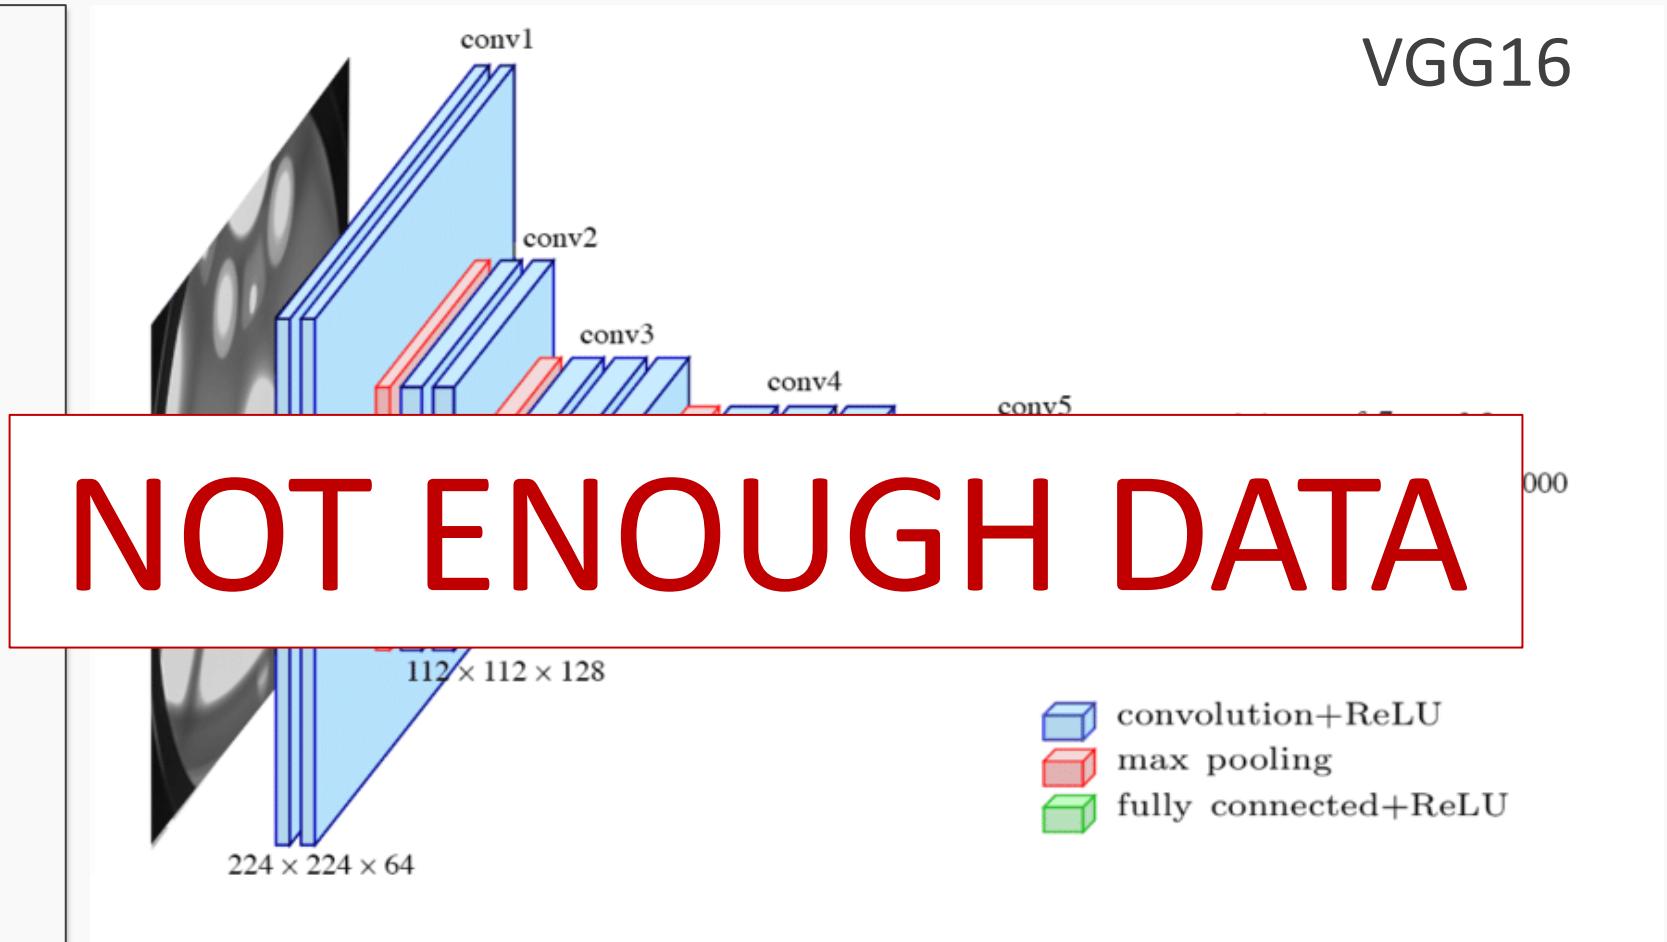
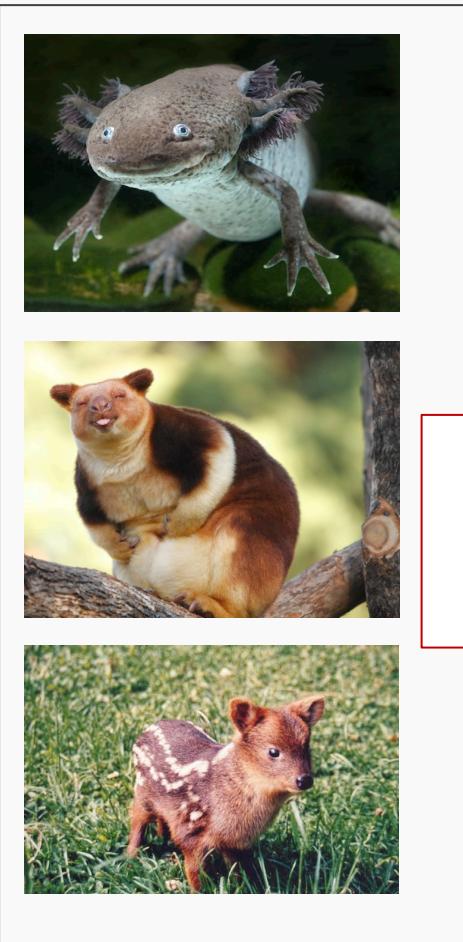


Number of parameters: 134,268,737

Data Set: Few hundred images

CS109B, PROTOPAPAS, GLICKMAN, TANNER

# Classify Rarest Animals

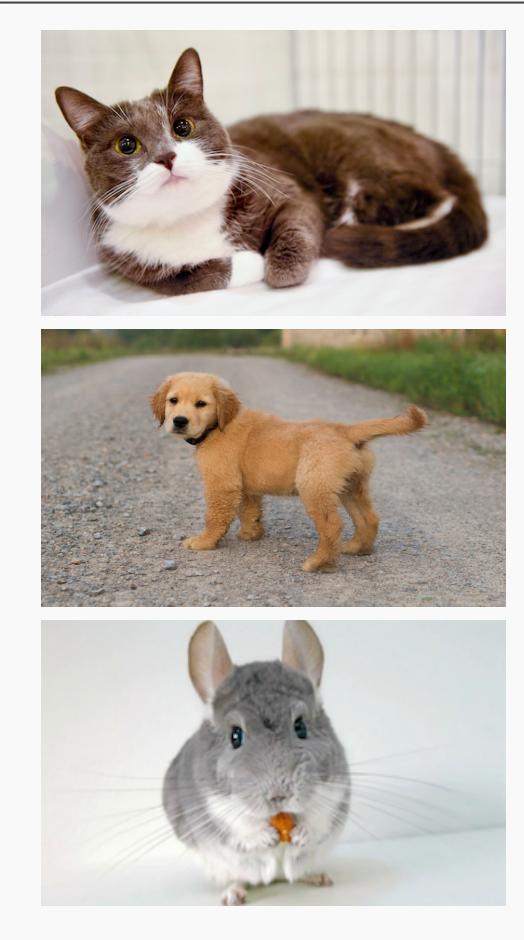


Number of parameters: 134,268,737

Data Set: Few hundred images

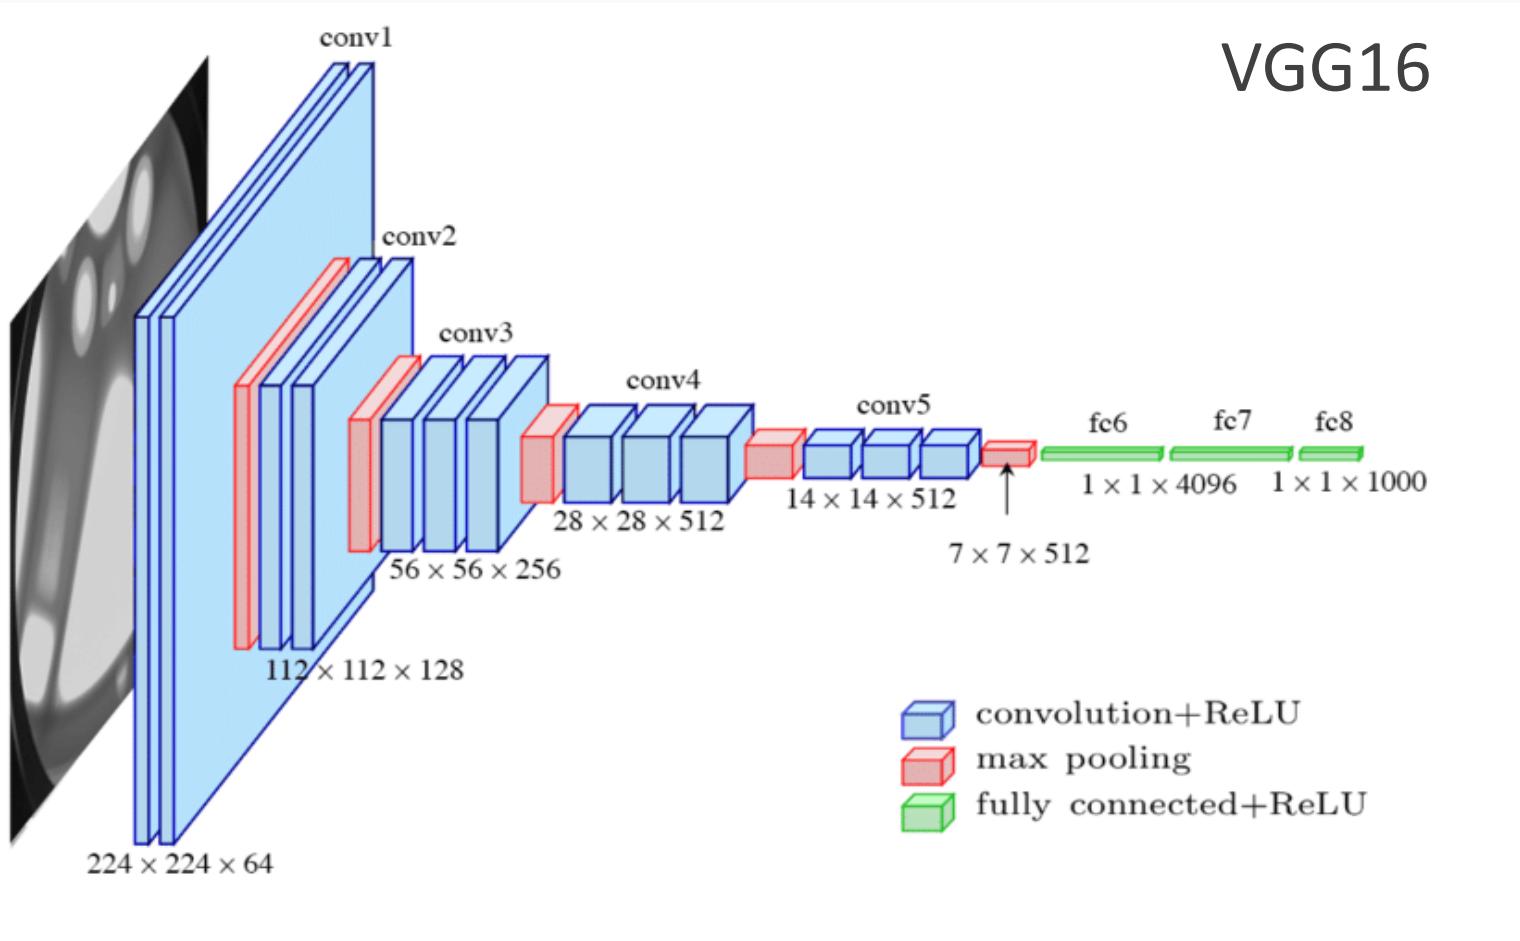
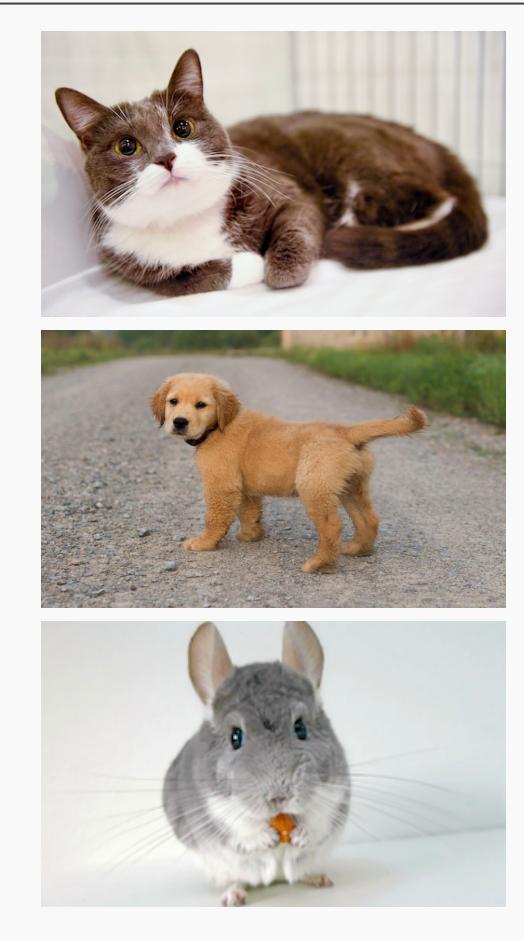
CS109B, PROTOPAPAS, GLICKMAN, TANNER

# Classify Cats, Dogs, Chinchillas etc



VGG16

# Classify Cats, Dogs, Chinchillas etc

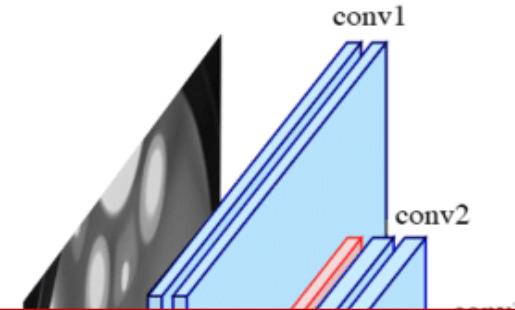
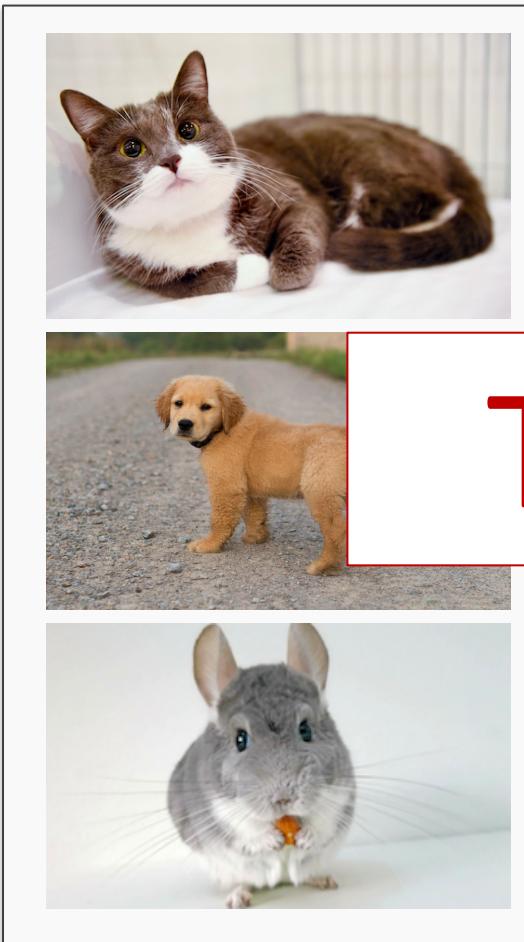


Number of parameters: 134,268,737

Enough training data. ImageNet approximate 1.2M

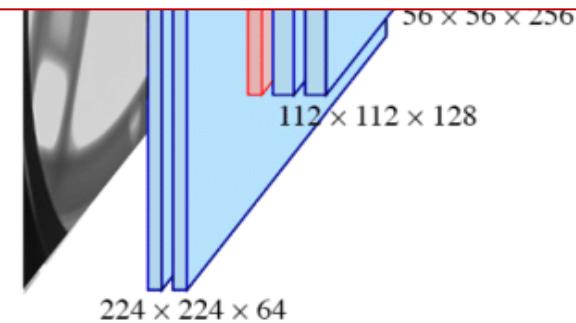
CS109B, PROTOPAPAS, GLICKMAN, TANNER

# Classify Cats, Dogs, Chinchillas etc



VGG16

**TAKES TOO LONG**



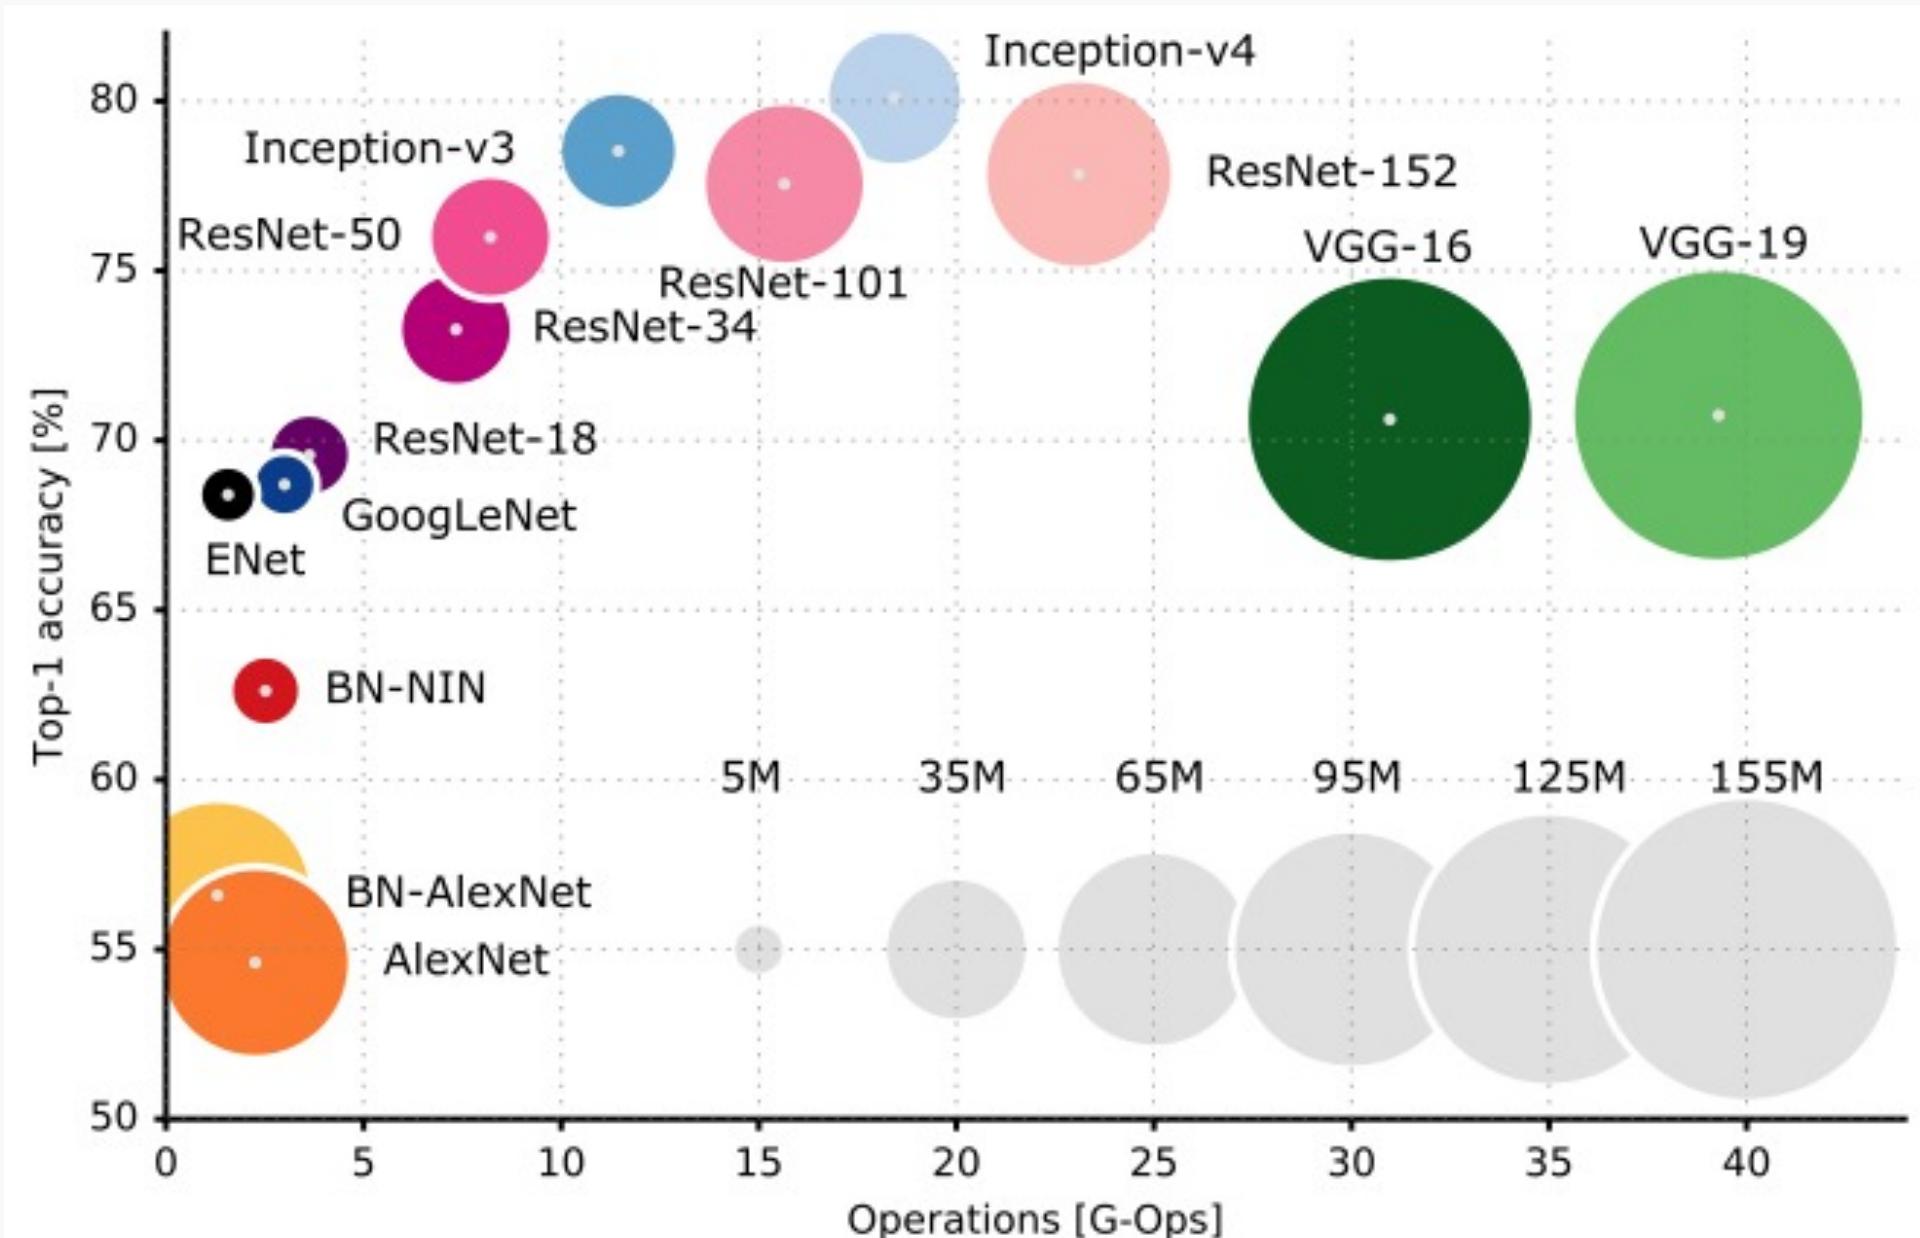
- convolution+ReLU
- max pooling
- fully connected+ReLU

Number of parameters: 134,268,737

Enough training data. ImageNet approximate 1.2M

CS109B, PROTOPAPAS, GLICKMAN, TANNER

# Training time for SOTAs



# Transfer Learning To The Rescue

---

How do you build an image classifier that can be trained in a few minutes on a CPU with very little data?



# Basic idea of Transfer Learning

**Wikipedia:**

Transfer learning (TL) is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

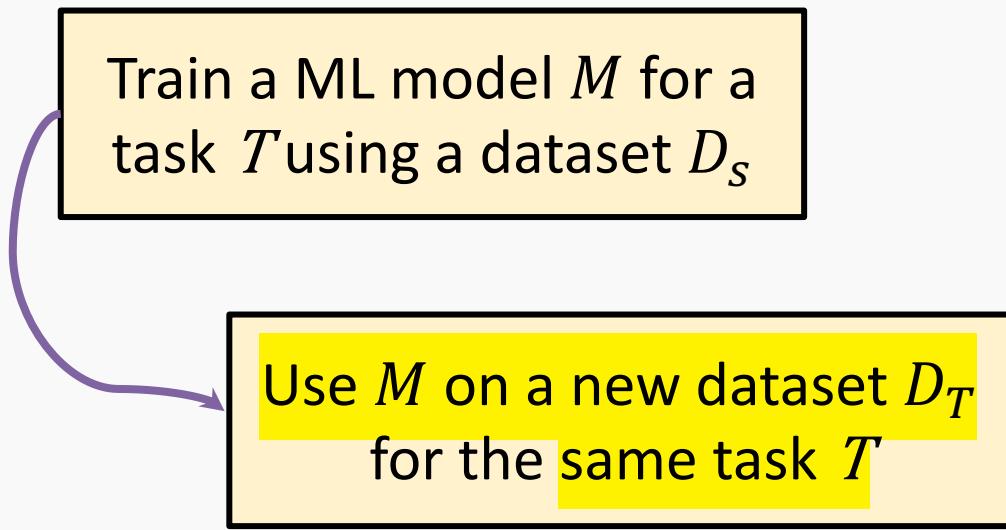
# Basic idea of Transfer Learning

Train a ML model  $M$  for a task  $T$  using a dataset  $D_s$

## Wikipedia:

**Transfer learning (TL)** is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

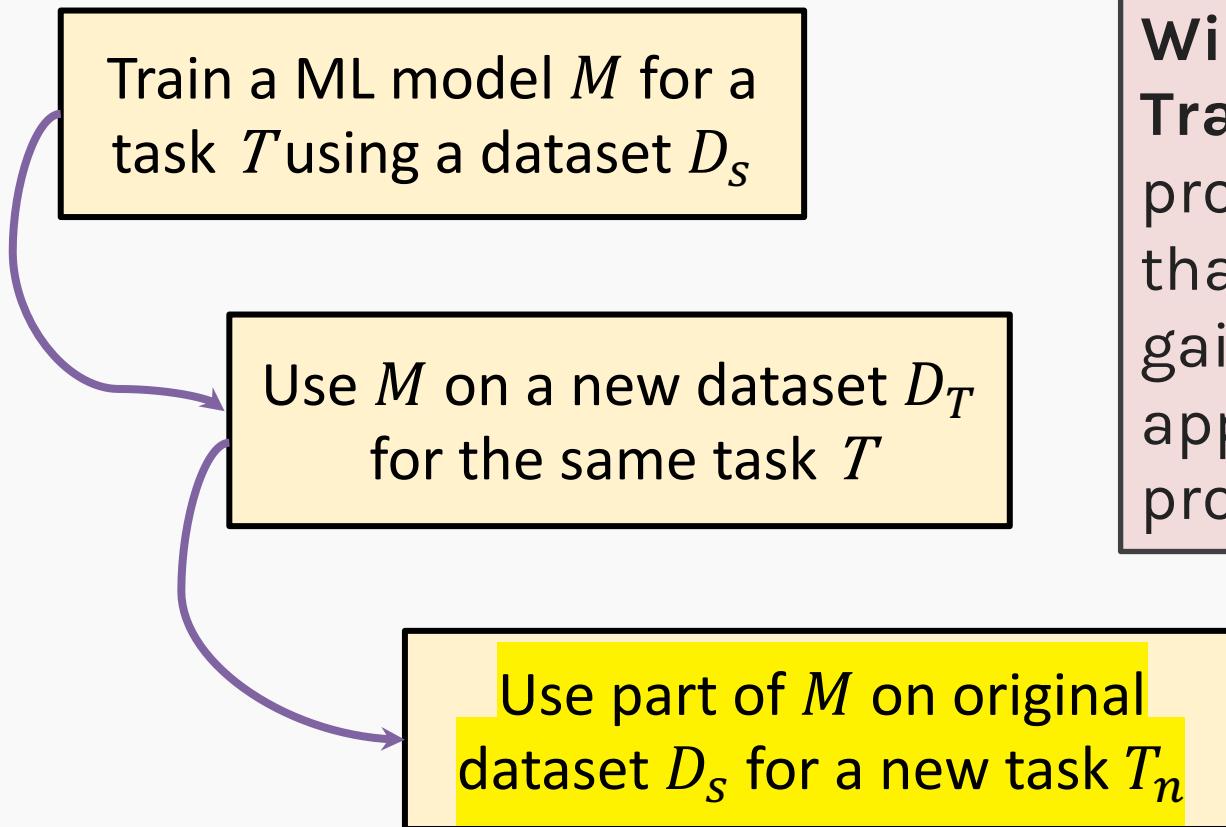
# Basic idea of Transfer Learning



## Wikipedia:

**Transfer learning (TL)** is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

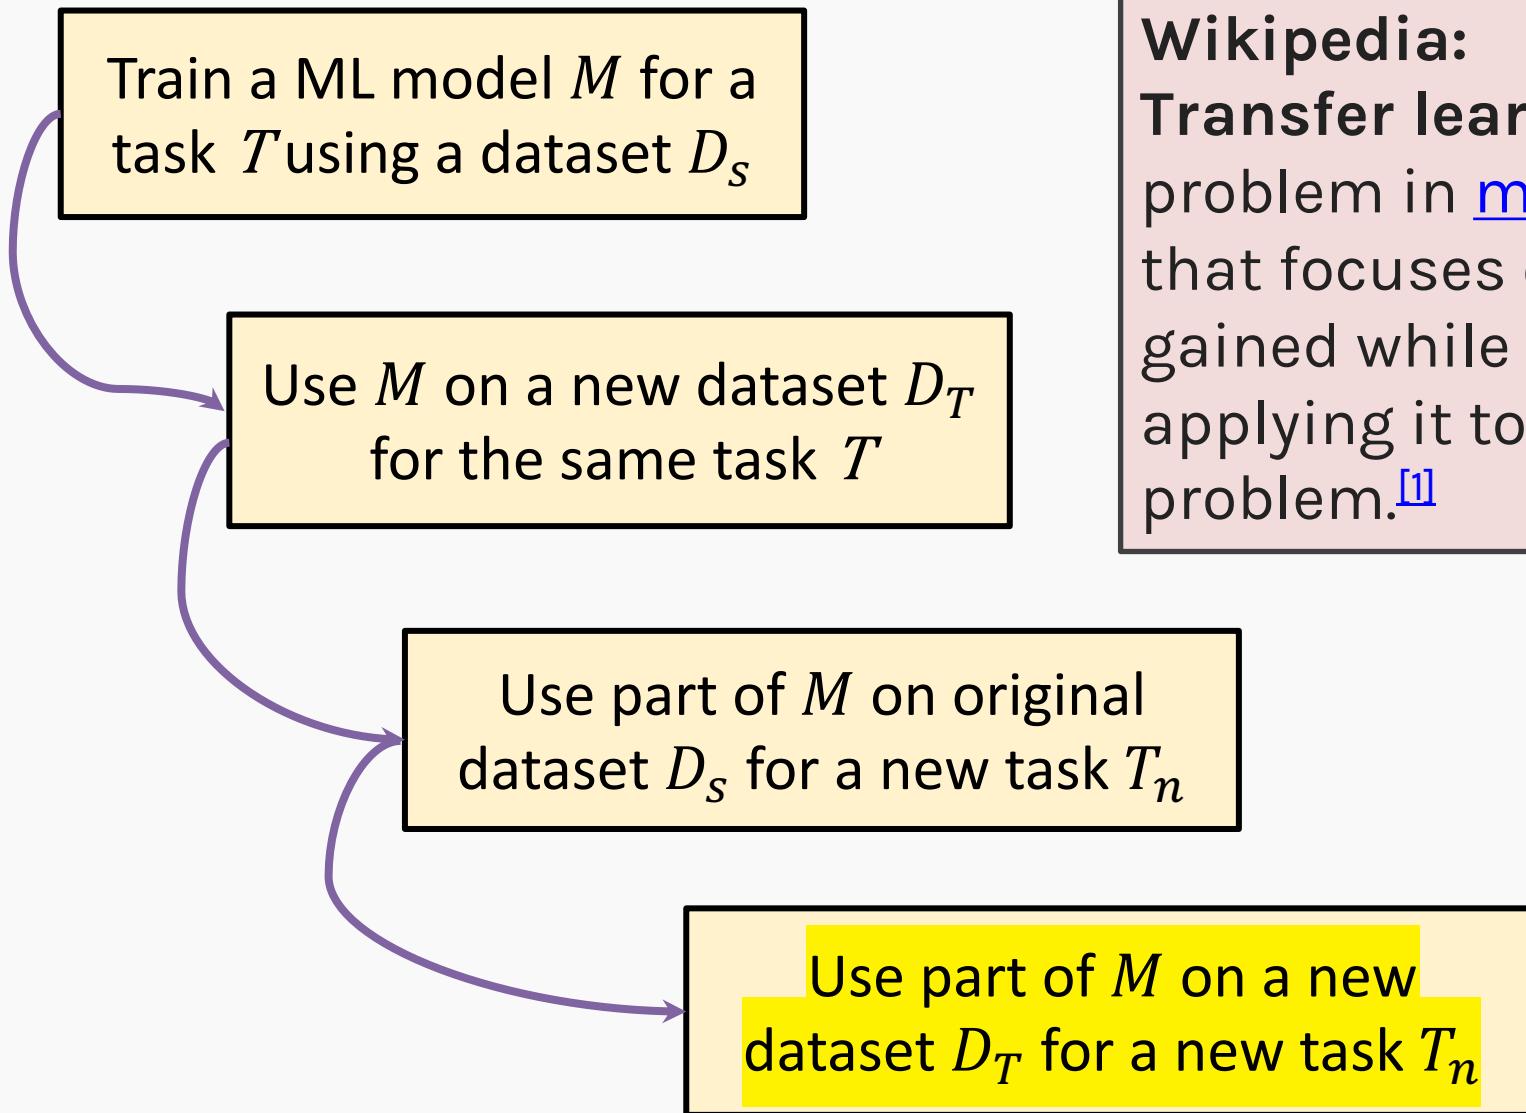
# Basic idea of Transfer Learning



## Wikipedia:

**Transfer learning (TL)** is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

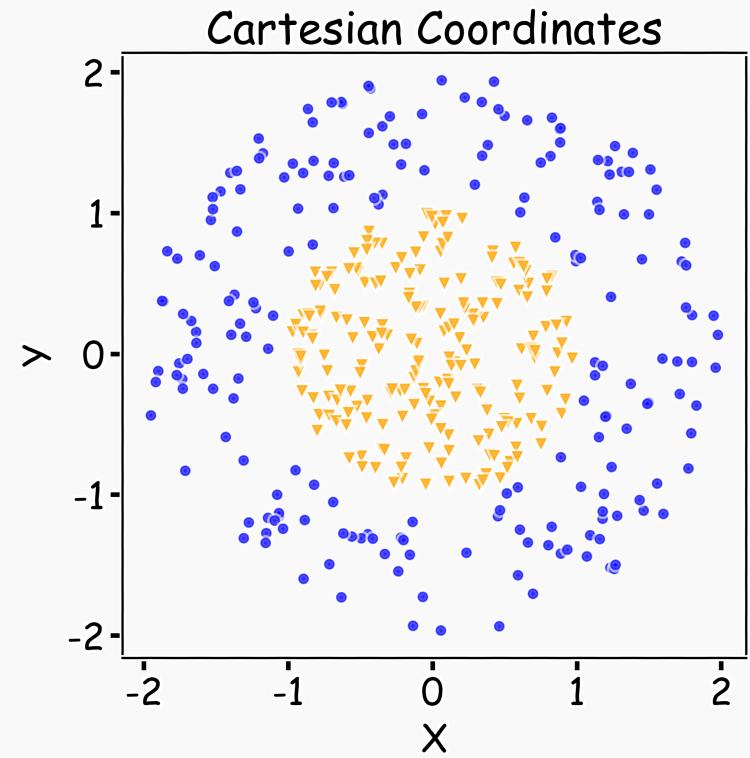
# Basic idea of Transfer Learning



## Wikipedia:

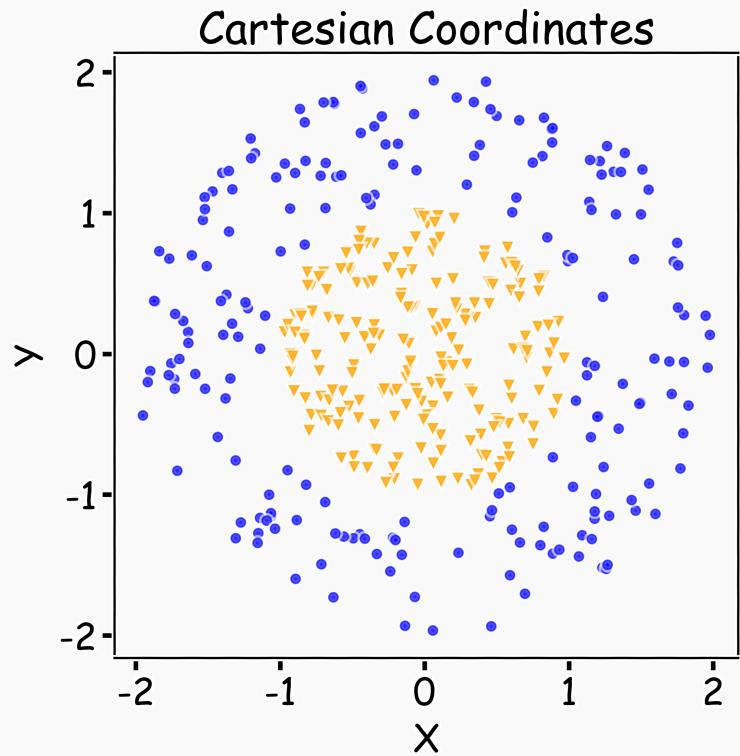
Transfer learning (TL) is a research problem in [machine learning](#) (ML) that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.<sup>[1]</sup>

# Key Idea: Representation Learning



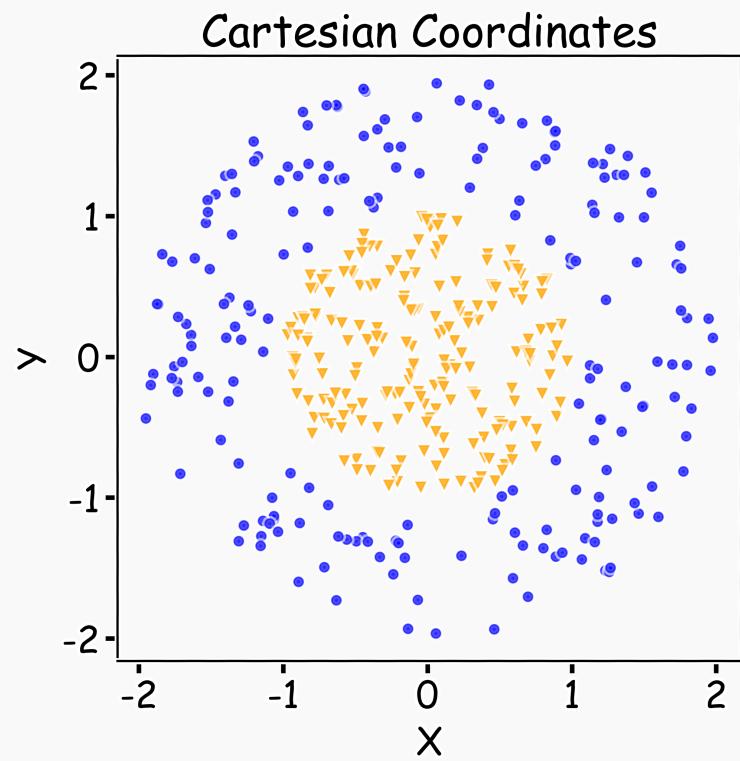
# Key Idea: Representation Learning

Relatively difficult task

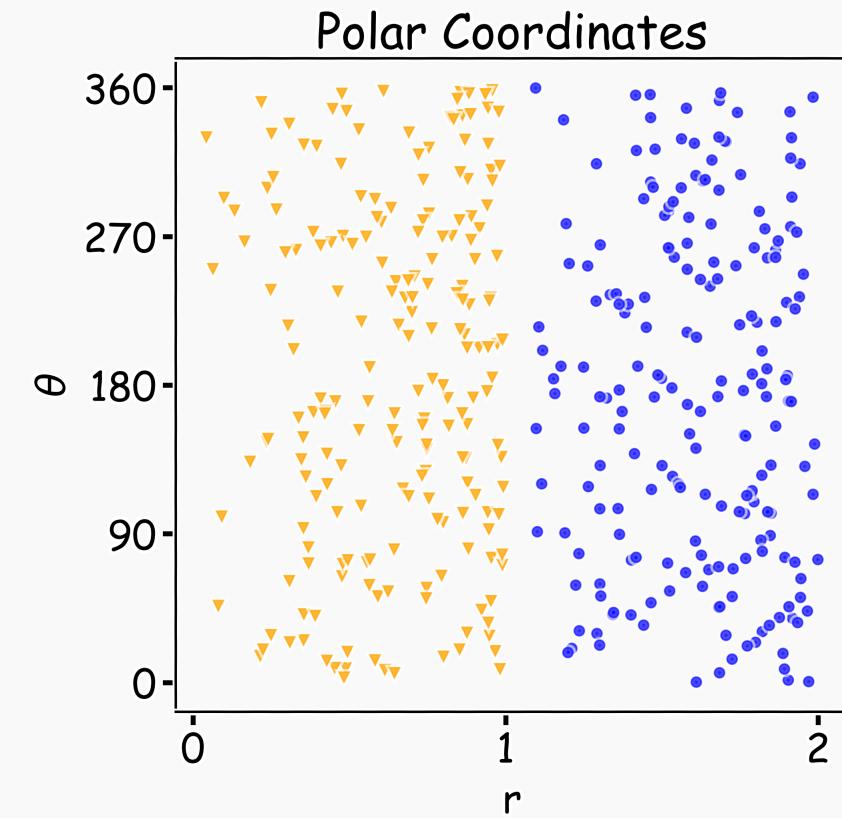


# Key Idea: Representation Learning

Relatively difficult task

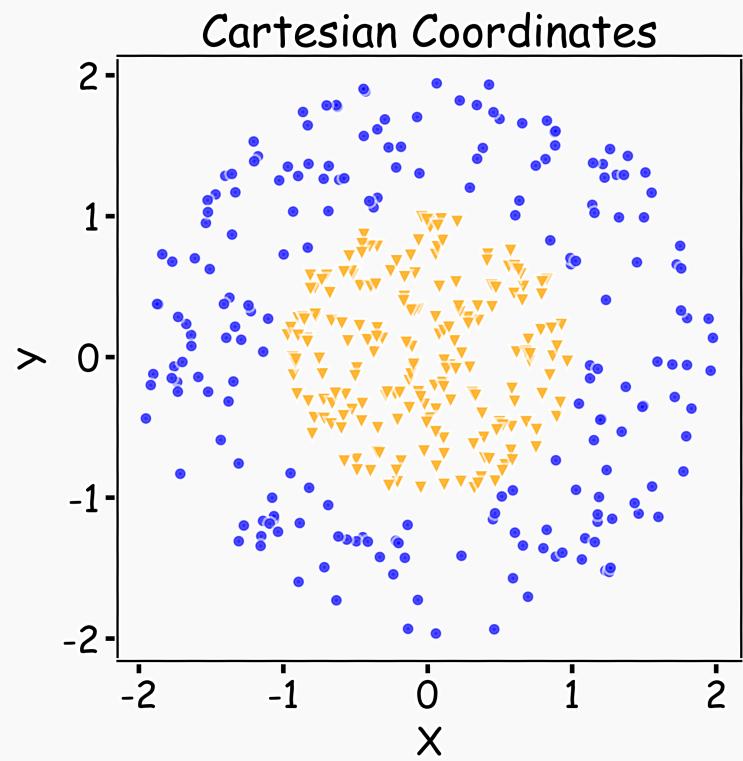


Transform:  
 $(X, Y) \rightarrow (r, \theta)$



# Key Idea: Representation Learning

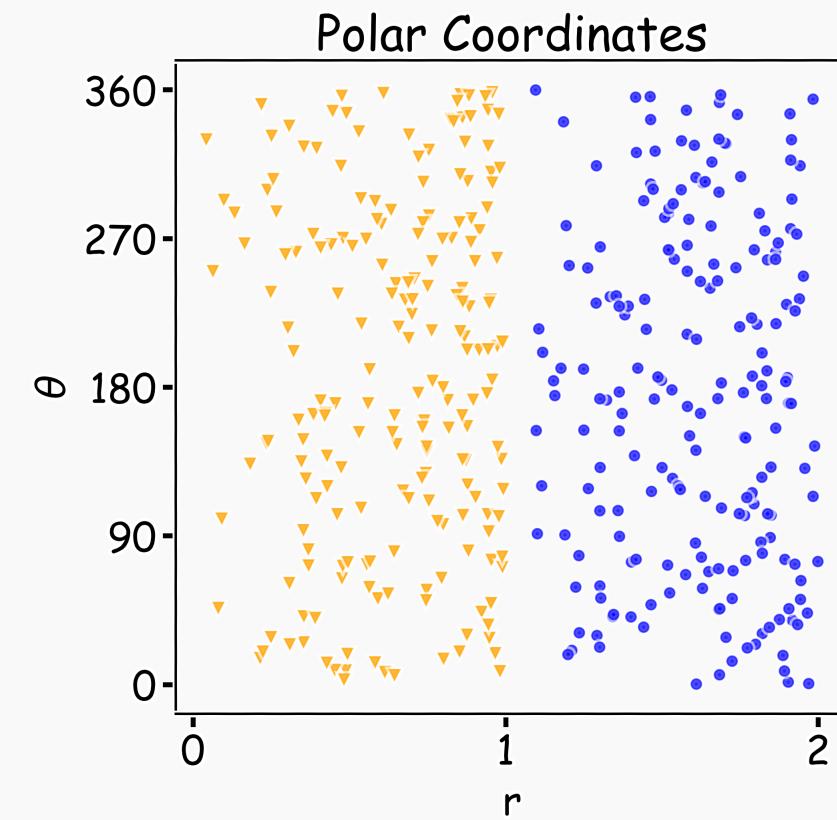
Relatively difficult task



Transform:  
 $(X, Y) \rightarrow (r, \theta)$

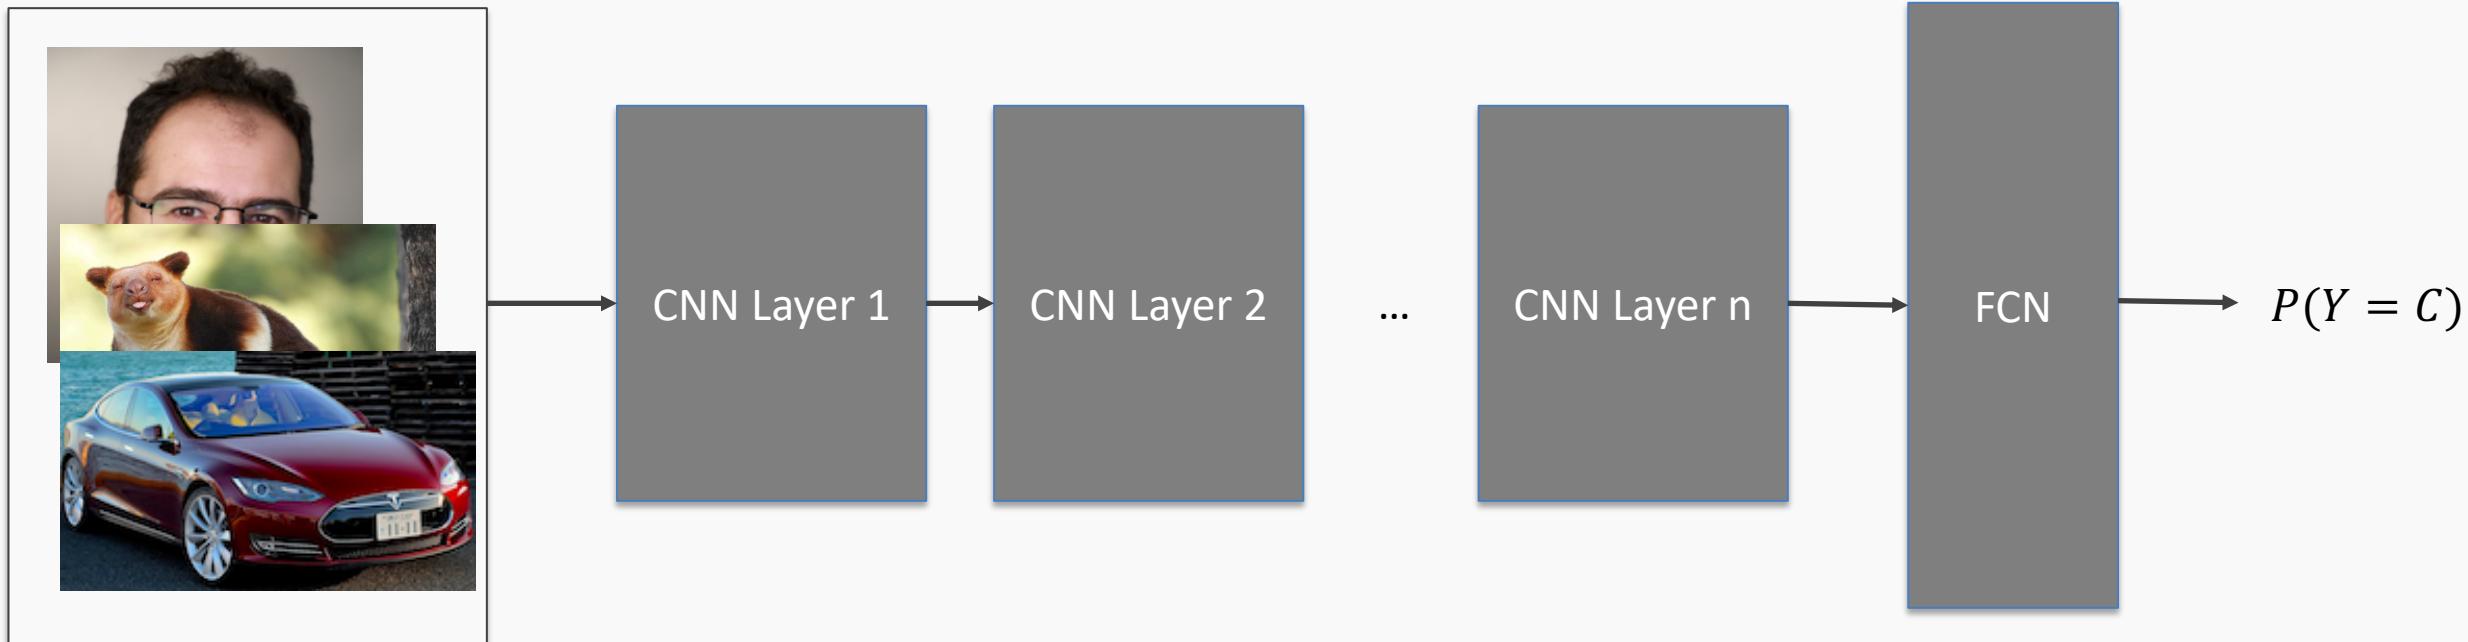


Easier task



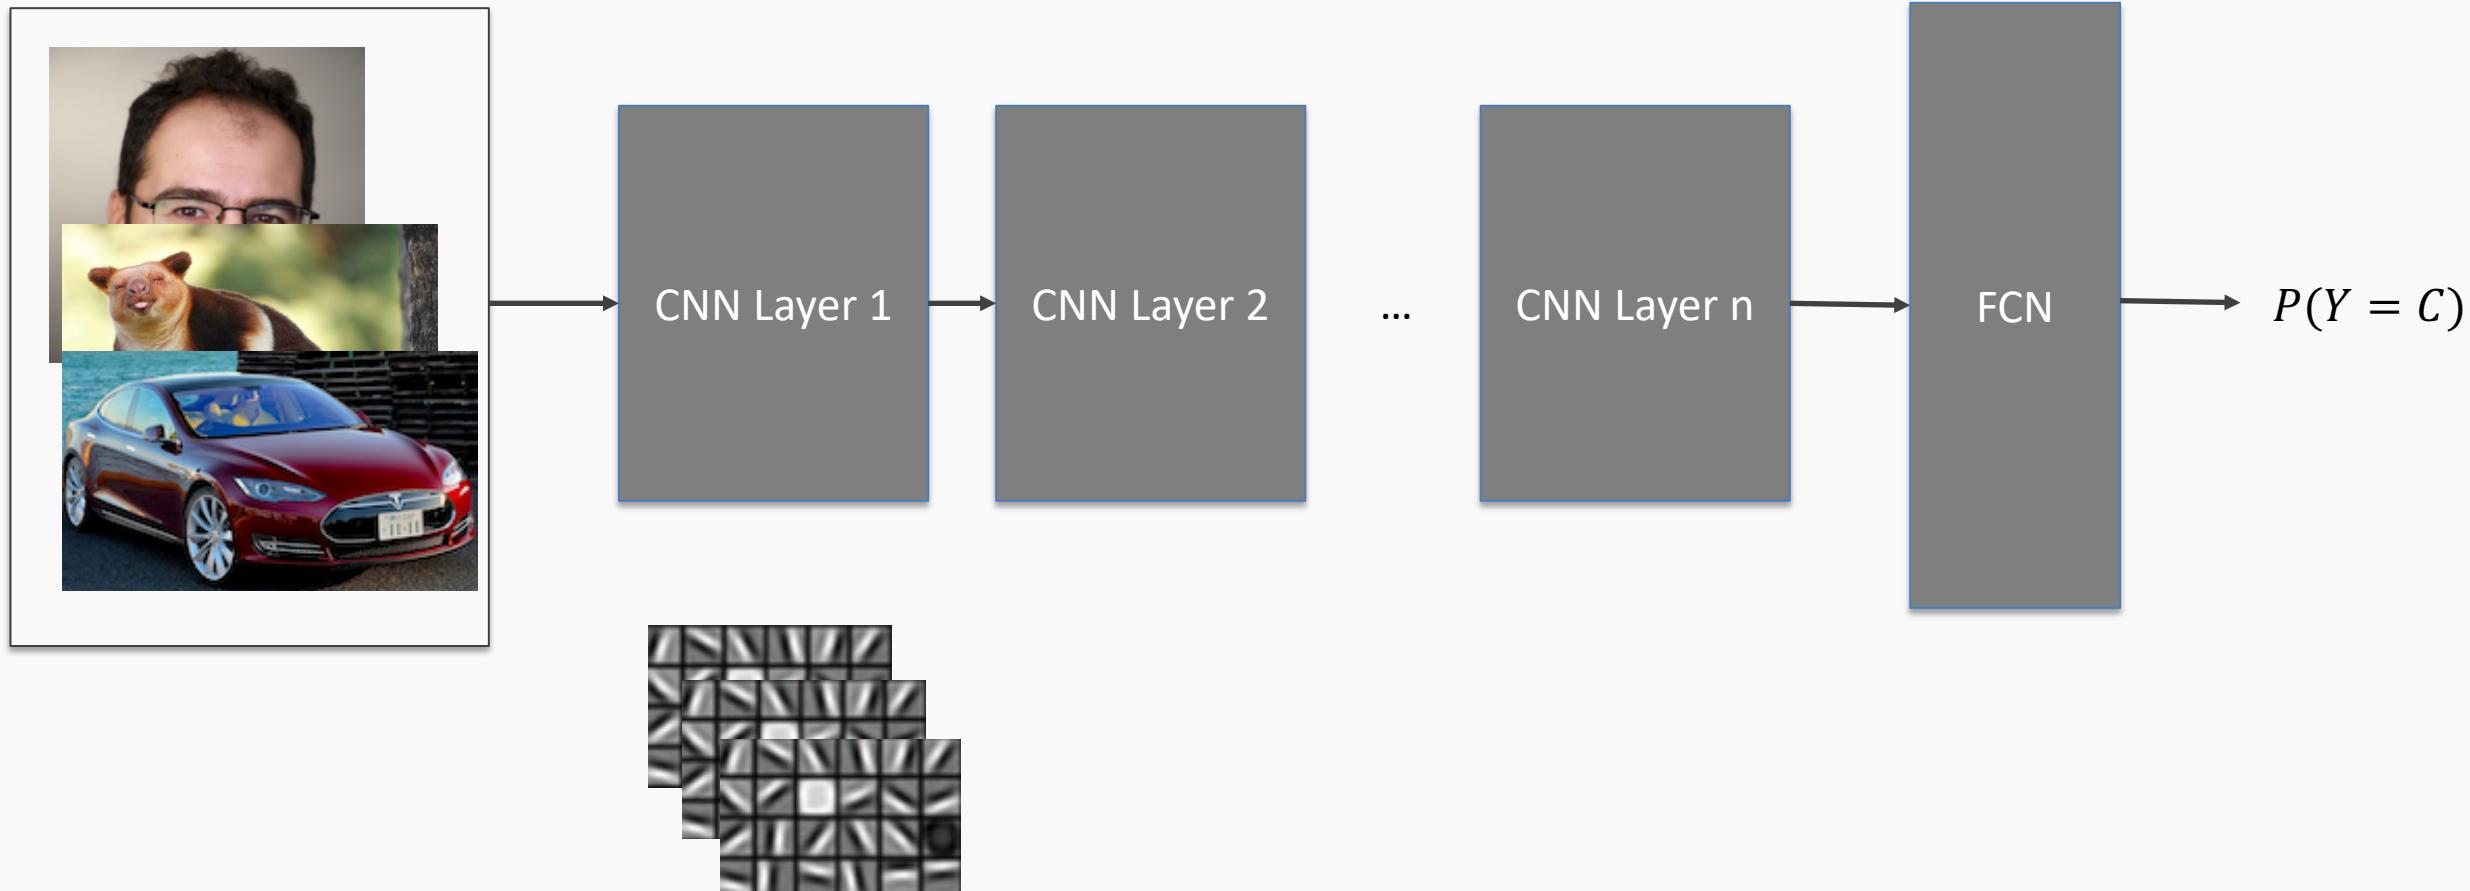
# Representation Learning

Task: classify cars, people, animals and objects



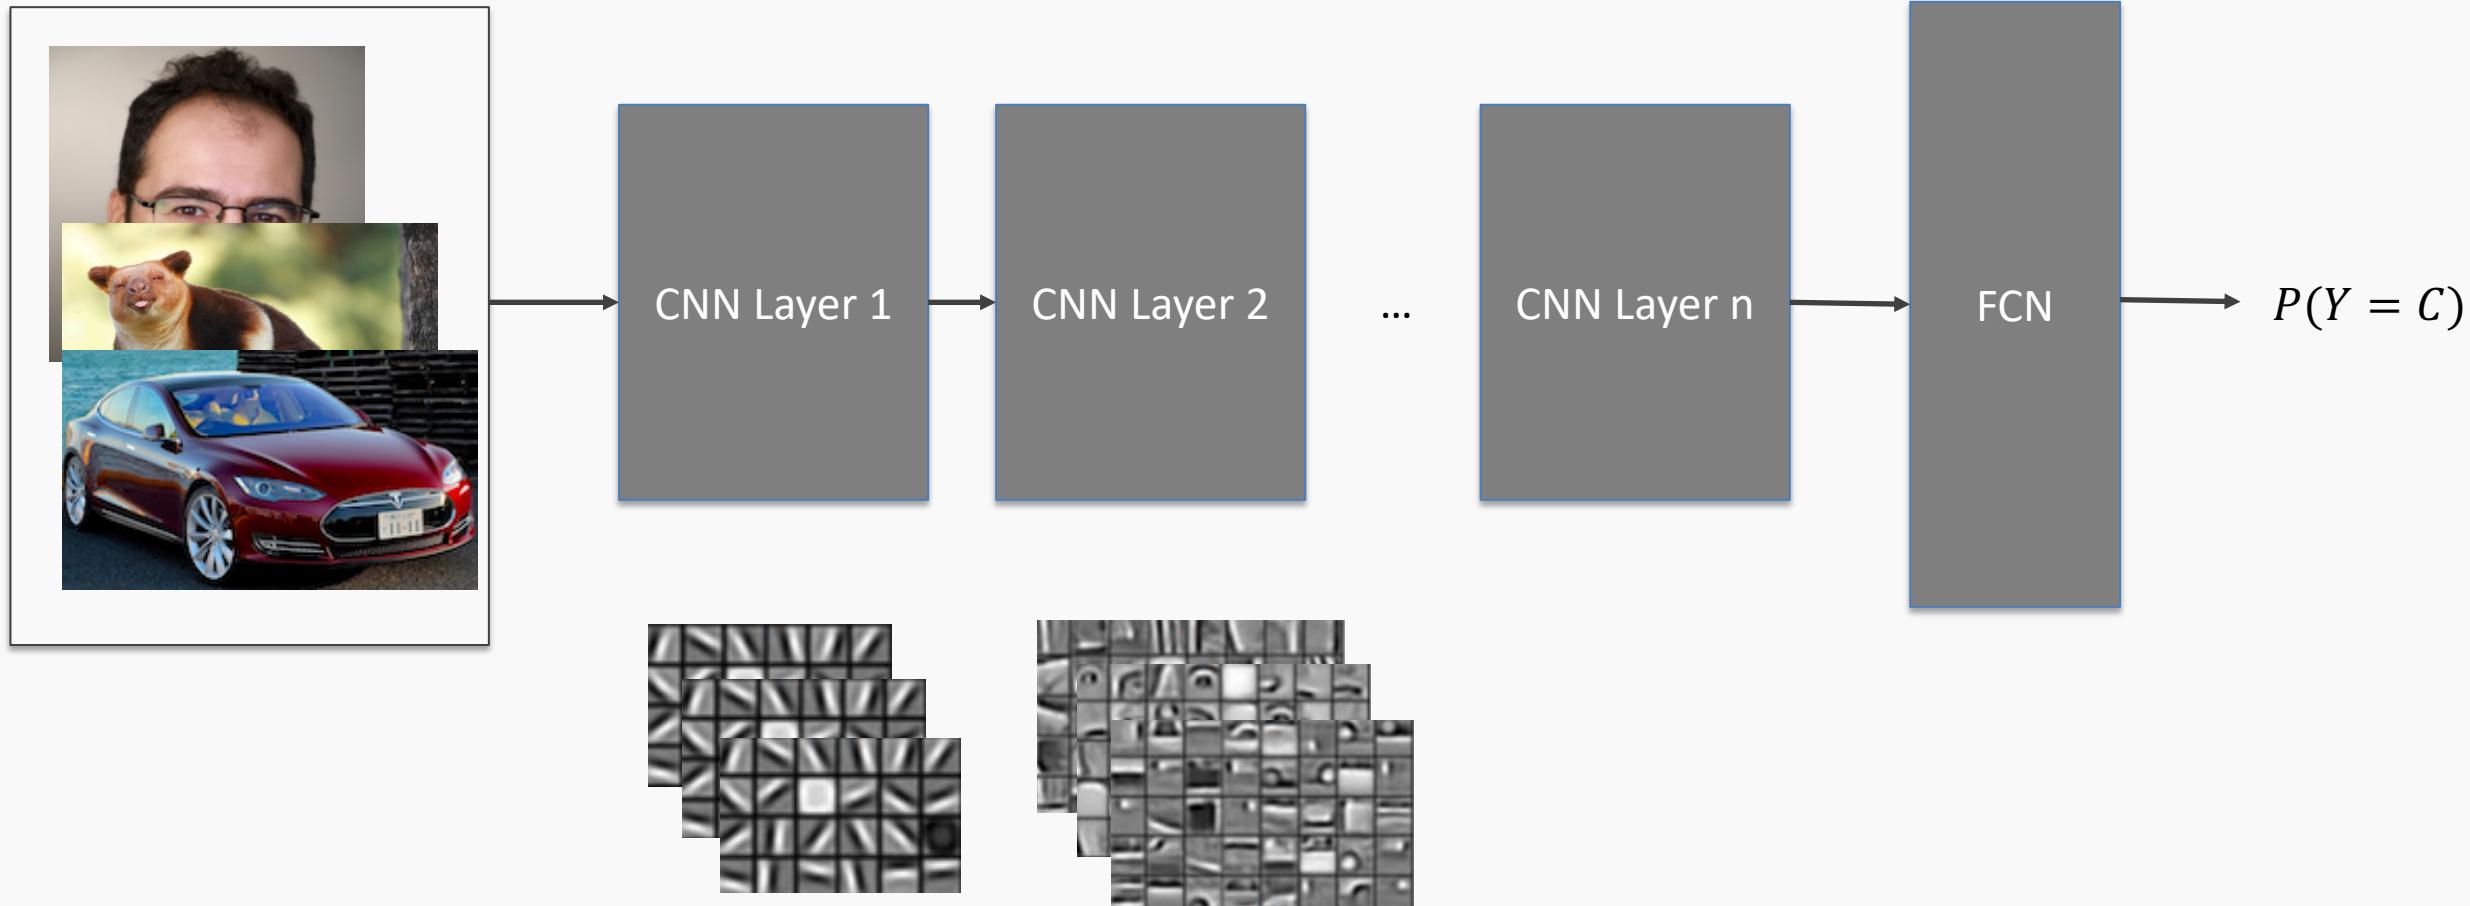
# Representation Learning

Task: classify cars, people, animals and objects



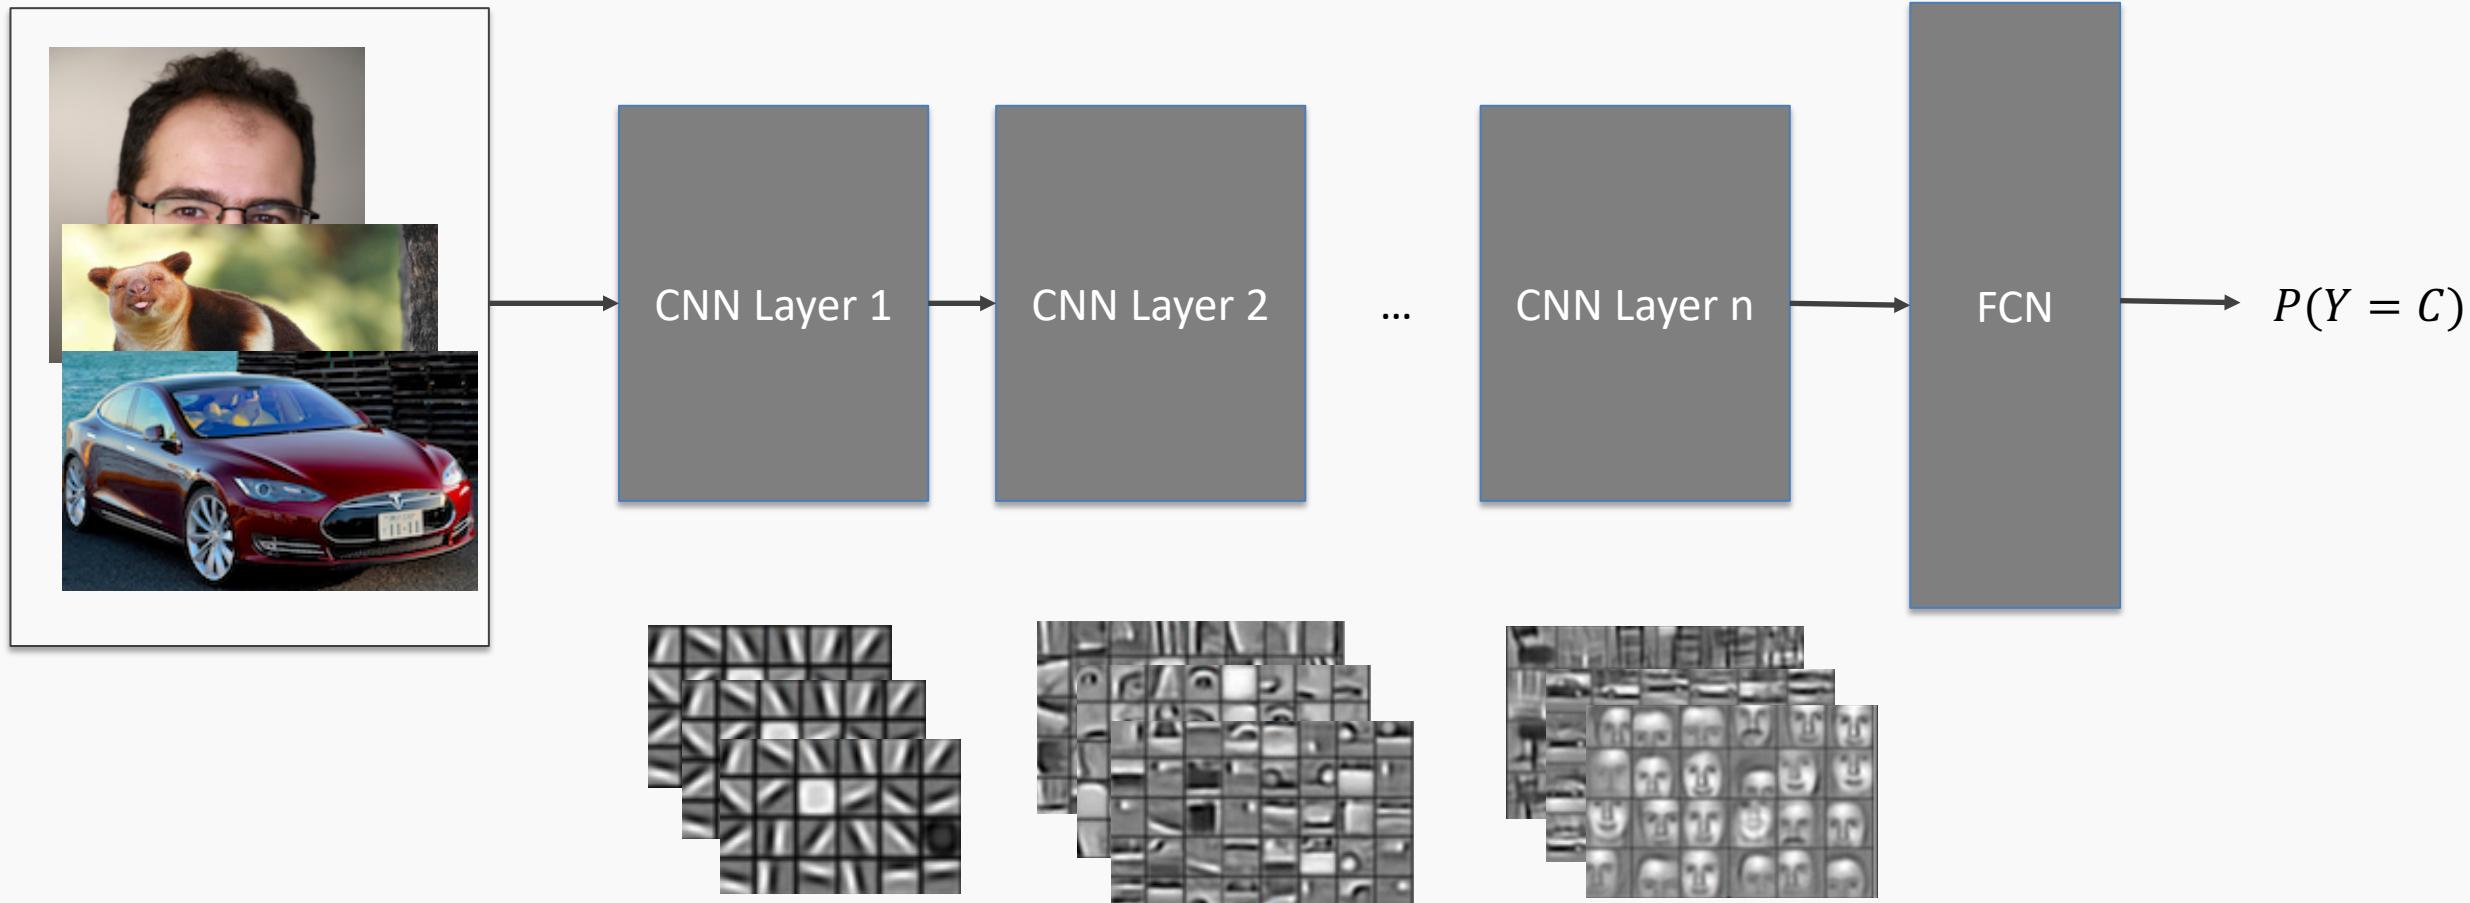
# Representation Learning

Task: classify cars, people, animals and objects



# Representation Learning

Task: classify cars, people, animals and objects



# Transfer Learning To The Rescue

---

How do you make an image classifier that can be trained in a few minutes on a CPU with very little data?



# Transfer Learning To The Rescue

---

How do you make an image classifier that can be trained in a few minutes on a CPU with very little data?

***Use pre-trained models***, i.e., models with known weights.



# Transfer Learning To The Rescue

How do you make an image classifier that can be trained in a few minutes on a CPU with very little data?

**Use pre-trained models**, i.e., models with known weights.

**Main Idea:** earlier layers of a network learn low level features, which can be adapted to new domains by changing weights at later and fully-connected layers.

# Transfer Learning To The Rescue

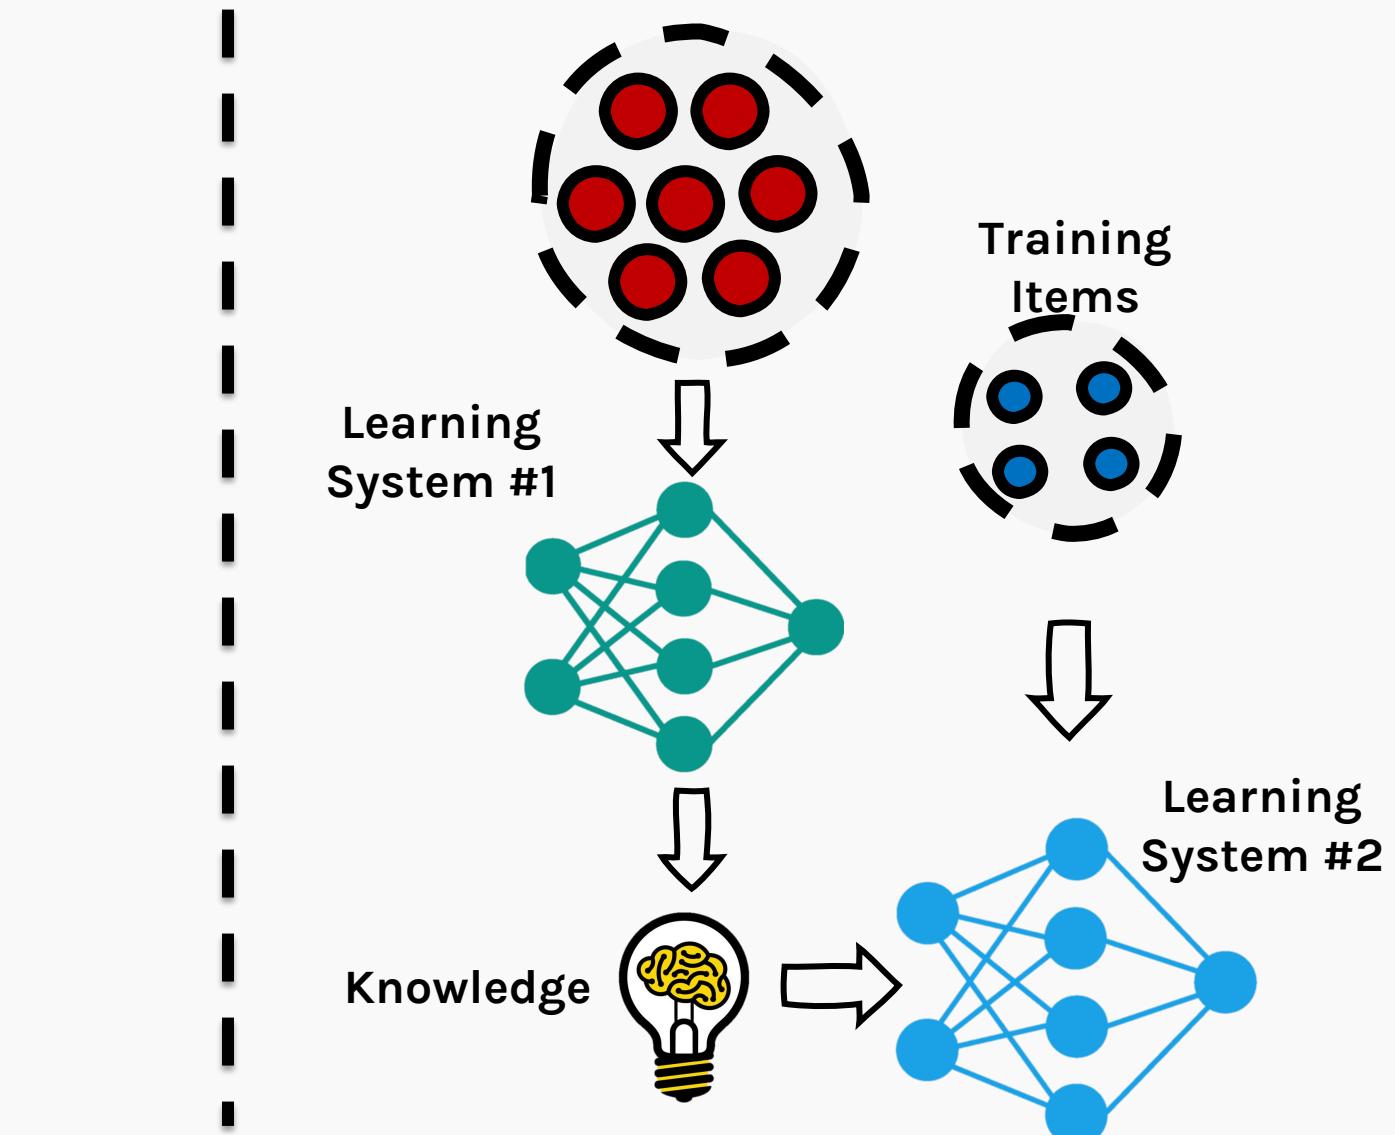
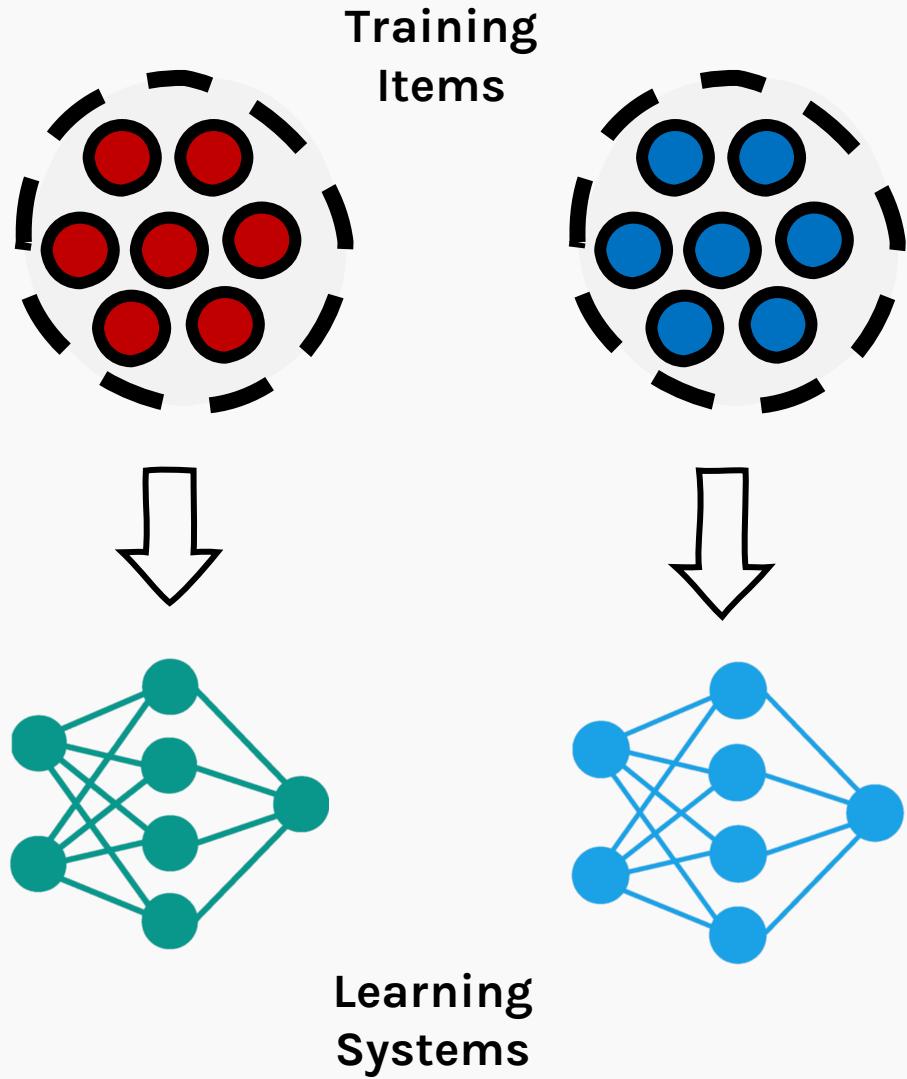
How do you make an image classifier that can be trained in a few minutes on a CPU with very little data?

**Use pre-trained models**, i.e., models with known weights.

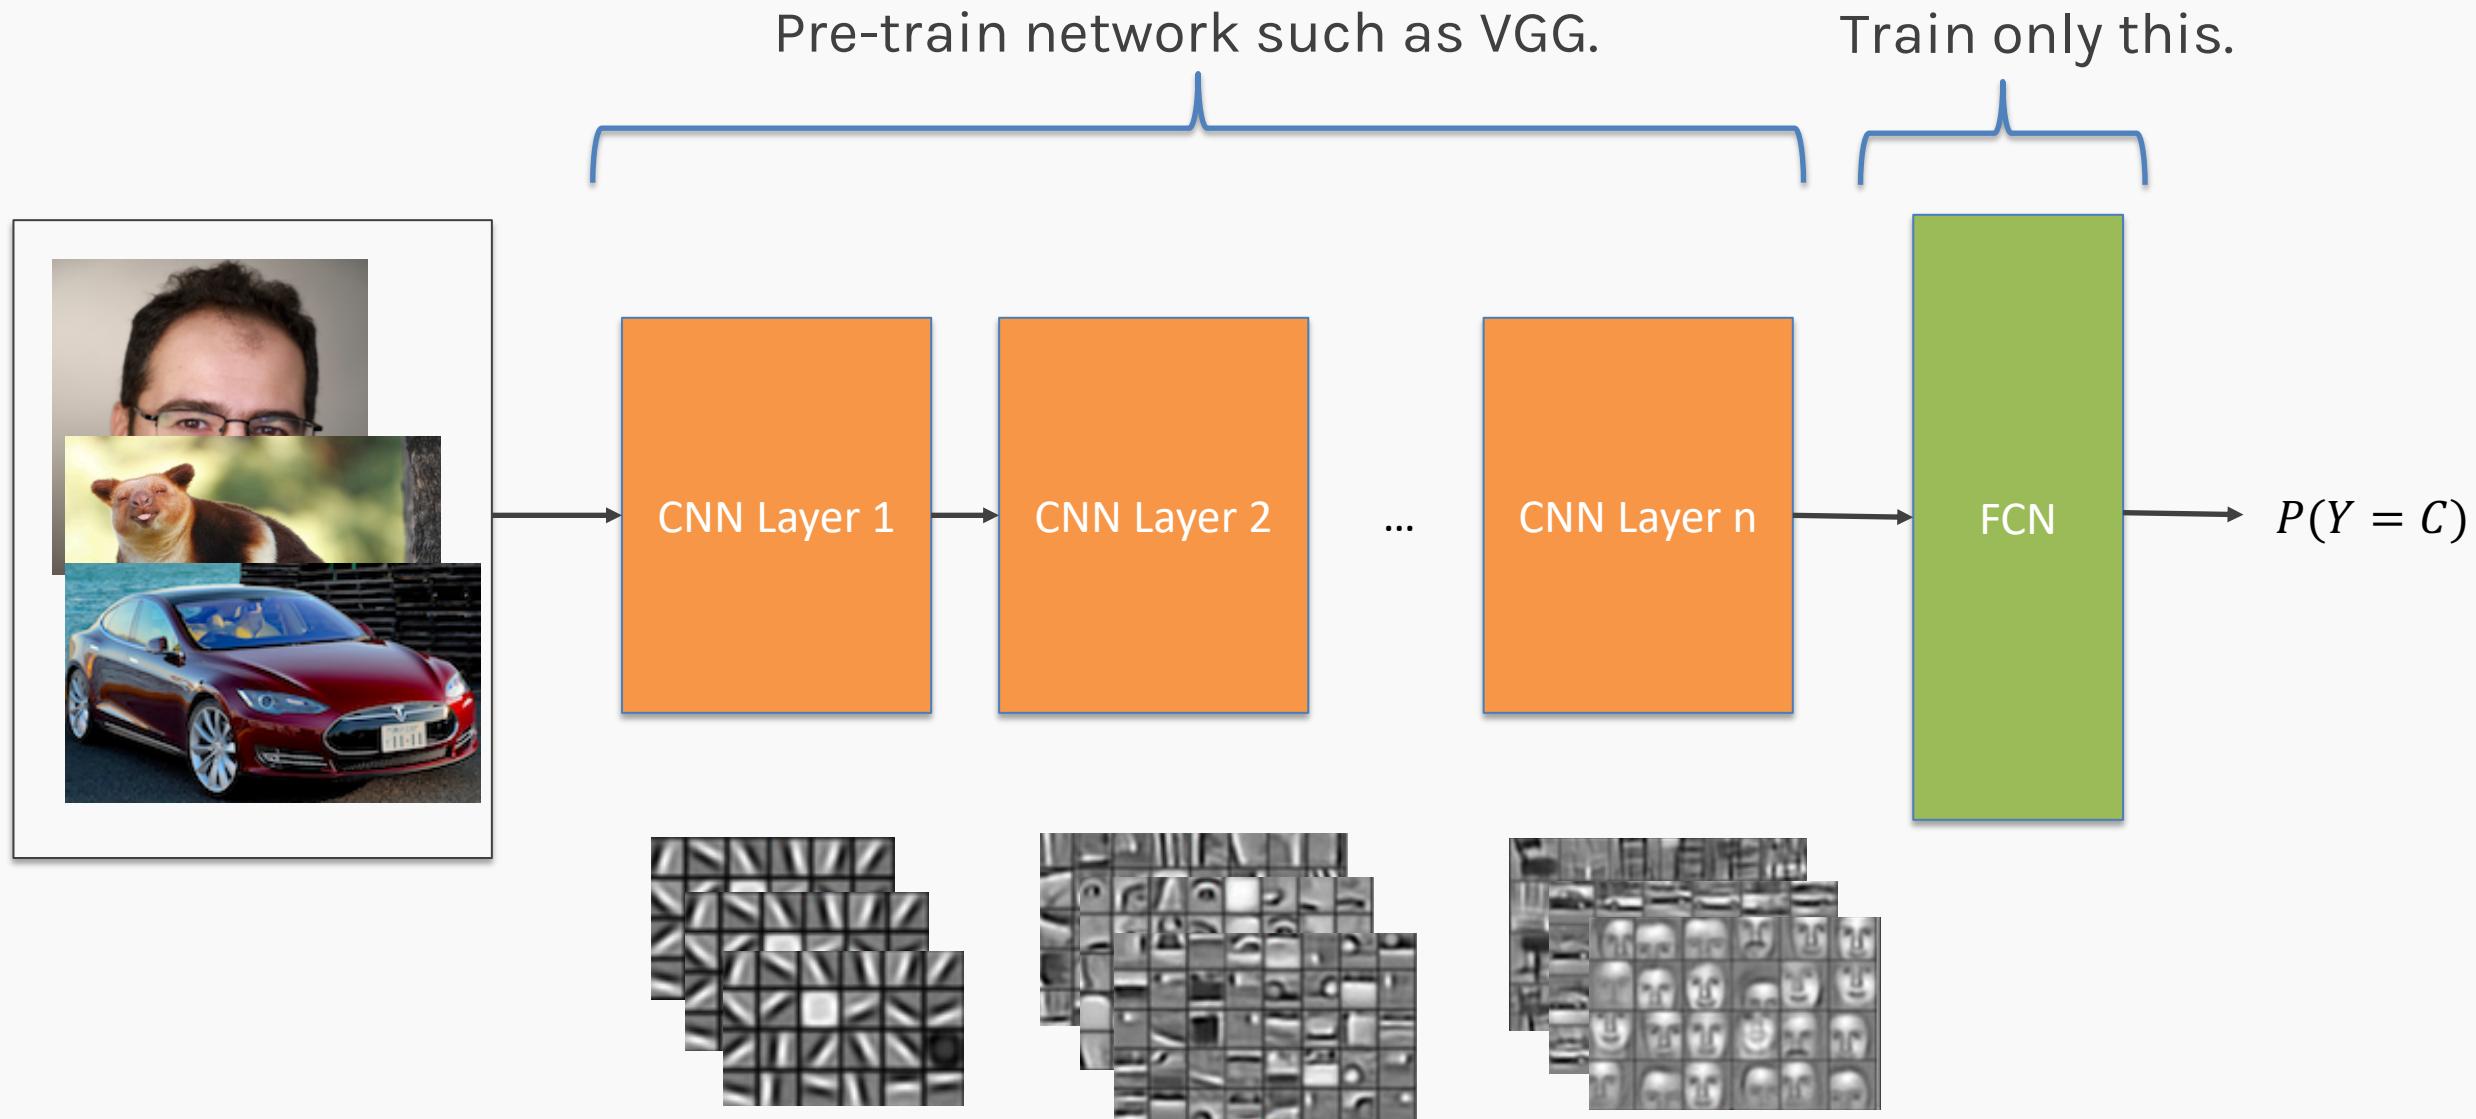
**Main Idea:** earlier layers of a network learn low level features, which can be adapted to new domains by changing weights at later and fully-connected layers.

**Example:** use ImageNet trained with any sophisticated huge network. Then retrain it on a few images.

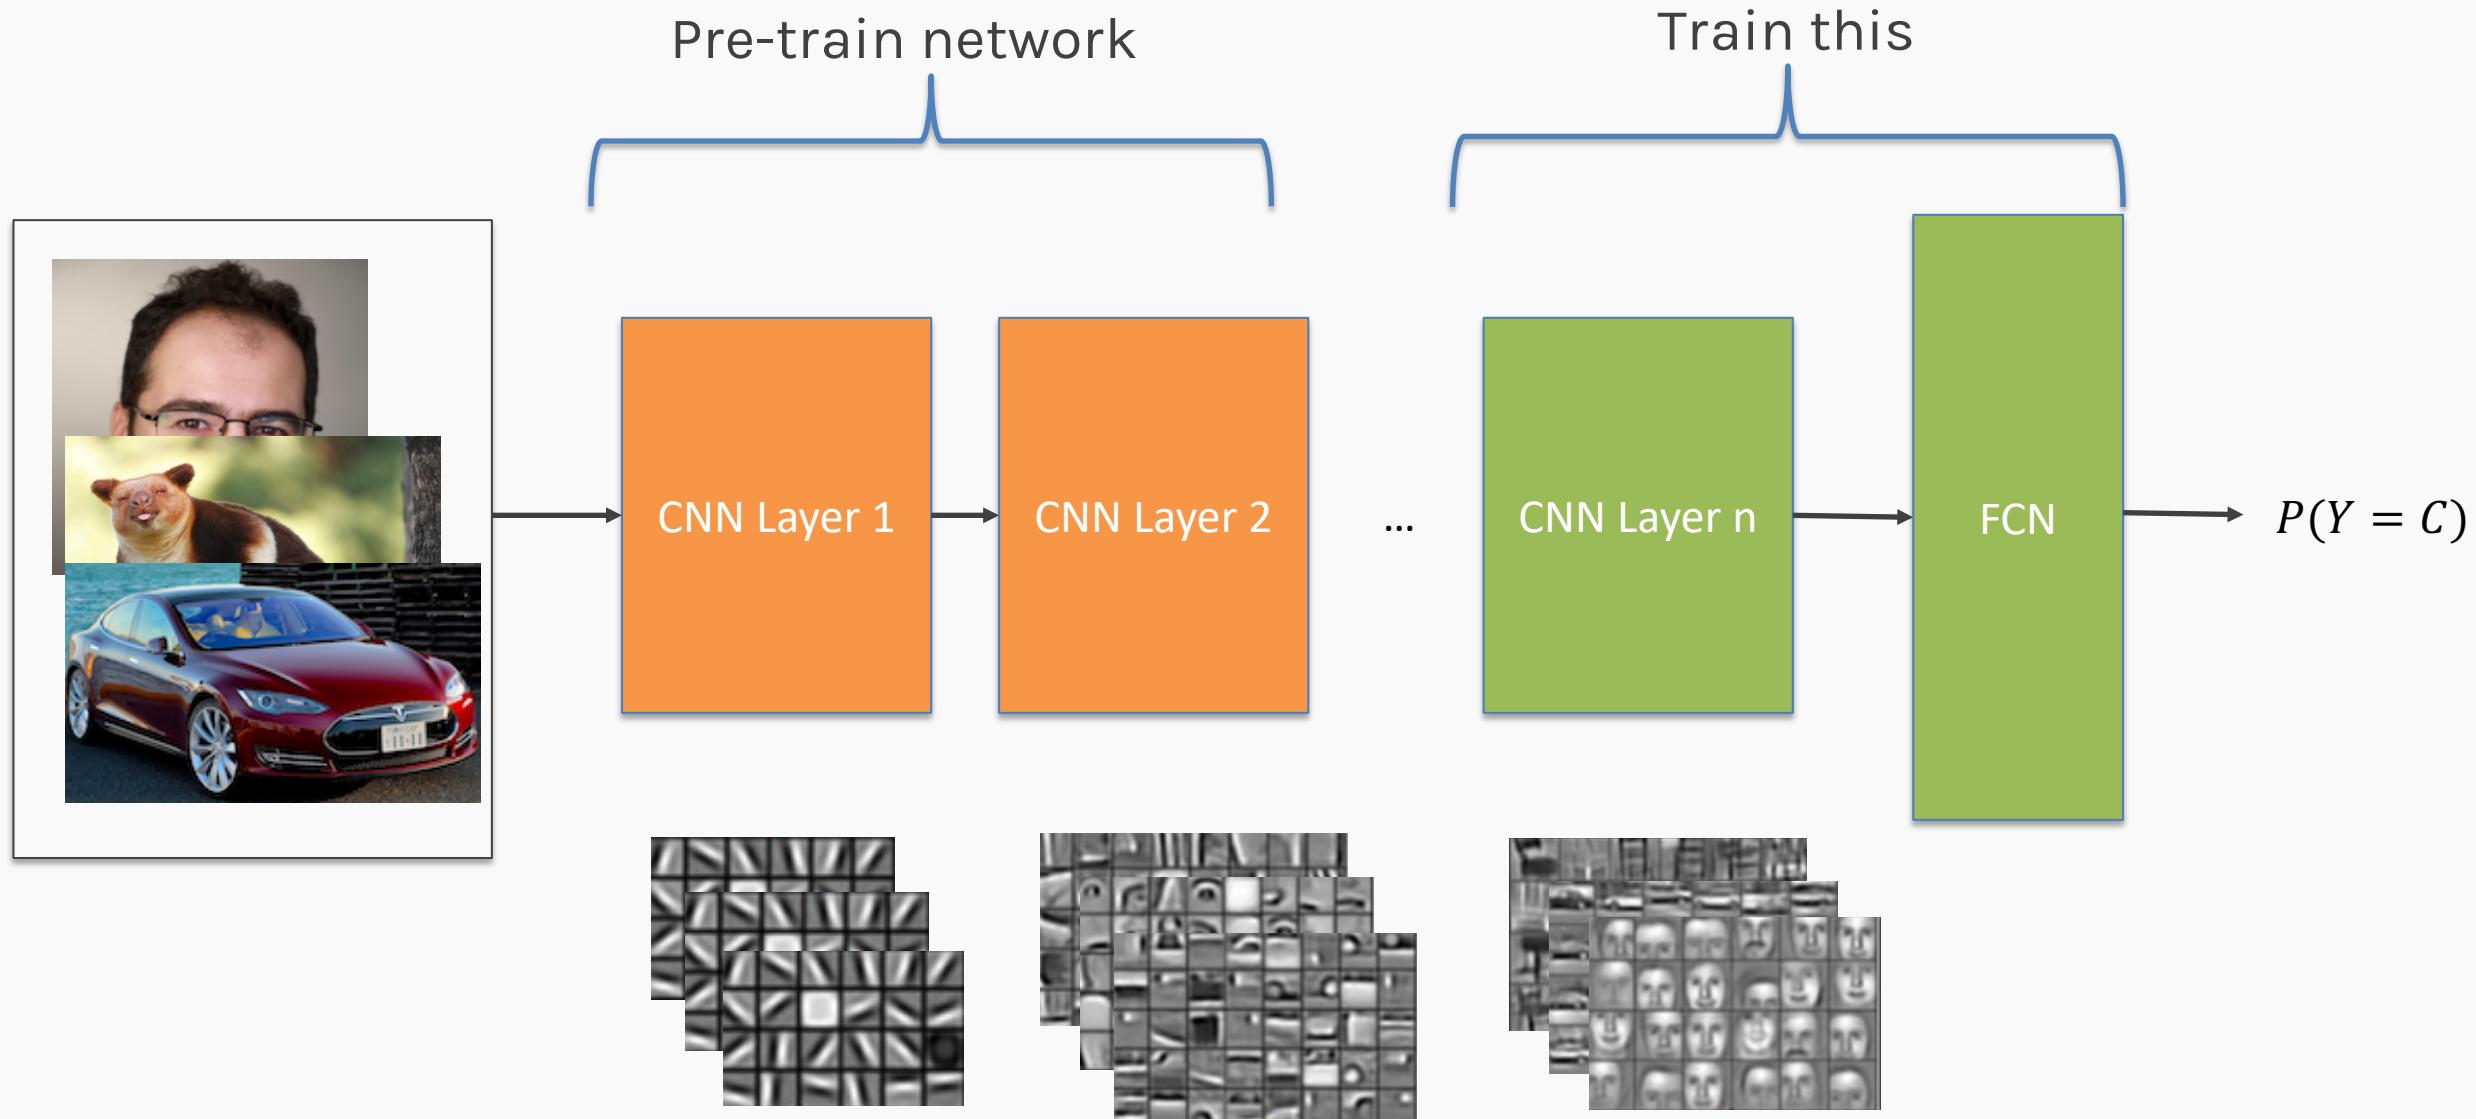
# Traditional Machine Learning vs Transfer Learning



# Transfer Learning: Only train dense layers

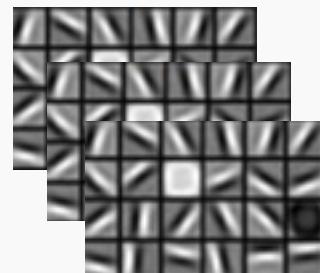
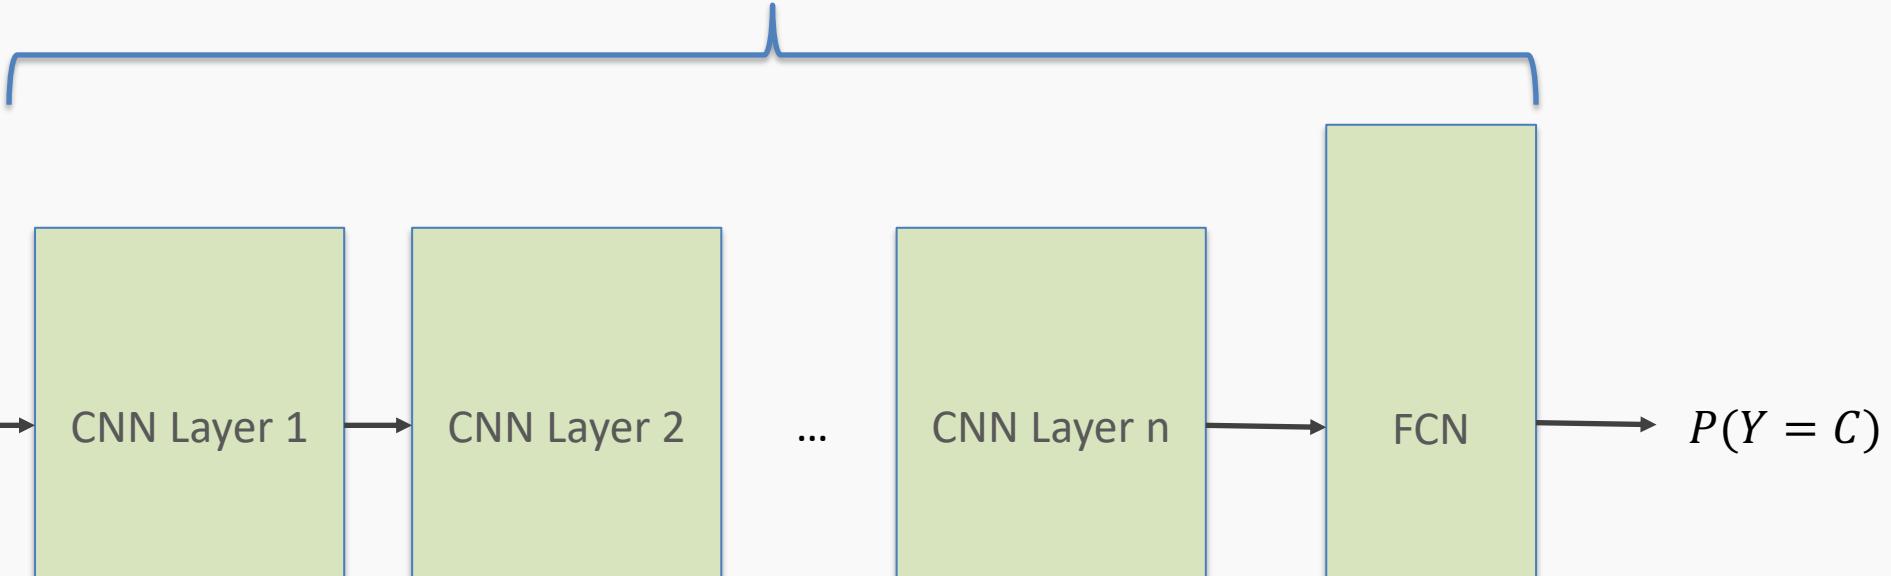
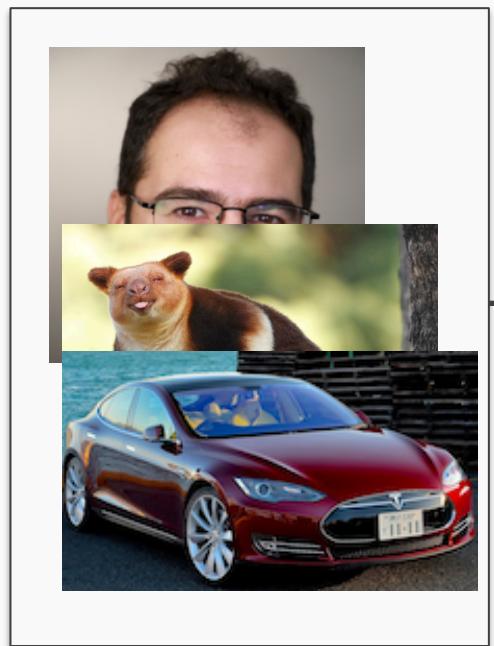


# Transfer Learning: Train last few Conv layers and dense layers



# Transfer Learning: **Fine tuning**

Train everything but start with weights that are trained already



## Transfer Learning for Deep Learning

### What people think:

- You can't do deep learning unless you have a million labeled examples.

### What people do:

- You can train on a nearby objective for which is easy to generate labels (ImageNet).
- You can transfer learned representations from a relate task.
- You can learn representations from **unlabeled** data.

## Transfer Learning for Deep Learning

### What people think:

- You can't do deep learning unless you have a million

Indeed, you can train a network for classification and use it for segmentation.

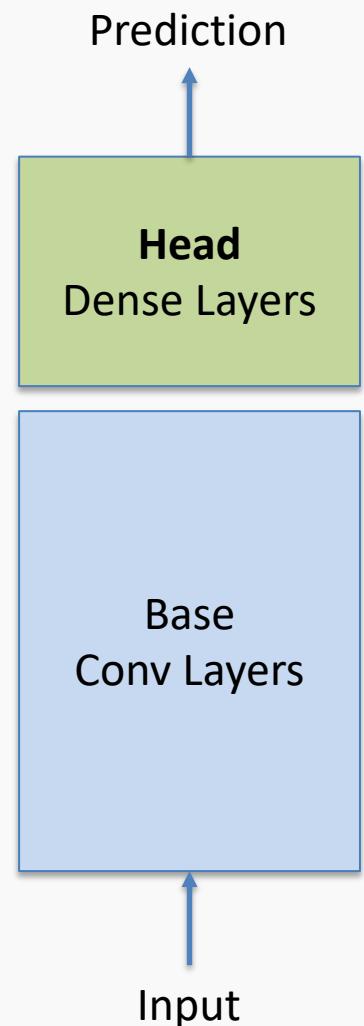
### What people do:

- You can train on a nearby objective for which is easy to generate labels (ImageNet).
- You can transfer learned representations from a related task.
- You can learn representations from unlabeled data.

We will revisit this idea when we will talk about autoencoders and language model.

# Transfer learning - Feature-Representation Extraction

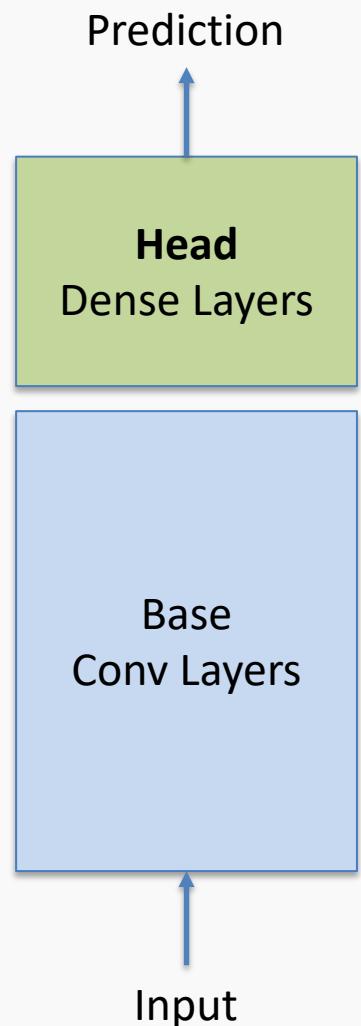
Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.



# Transfer learning – Feature-Representation Extraction

Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.

**Example:** Classify cats and dogs from your own images:



# Transfer learning – Feature-Representation Extraction

Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.

**Example:** Classify cats and dogs from your own images:

- Keep (frozen) convolutional **base** from big model, since convolutional base is “generic”.



# Transfer learning – Feature-Representation Extraction

Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.

**Example:** Classify cats and dogs from your own images:

- Keep (frozen) convolutional **base** from big model, since convolutional base is “generic”.
- Throw away **the fully connected layers (head)** since **these layers have no notion of space.**



# Transfer learning – Feature-Representation Extraction

Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.

**Example:** Classify cats and dogs from your own images:

- Keep (frozen) convolutional **base** from big model, since convolutional base is “generic”.
- Throw away **the fully connected layers (head)** since these layers have no notion of space.
- Introduce new dense layers. It can be new architecture or the same as the original model.



# Transfer learning – Feature-Representation Extraction

Use representations learned by big net to extract features from new samples, which are then fed to a new classifier.

**Example:** Classify cats and dogs from your own images:

- Keep (frozen) convolutional **base** from big model, since convolutional base is “generic”.
- Throw away **the fully connected layers (head)** since these layers have no notion of space.
- Introduce new dense layers. It can be new architecture or the same as the original model.

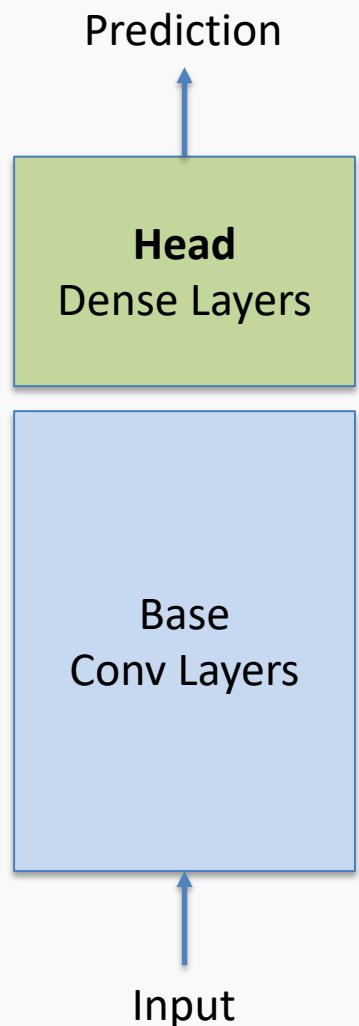
**Note:** Classify cats and dogs from your own images. Since there are both dogs and cats in ImageNet you could get **away** with using the head **FC** layers as well (instance TL). But by throwing it away you can learn more from other dog/cat images.



# Transfer learning – Fine-tuning

Up to now we have frozen the entire convolutional base.

Remember that earlier layers learn highly generic feature maps  
(edges, colors, textures).

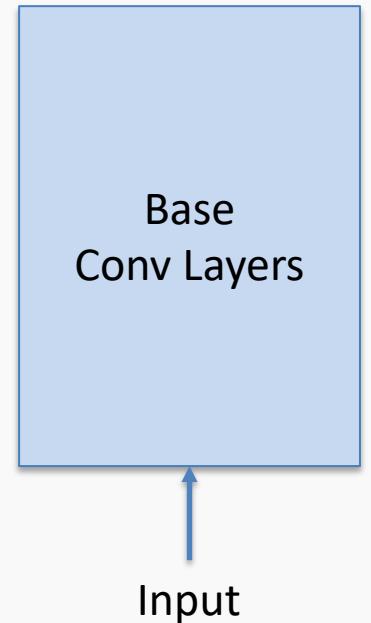


# Transfer learning – Fine-tuning

Up to now we have frozen the entire convolutional base.

Remember that earlier layers learn highly generic feature maps (edges, colors, textures).

- Throw away **the fully connected layers (head)** since these layers have no notion of space.

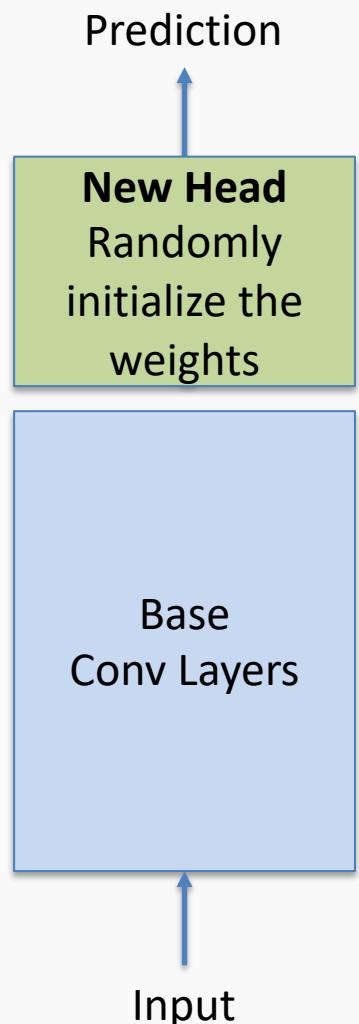


# Transfer learning – Fine-tuning

Up to now we have frozen the entire convolutional base.

Remember that earlier layers learn highly generic feature maps (edges, colors, textures).

- Throw away **the fully connected layers (head)** since these layers have no notion of space.
- Introduce new dense layers. It can be new architecture or the same as the original model. Train the new head while keeping the base fixed.



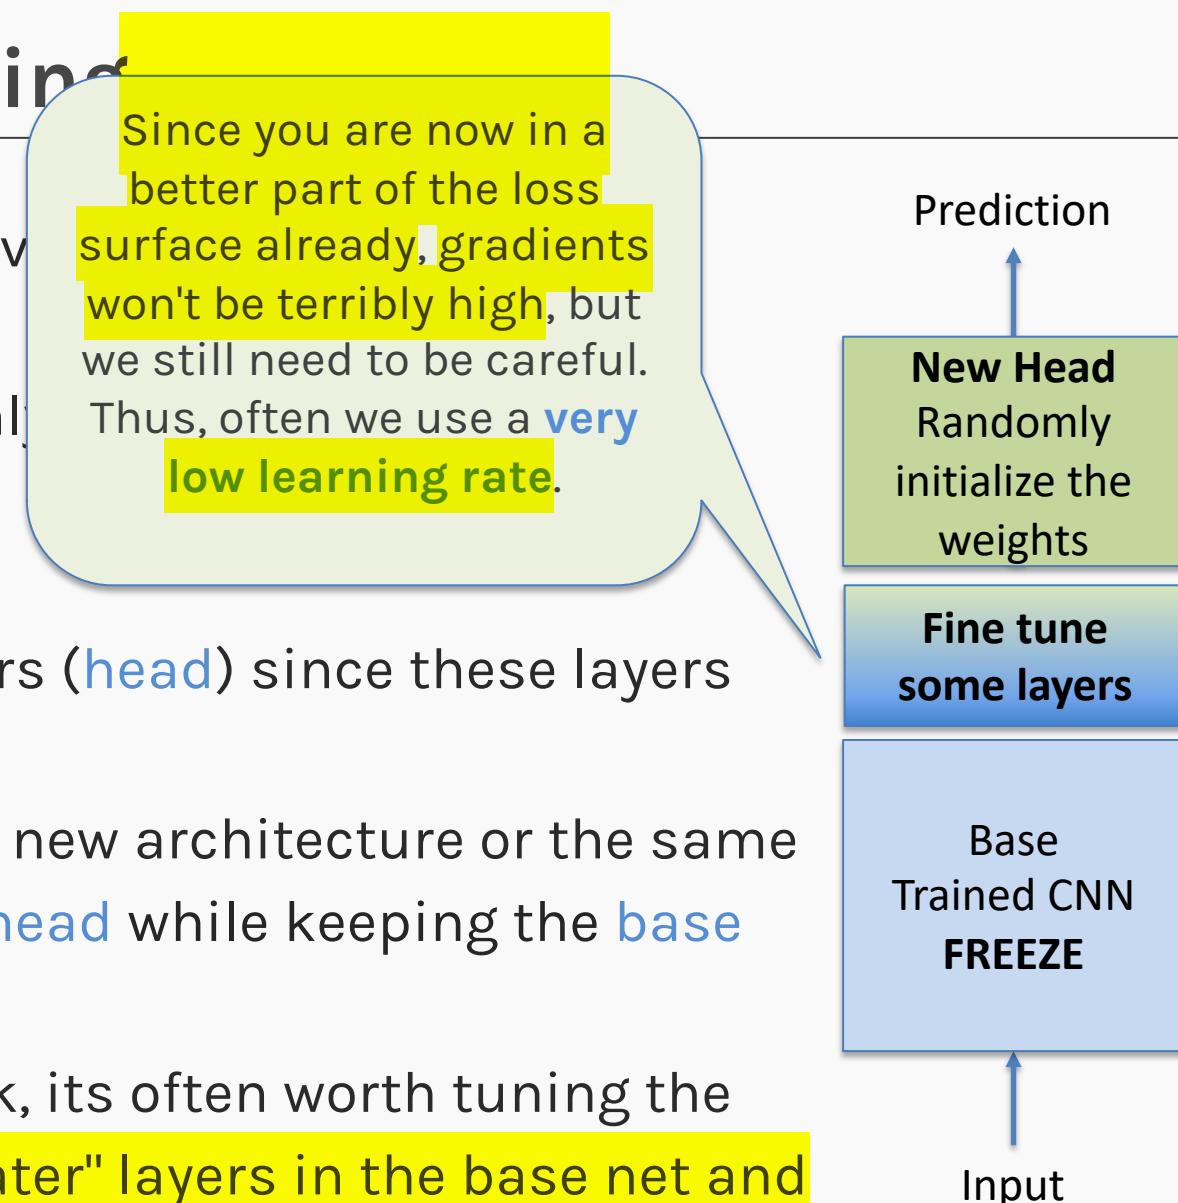
# Transfer learning – Fine-tuning

Up to now we have frozen the entire convolutional network.

Remember that earlier layers learn highly abstract features (edges, colors, textures).

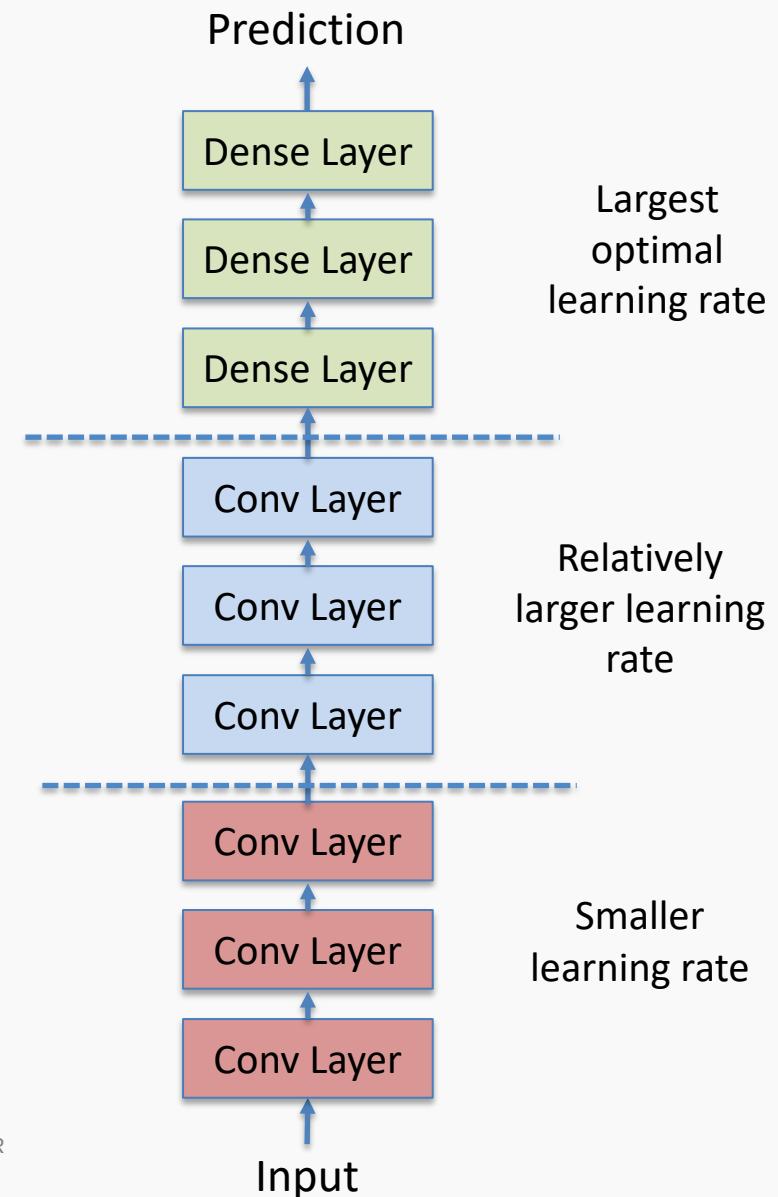
- Throw away **the fully connected layers (head)** since these layers have no notion of space.
- Introduce new dense layers. It can be new architecture or the same as the original model. **Train** the new **head** while keeping the **base fixed**.
- To particularize the model to our task, it's often worth tuning the later layers as well. **Unfreeze** some "later" layers in the base net and now train the base net and FC net together.

**Note:** But we must be very careful not to have big gradient updates.



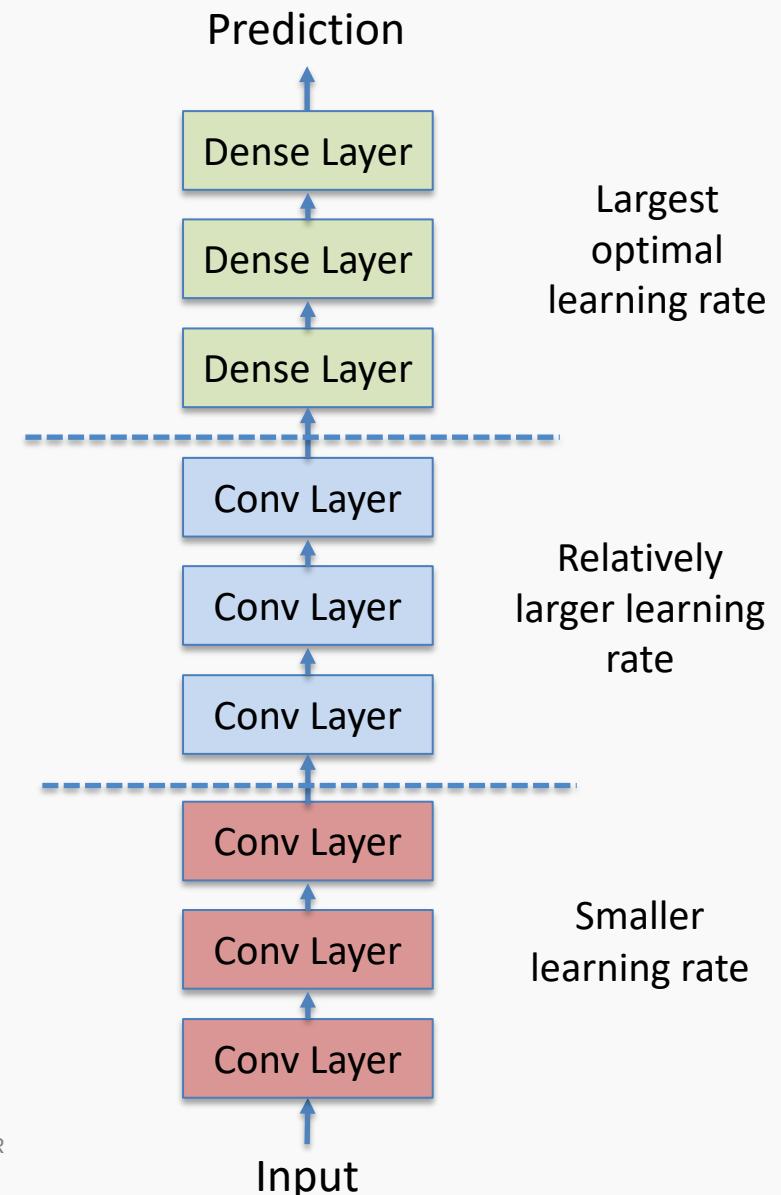
# Transfer Learning for Deep Learning: Differential Learning Rates

- A low learning rate can take a lot of time to train on the "later" layers. Since we trained the FC head earlier, we could probably retrain them at a higher learning rate.



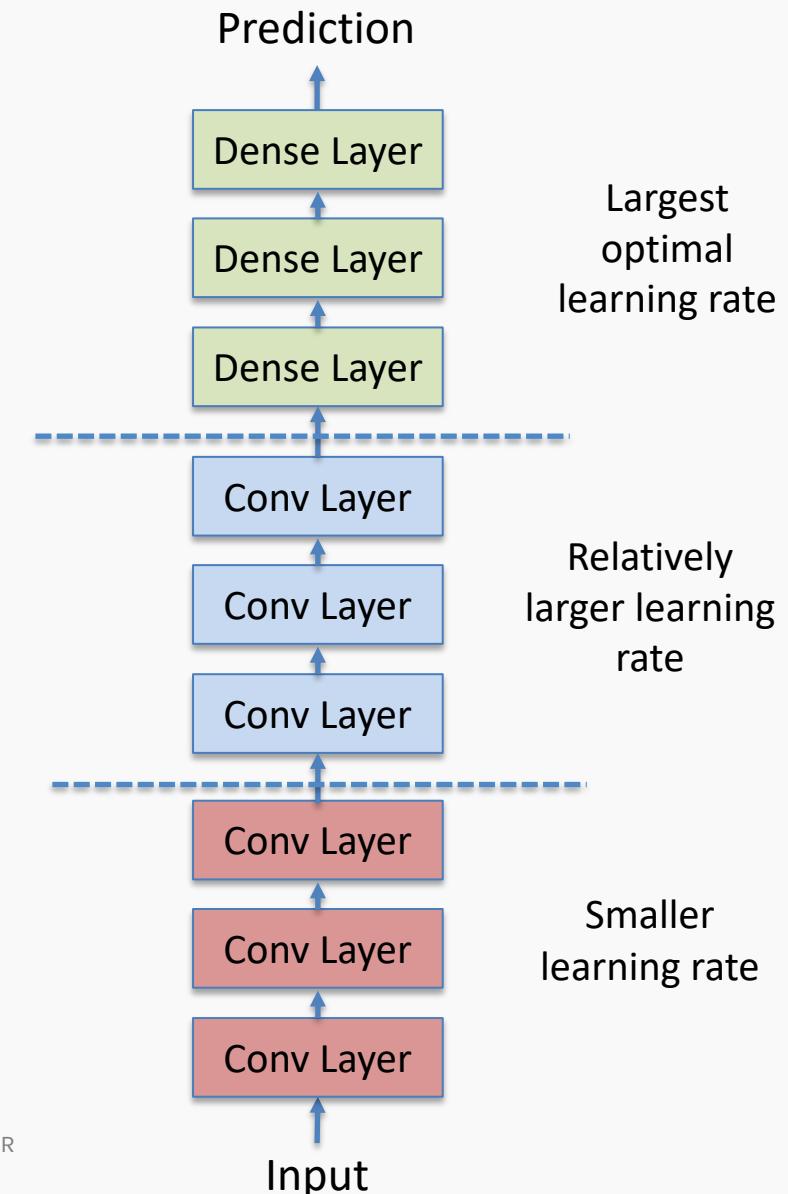
# Transfer Learning for Deep Learning: Differential Learning Rates

- A low learning rate can take a lot of time to train on the "later" layers. Since we trained the FC head earlier, we could probably retrain them at a higher learning rate.
- General Idea: **Train different layers at different rates.**



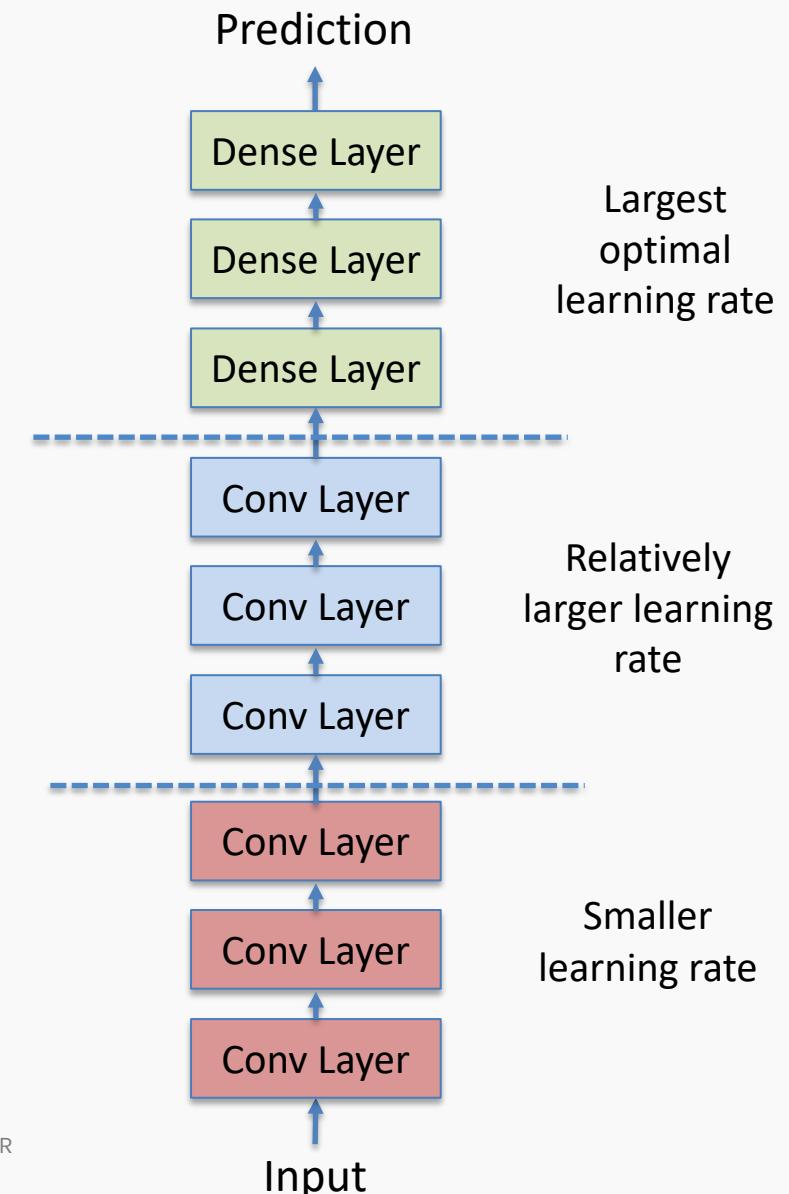
# Transfer Learning for Deep Learning: Differential Learning Rates

- A low learning rate can take a lot of time to train on the "later" layers. Since we trained the FC head earlier, we could probably retrain them at a higher learning rate.
- General Idea: **Train different layers at different rates.**
- Each "earlier" layer or layer group (the color-coded layers in the image) can be trained at 3x-10x smaller learning rate than the next "later" one.

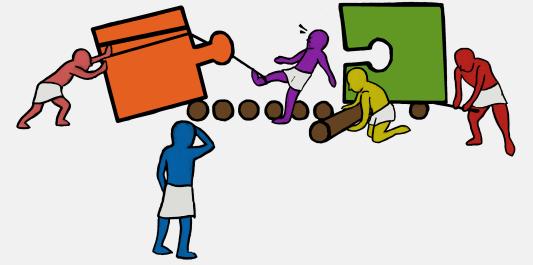


# Transfer Learning for Deep Learning: Differential Learning Rates

- A low learning rate can take a lot of time to train on the "later" layers. Since we trained the FC head earlier, we could probably retrain them at a higher learning rate.
- General Idea: **Train different layers at different rates.**
- Each "earlier" layer or layer group (the color-coded layers in the image) can be trained at 3x-10x smaller learning rate than the next "later" one.
- One could even train the entire network this way until we overfit and then step back some epochs.



# Exercise: Transfer Learning



The goal of this exercise is to use Transfer Learning to achieve near-perfect accuracy for a highly customized task. The task at hand is to distinguish images of people with Sunglasses or Hat.

